

CS 310 Operating Systems

Lecture 2: Operating System Concepts - 1

Ravi Mittal
IIT Goa

Acknowledgements !

- Contents of this class presentation have been taken from various sources. Thanks are due to the original content creators:
 - Class presentation: University of California, Berkeley: David Culler, Anthony D. Joseph, John Kubiatawicz, AJ Shankar, George Necula, Alex Aiken, Eric Brewer, Ras Bodik, Ion Stoica, Doug Tygar, and David Wagner.
- Book: Operating Systems: Principles and Practice (2nd Edition) Anderson and Dahlin, Volume 1

Read the following:

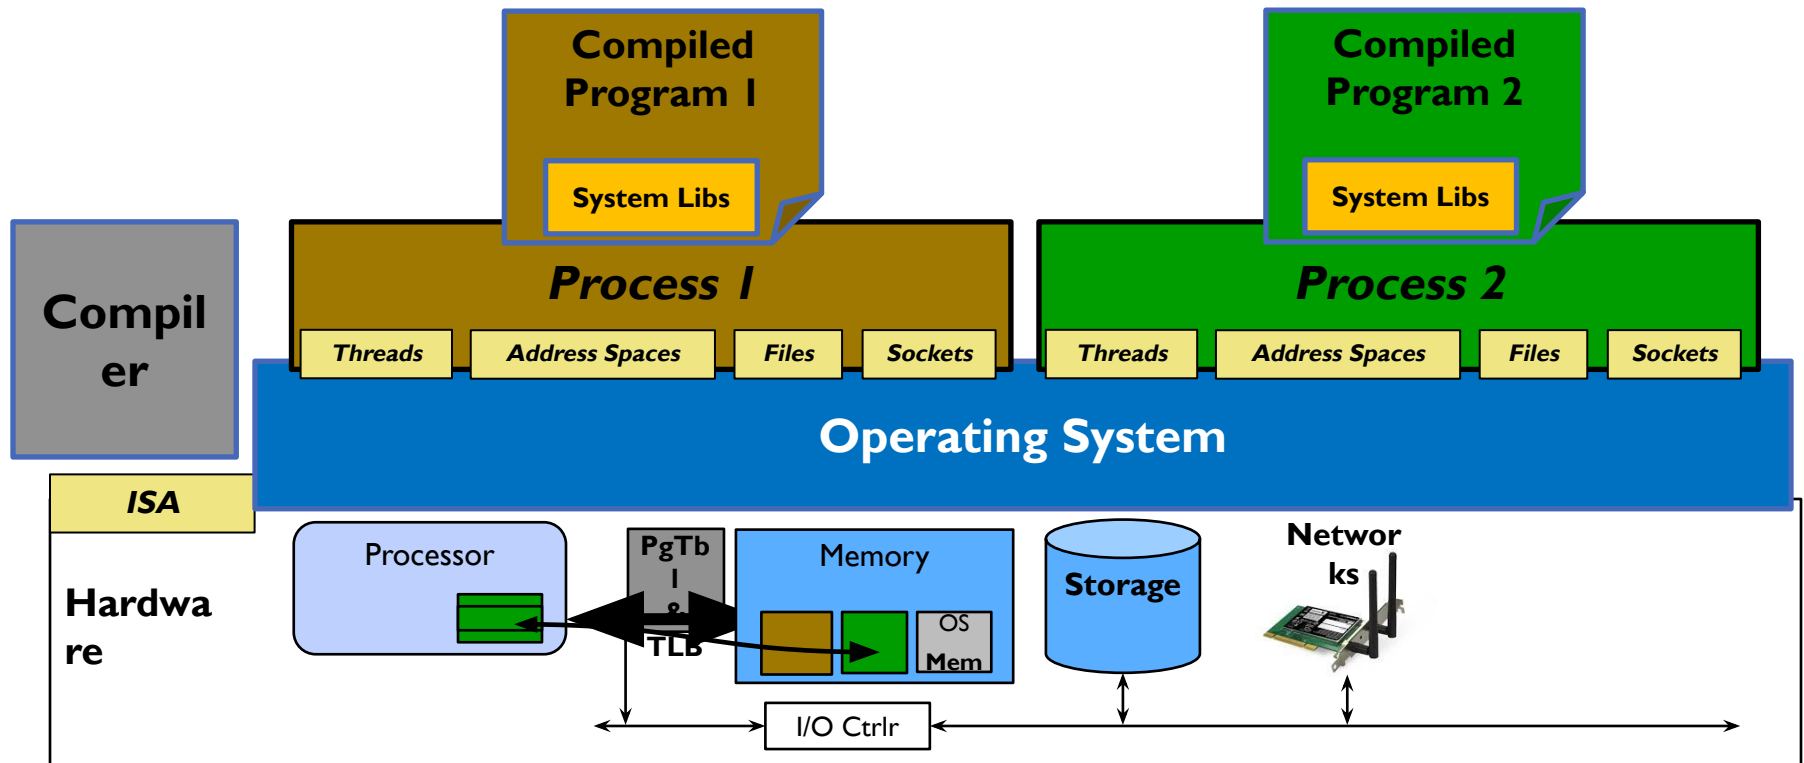
- Operating Systems: Principles and Practice (2nd Edition)
Anderson and Dahlin
 - Volume 1, Kernel and Processes
 - Chapter 2

We will study..

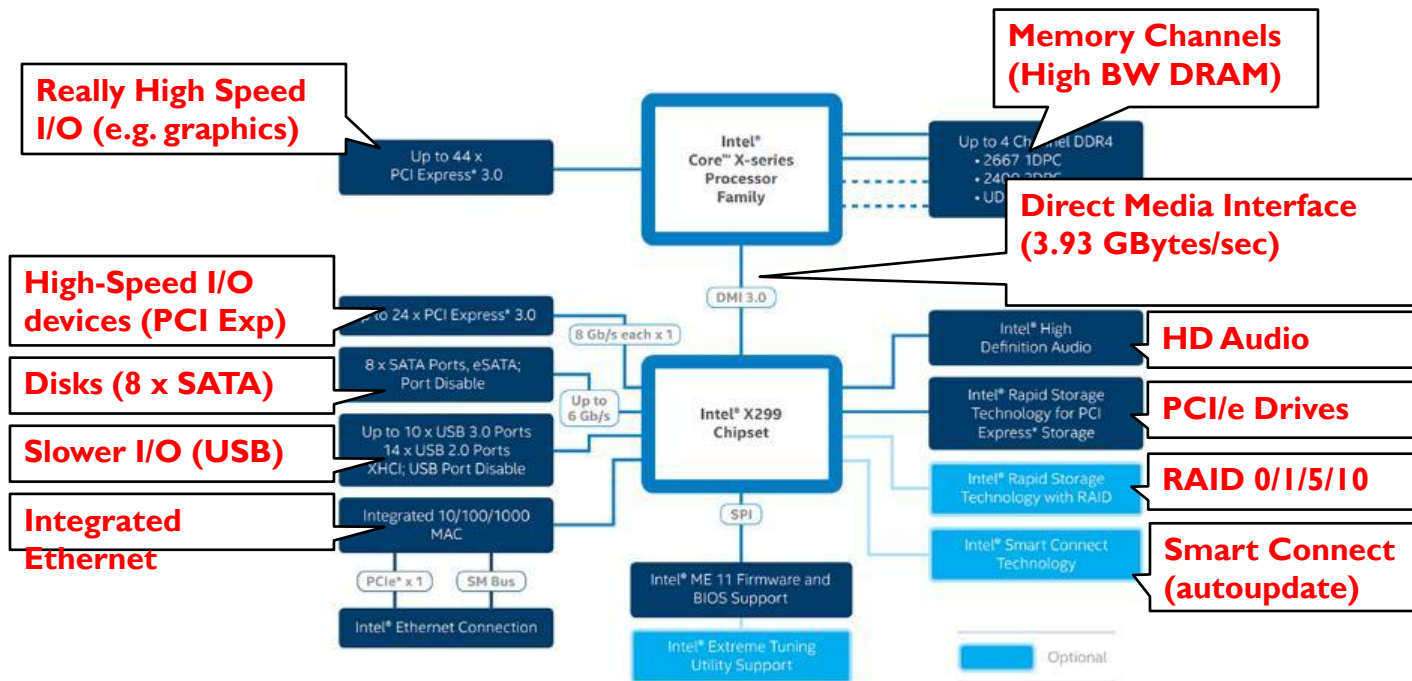
- Introduction
- The first concept: Processes
- The Second Concept: Address Space
- The Third Concept: Threads
- The Fourth Concept: Dual Mode of Operation

Last Class

Recall: OS Protection

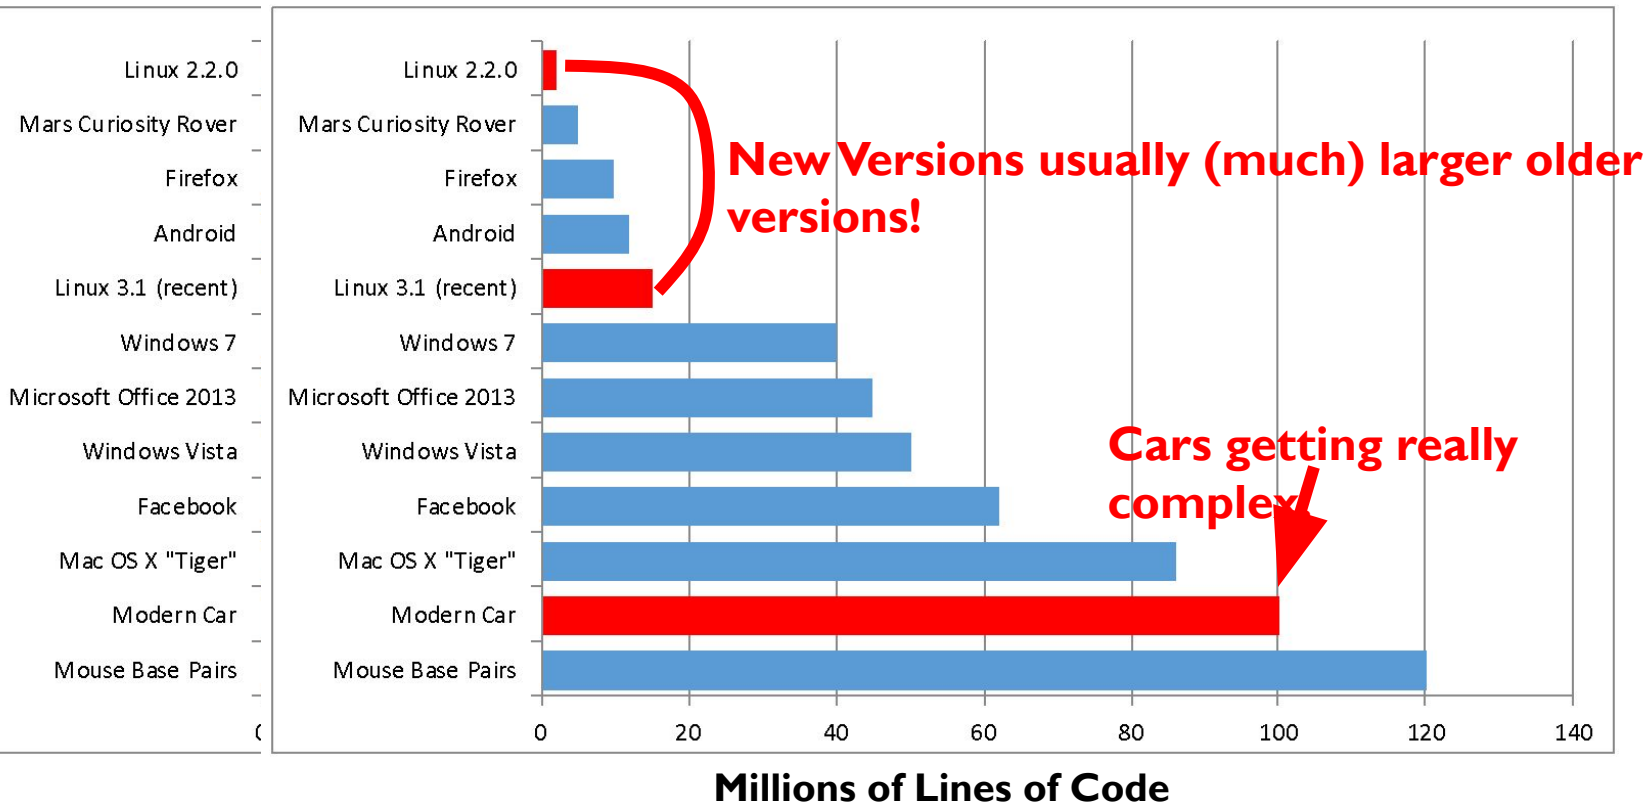


Recall: HW Functionality \Rightarrow great complexity!



Intel Skylake-X I/O Configuration

Recall: Increasing Software Complexity




(source <https://informationisbeautiful.net/visualizations/million-lines-of-code/>)

Linux 5.3 has approx. 27 million lines of code


Complexity leaks into OS if not properly designed

- Buggy device drivers
- Holes in security model or bugs in OS lead to instability and privacy breaches
 - Meltdown (2017)
 - Spectre (2017)



OS Abstracts Underlying Hardware to help Tame Complexity

- Processor → Thread
- Memory → Address Space
- Disks, SSDs, ... → Files
- Networks → Sockets
- Machines → Processes



Design Goals of an Operating System

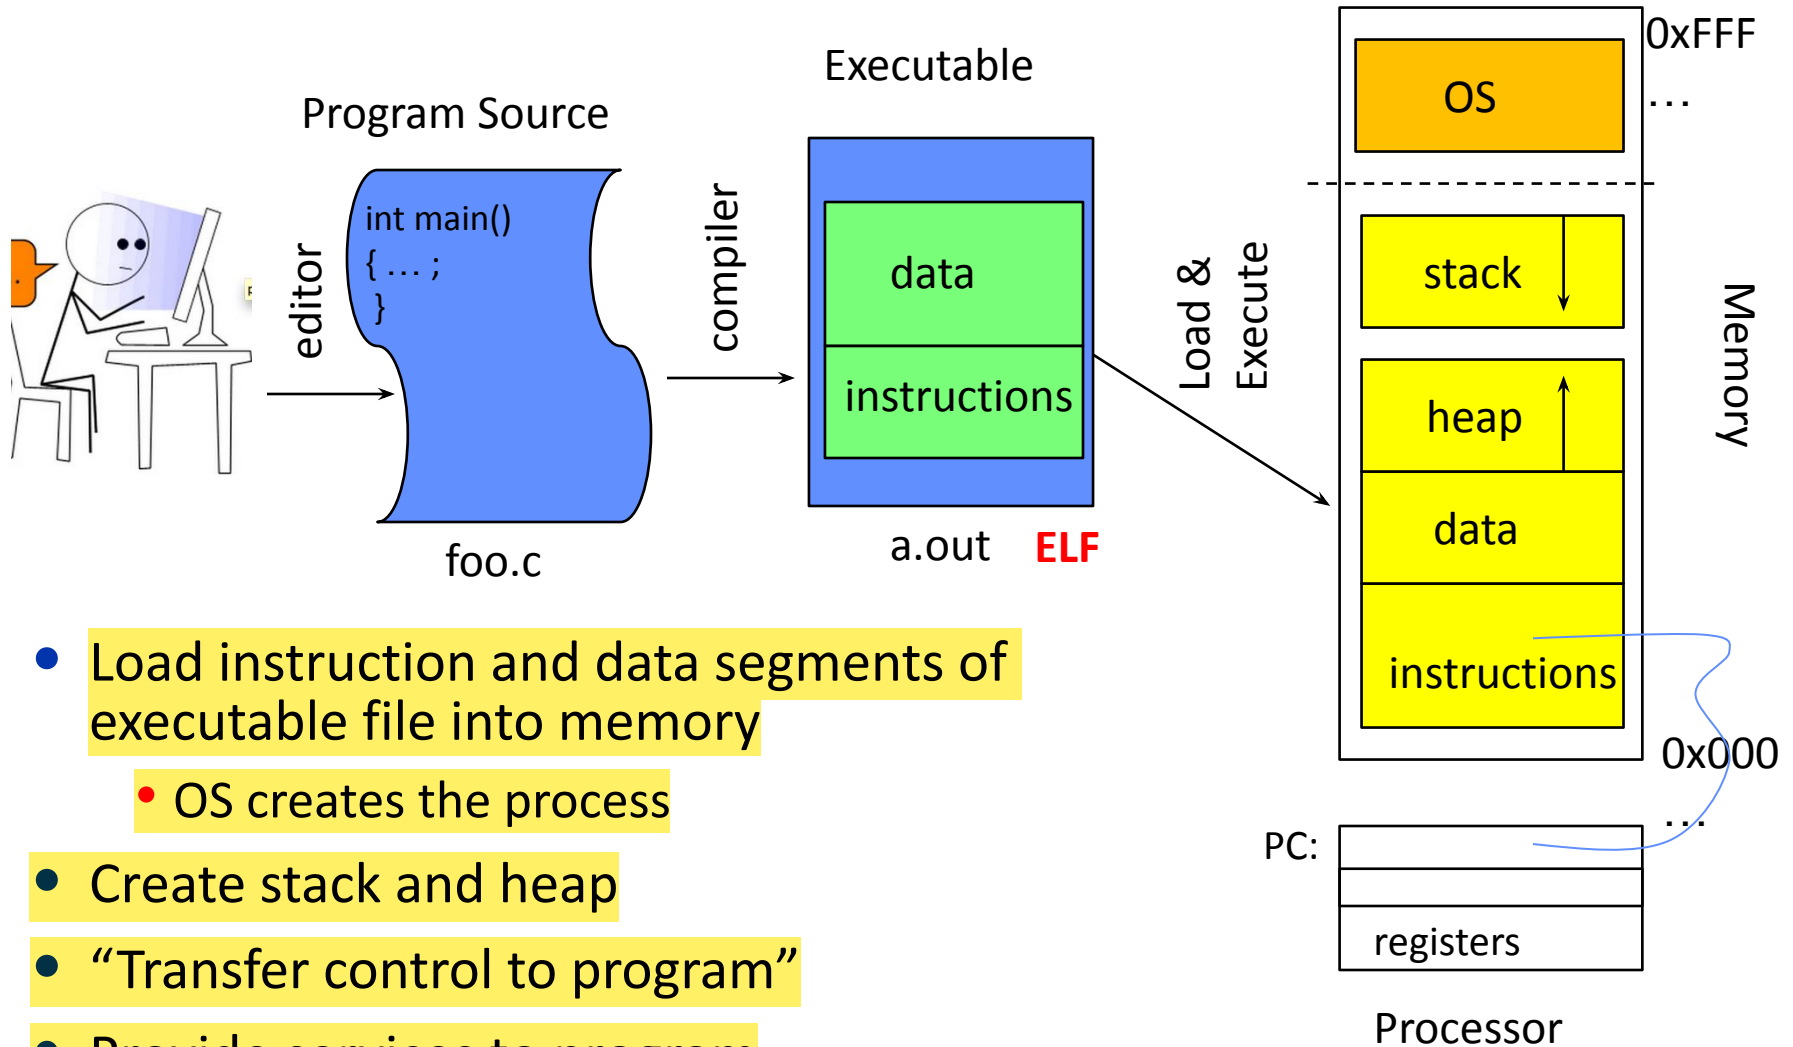
- Abstraction of the hardware resources
- Convenience of users
- Efficiency of usage of CPU, Memory etc
- Isolation between multiple processes

Four important concepts..

Four Fundamental OS Concepts

- **Process: an instance of a running program**
 - Protected Address Space + One or more Threads
- **Address space**
 - Set of memory addresses accessible to program (for read or write)
- **Thread: Execution Context**
 - Fully describes program state
- **Dual mode operation / Protection**
 - User / Kernel mode

OS Bottom Line: Run Programs



ELF: Executable and Linking Format

Operating System Needs are changing (1)

- Network Speeds:
 - 1980: 300 bps /\$
 - 2000: ~ 256 Kbps /\$
 - 2020: ~ 20 Mbps/\$
- So, in the past 40 years, the speed of networking has increased by a factor of ~ 67000 times
- We now have Internet connection of 1 Gbps

Operating System Needs are changing (2)

- Number of Cores / CPU
 - 1980: 1 core / CPU
 - 2000: 1 core / CPU
 - 2020: 8+ cores / CPU □ 64 cores / Server-CPU
- So, in the past 20 years, the number of cores have become 10 times or more
- How OS handles tasks on so many cores?
 - Process/job management
 - Power Management
 - Scheduling, Availability, Security

Operating System Needs are changing (3)

- **Cost per megaflop/sec:**
 - 1980: ~ \$100,000 /megaflop/sec
 - 2000: ~ \$25 /megaflop/sec
 - 2020: ~ \$0.20 /megaflop/sec
- So, in the past 40 years, the cost of per million operations has decreased by a factor of ~ 500,000 times

Operating System Needs are changing (4)

- **RAM Capacity B/\$:**
 - 1980: $\sim 2 \text{ KiB} / \$$
 - 2000: $\sim 2 \text{ MiB} / \$$
 - 2020: $\sim 2 \text{ GiB} / \$$
- So, in the past 40 years, the cost per byte of RAM has decreased by a factor of $\sim 1,000,000$ times

Operating System Needs are changing (5)

- **Ratio of Computers to Users**
 - 1980: ~ 100 users : 1 computer
 - 2000: ~ 1 user : 1 computer
 - 2020: ~ 1 user : many computers
- So, in the past 40 years, the number of users to computer ratio has increased by a factor of at least 200 times
- Paradigm shift in which OS are designed for evolving computing technology

Operating System Challenges

- Reliability
- Availability
- Security
- Privacy
- Portability
- Performance

OS System needs are changing

Legacy Needs

- Runs one application at a time
- Manage time quotas for many users
- Uses submit jobs and get results days later

Modern Needs

- Multiprogramming across many cores and many concurrent users
- Interactive jobs; completing all jobs asap
- Optimize for users not for computer's resource time

Future Needs

- Manage and use an ever-increasing number of processors/computers
- Peta-scale storage, data centers, cloud
- Optimize for diversity

The first OS concepts - Process

What is a process ?


- A process is a program in execution
- The process is the OS's abstraction for execution
- Execution environment with restricted rights
 - (Protected) Address Space with One or More Threads
 - Owns memory (address space), file descriptors, sockets
 - Encapsulate one or more threads sharing process resources
- Simplest (classic) case: a sequential process
 - A single thread of execution (an abstraction of the CPU)
- A sequential process is
 - The unit of execution
 - The unit of scheduling
 - The dynamic (active) execution context
 - vs. the program – static, just a bunch of bytes

Process – Two key abstractions

- Process is not the binary source code
 - It is an instance of running program
- Process provides two key abstractions
 - Memory
 - Each process assumes entire system memory to itself
 - Execution
 - Provides abstraction of continued operation
 - There may be hundreds of processes in a system
 - Each process gets impression of continuous operation

What is a Process?

- Why processes?
 - Protected from each other! OS Protected from them
 - Processes provides memory protection




Protection and Isolation

- Processes provide protection and isolation
 - Reliability: bugs can only overwrite memory of process they are in
 - Security and privacy: malicious or compromised process can't read or write other process' data
- Mechanisms:
 - Address translation: address space only contains its own data

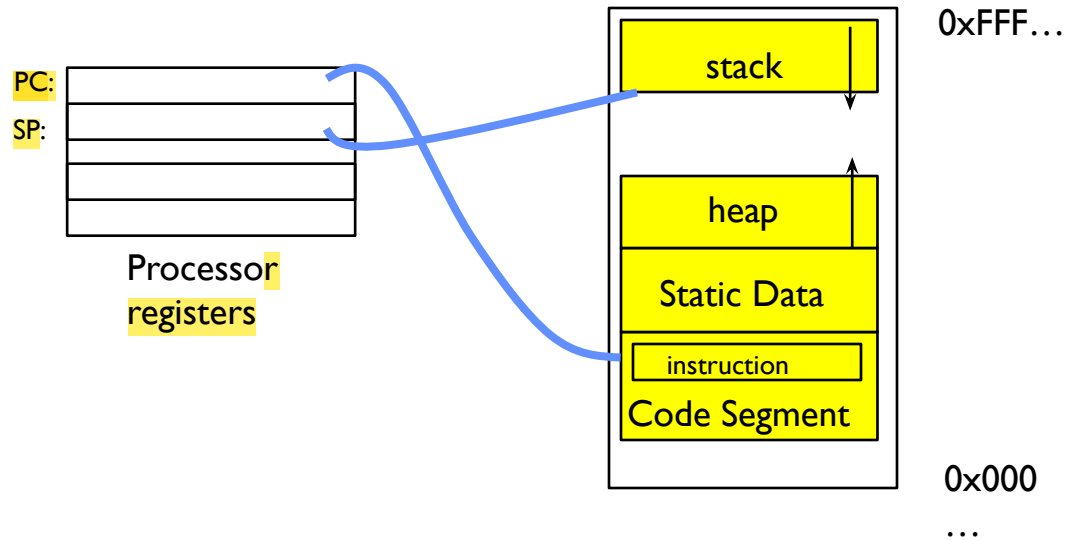
The Second OS concepts – Address Space

Second OS Concept: Address Space

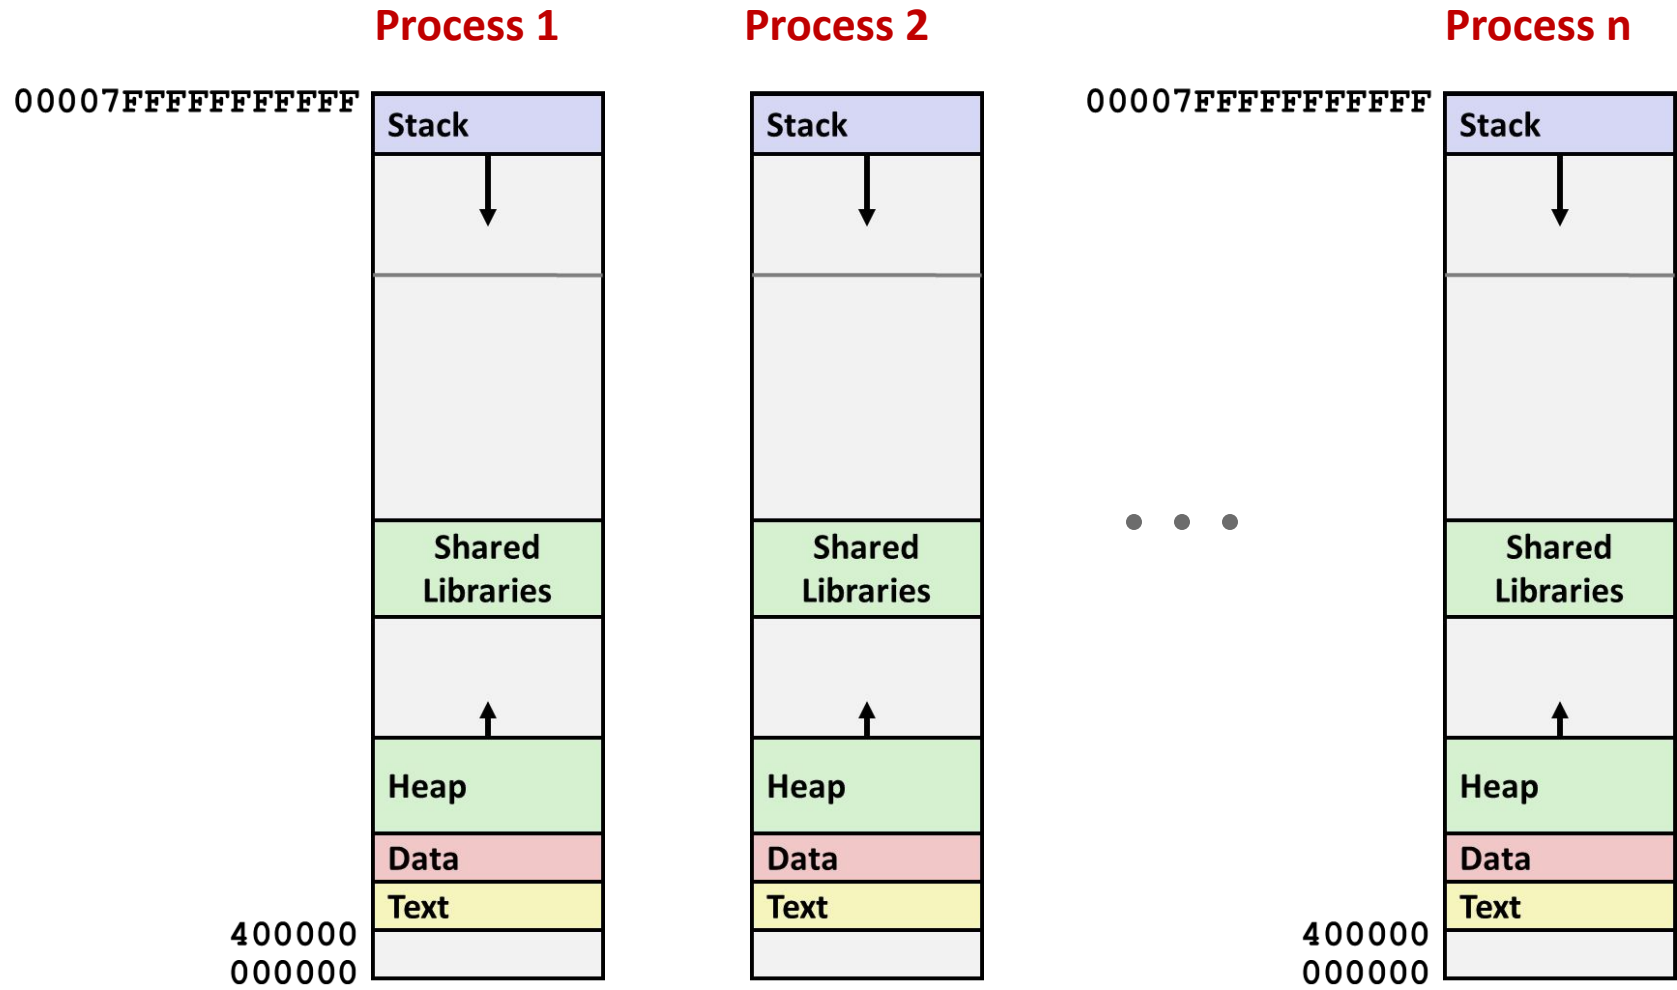
- Address space \Rightarrow the set of accessible addresses + state associated with them:
 - For 32-bit processor: $2^{32} = 4$ billion (10^9) addresses
 - For 64-bit processor: $2^{64} = 18$ quintillion (10^{18}) addresses
- The address space of a process contains all of the memory state of the running program
- What happens when you read or write to an address?
 - Perhaps acts like regular memory
 - Perhaps causes I/O operation
 - (Memory-mapped I/O)
 - Perhaps causes exception (fault)
 - Communicates with another program



Address Space: In a Picture

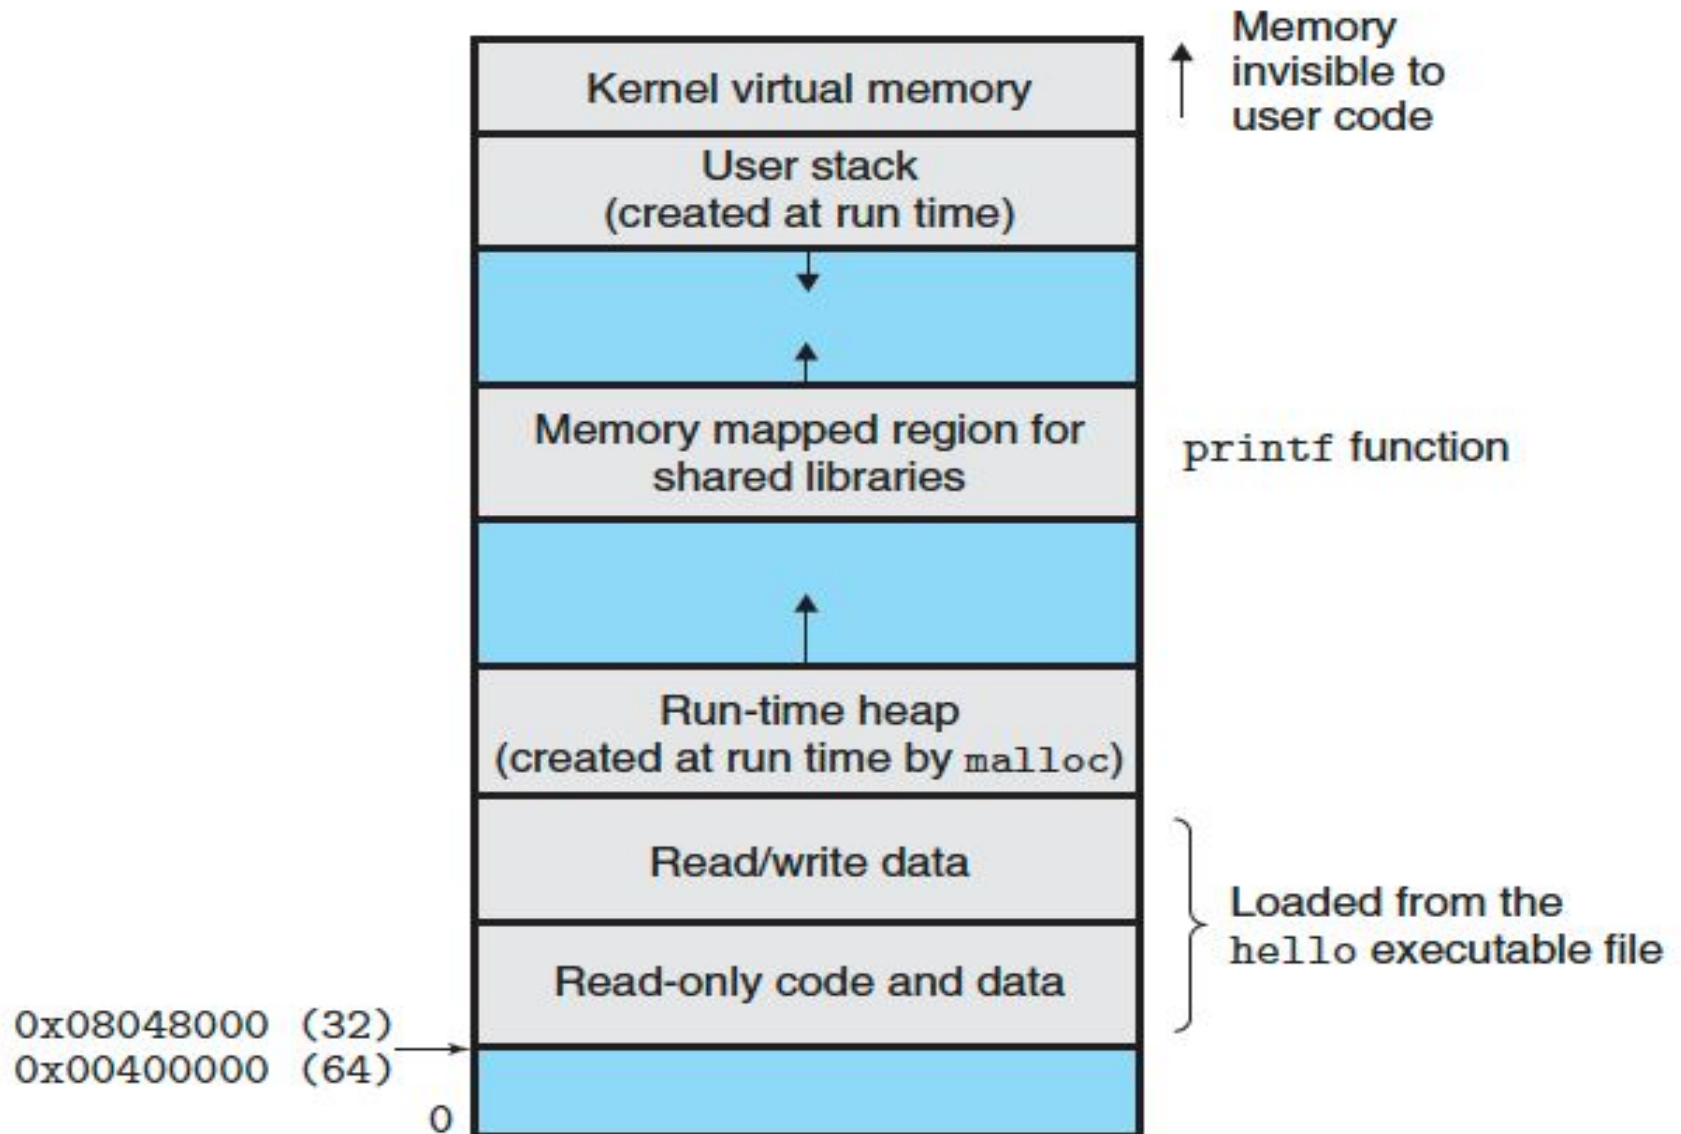


Recall: Each process can have the same address space



Solution: Virtual Memory

Address Space



Linux Memory layout of a process

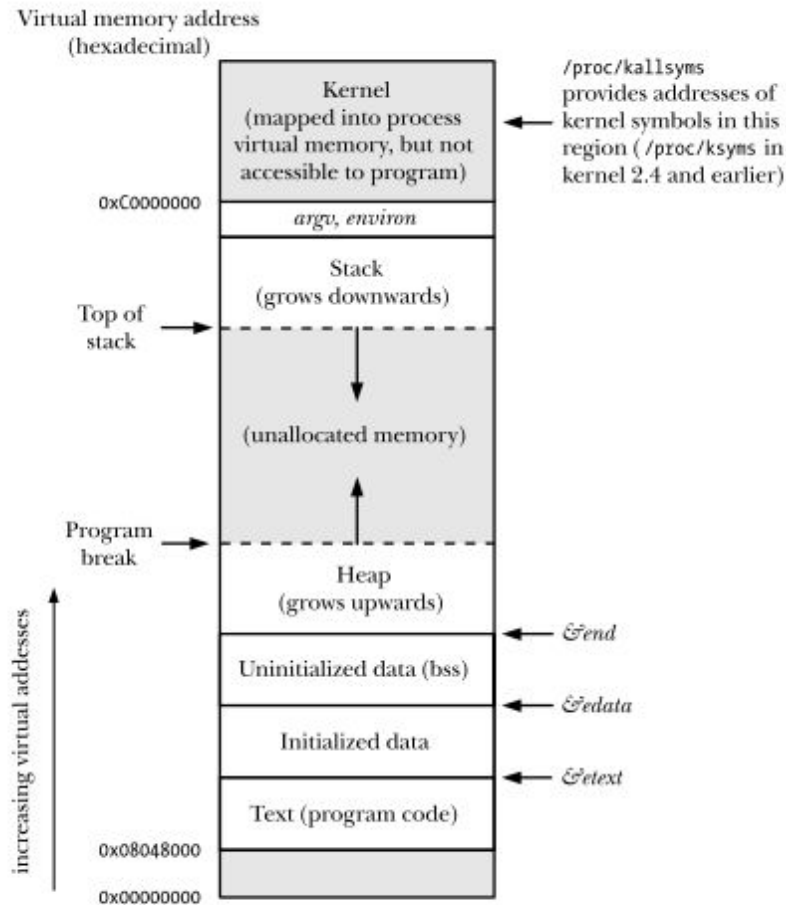



Figure 6-1: Typical memory layout of a process on Linux/x86-32

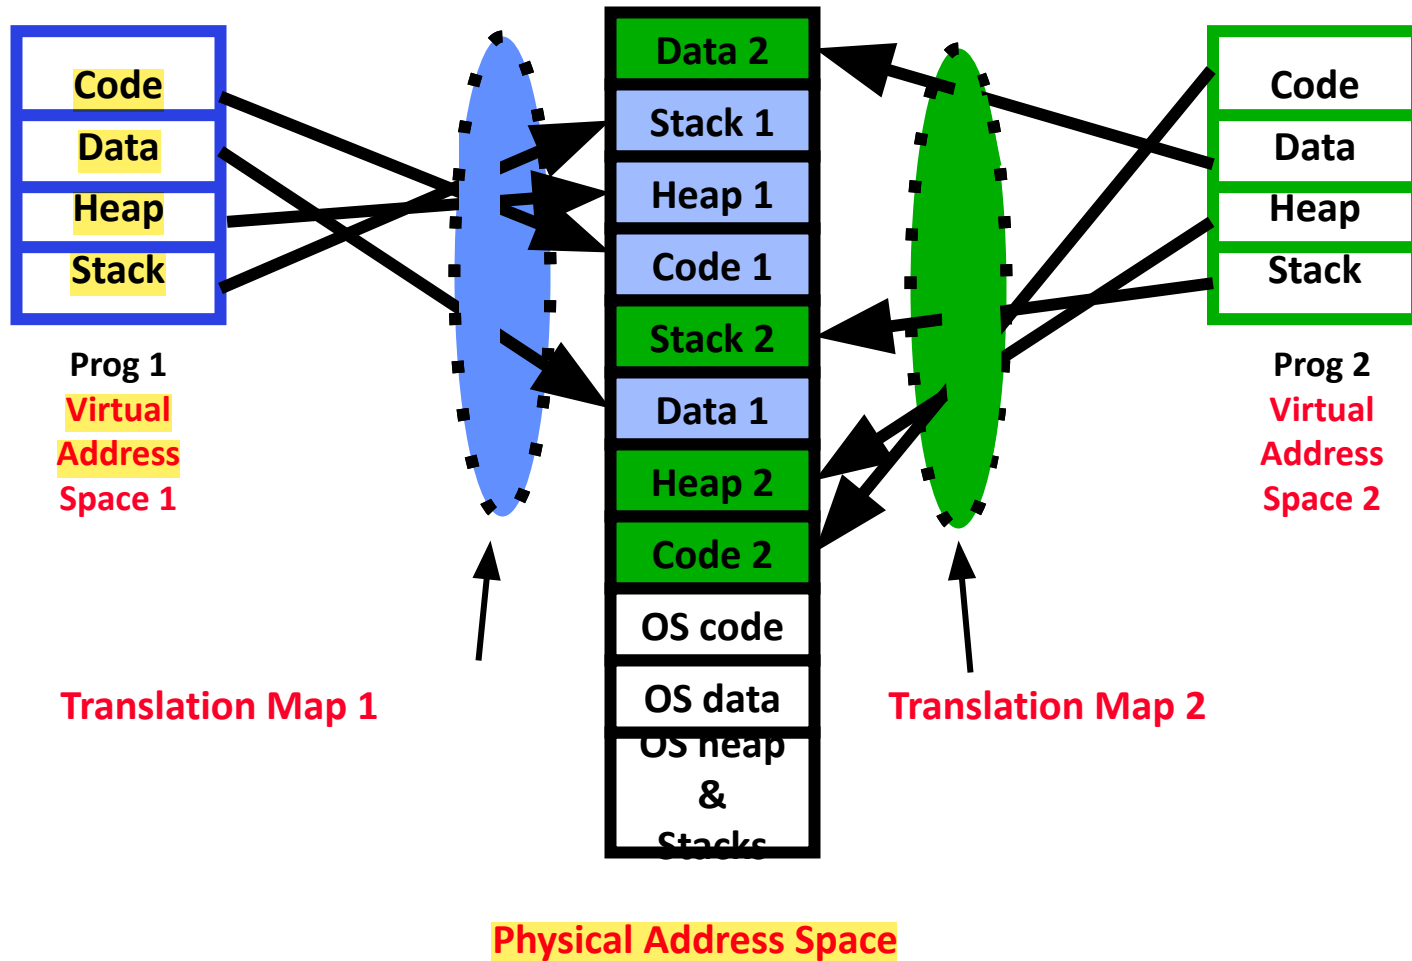
- Unix: Kernel space is mapped in high - but inaccessible to user processes

Address Space Translation

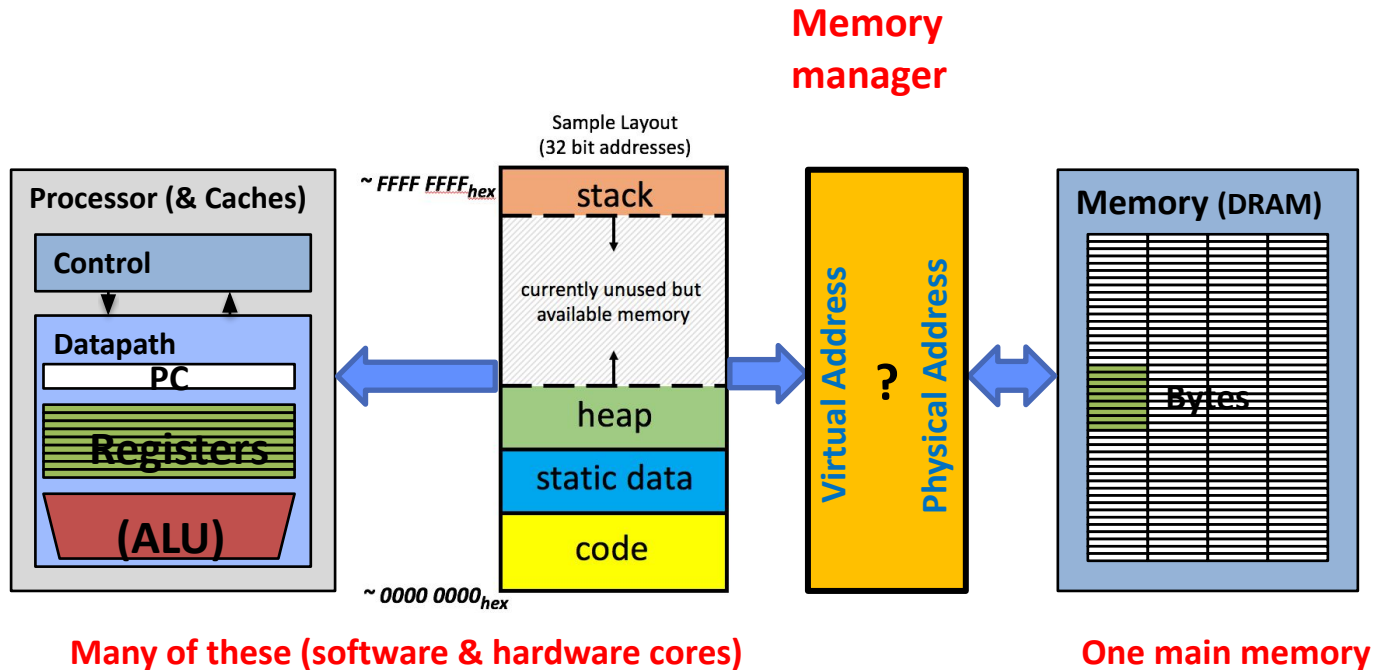
- Program operates in an address space that is distinct from the physical memory space of the machine



Translation through Page Table



Recall: Processors operate in Virtual Address Space



- Each Process uses it's own address space
- Processes use virtual addresses
- Many processes, all using same (conflicting) addresses
- Memory uses physical addresses (also, e.g., 0 ... 0xffff,ffff)
- **Memory manager maps virtual to physical addresses**

Lecture Summary

- The four important concepts of Operating Systems include
 - Process
 - Address Space
 - Thread
 - Dual mode of operation / Protection
- Process is a program in execution
- To create a process, the OS does the following
 - Creates space in MM to hold
 - Code (Text), Stack, Heap, and Static
 - Starts running a process by executing the first instruction pointed by the PC
- Address space \Rightarrow the set of accessible addresses + state associated with them