

CS310 Operating Systems

Lecture 6: System Calls


Ravi Mittal
IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, Profs. Natacha Crooks and Anthony D. Joseph, University of California, Berkeley
 - Book: Operating System Concepts, 10th Edition, by Silberschatz, Galvin, and Gagne

We will study..

- Kernel Stack – another system data structure
- System Calls
- Application Programming Interfaces
- A tree of processes in Linux

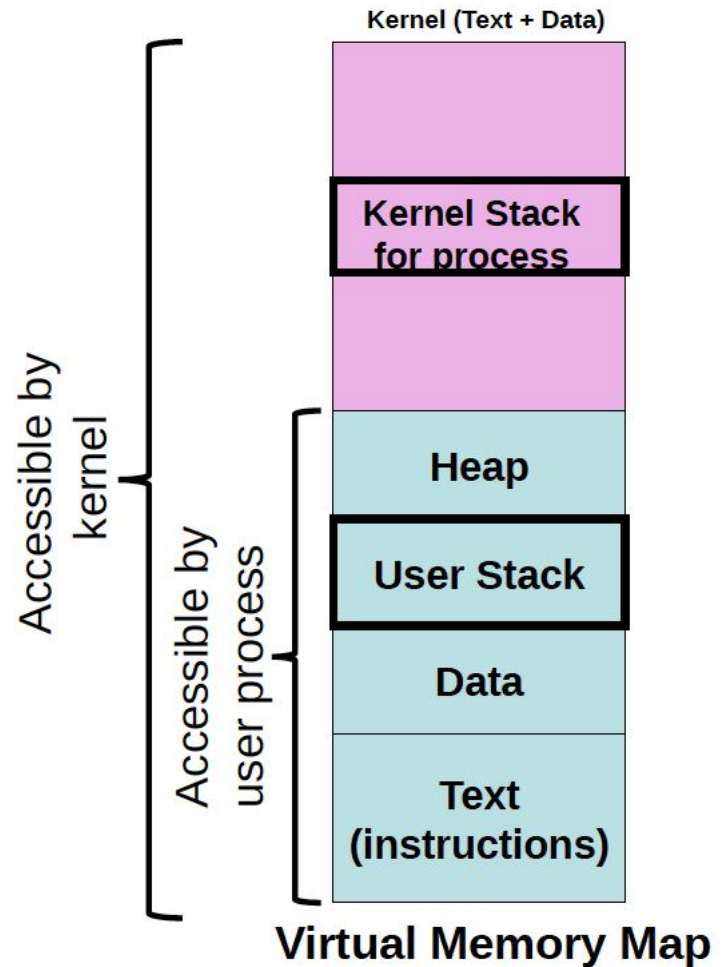


Read the following:

- Book: Operating Systems: Principles and Practice (2nd Edition) Anderson and Dahlin
 - Volume 1, Kernel and Processes
 - Chapter 2.4: About Kernel Stack
- Book: Operating Systems: Three Easy Pieces, by Remzi and Andrea Arpaci-Dusseau
 - Chapter 5: Process APIs

Kernel Stack

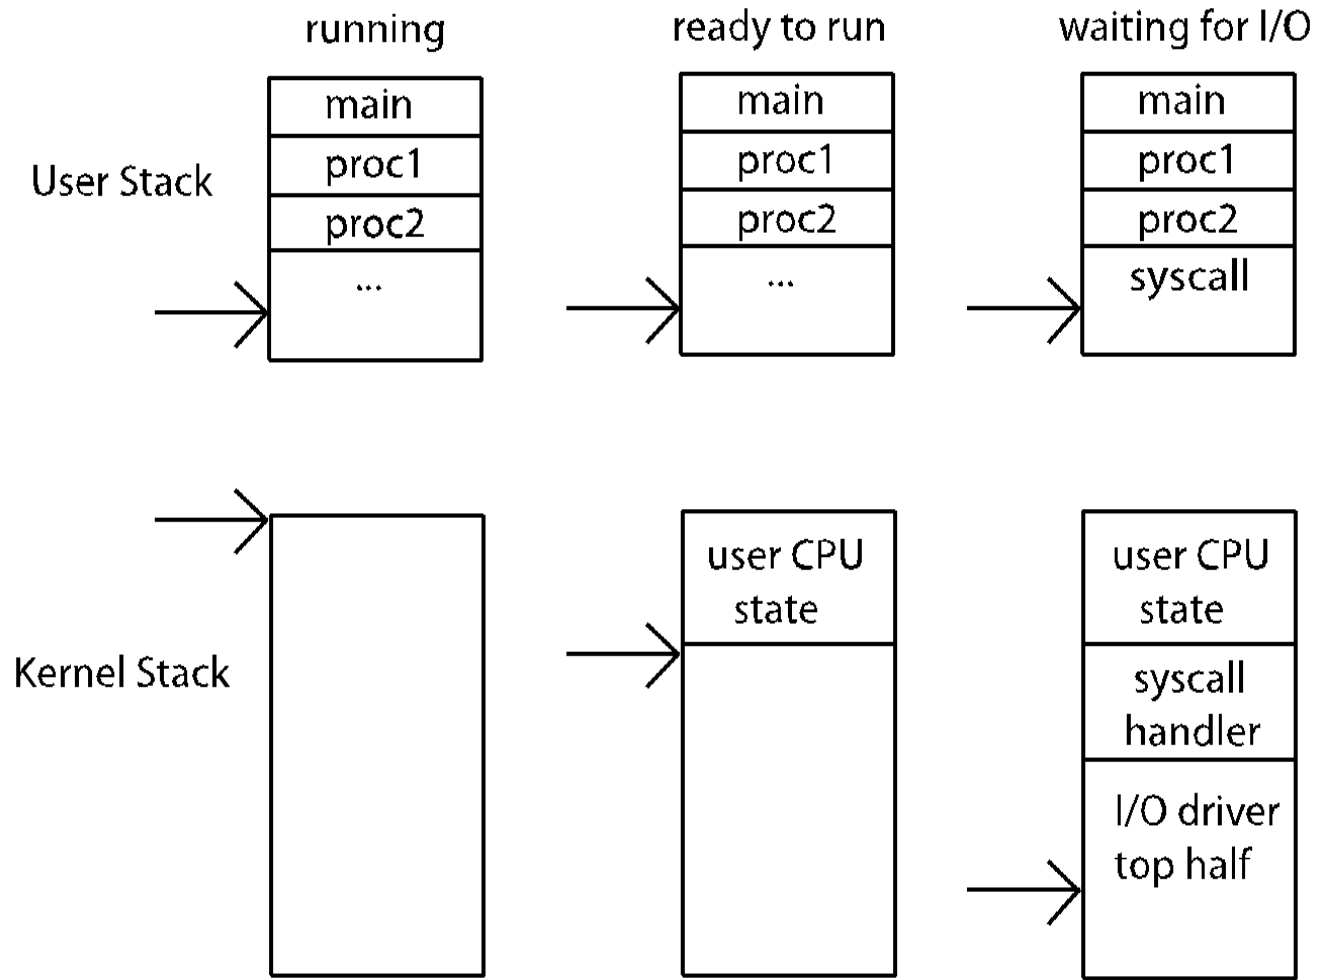
- For each user level process, most OSs allocate two stacks for
 - User stack
 - Kernel stack
- Kernel needs space to work
 - Can't put anything on the user stack (?)



Kernel Stack

- This simplifies switching to new process
- When a process is running in user mode, the kernel stack is empty
- The pointer to process's kernel stack is stored in PCB
- If the process is running on the processor on kernel mode, due to an interrupt or system call, it's kernel stack contains
 - Registers (saved from the suspended process)
 - PC, Stack pointer
 - Flags (Processor status word in x86)
- A special privileged hardware register (stack pointer) points to a region of kernel memory called the **Kernel stack**

Kernel Stack




Kernel Stack contents

- When the process is running, does the kernel stack have anything useful on it?
 - You might think yes – it called into the user program.
 - Actually no
- When a process is ready to run, but not running
 - It has its user stack as before,
 - We need a place to store the state that had been in the CPU when it stopped running
- When a process is waiting for I/O
 - The kernel stack contains the suspended computations
 - Eg Syscall handler state, I/O driver state etc
 - User CPU state: registers, PC, SP etc

System Calls


Function Calls vs System Calls

Function Call:




- Caller and callee in the same process
 - Same user
 - Same “domain of trust”

System Call



- OS is trusted; User process is not
- OS code runs privileged with complete access to all system resources

System Calls (“Syscalls”)



System Calls

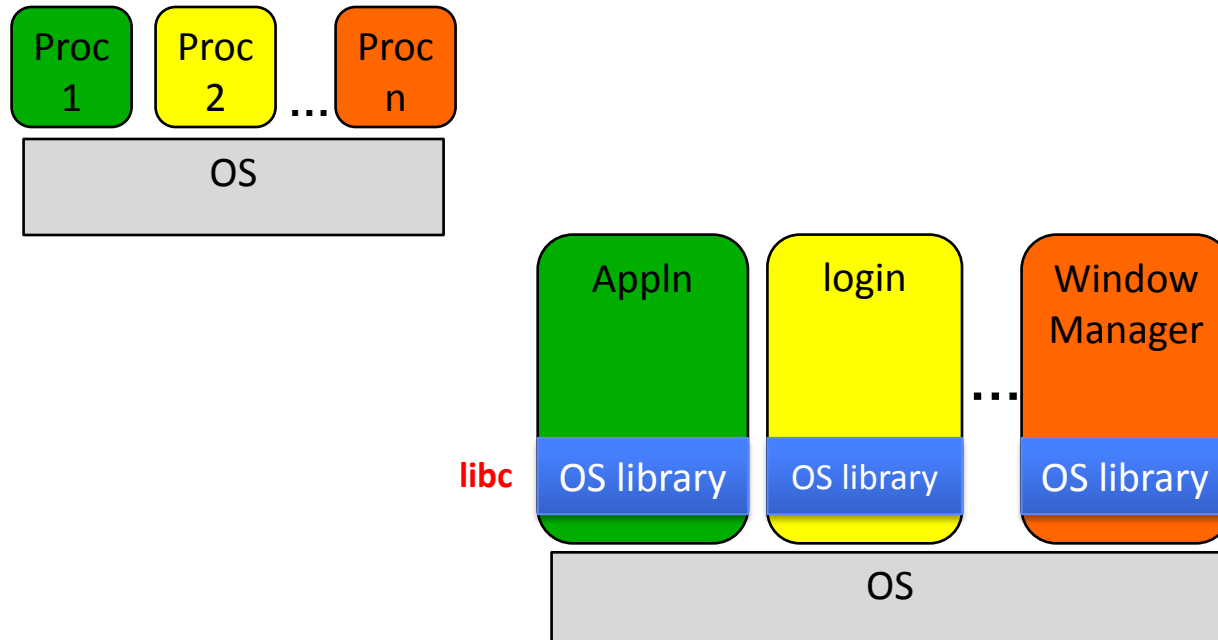
- API : Application Programming Interface
 - Function available to write user programs
- System calls are the interface of the OS for
 - Processes
 - Creating, existing, waiting, and terminating
 - Memory
 - Allocation and deallocation
 - Files and Folders
 - Opening, reading, writing, closing
 - Inter Process Communication
- We will study process system calls in later lectures

Should we rewrite programs for each OS?

- **POSIX API**: a standard set of system calls that an OS must implement
 - Programs written to the POSIX API can run on any POSIX compliant OS
 - Most modern OSes are POSIX compliant – Ensures program portability
- **Program language libraries** hide the details of invoking system calls
 - The **printf** function in the C library calls the write system call to write to screen
 - User programs usually do not need to worry about invoking system calls

C  **libc**  **syscall**

OS Library Issues Syscalls



APIs

- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use

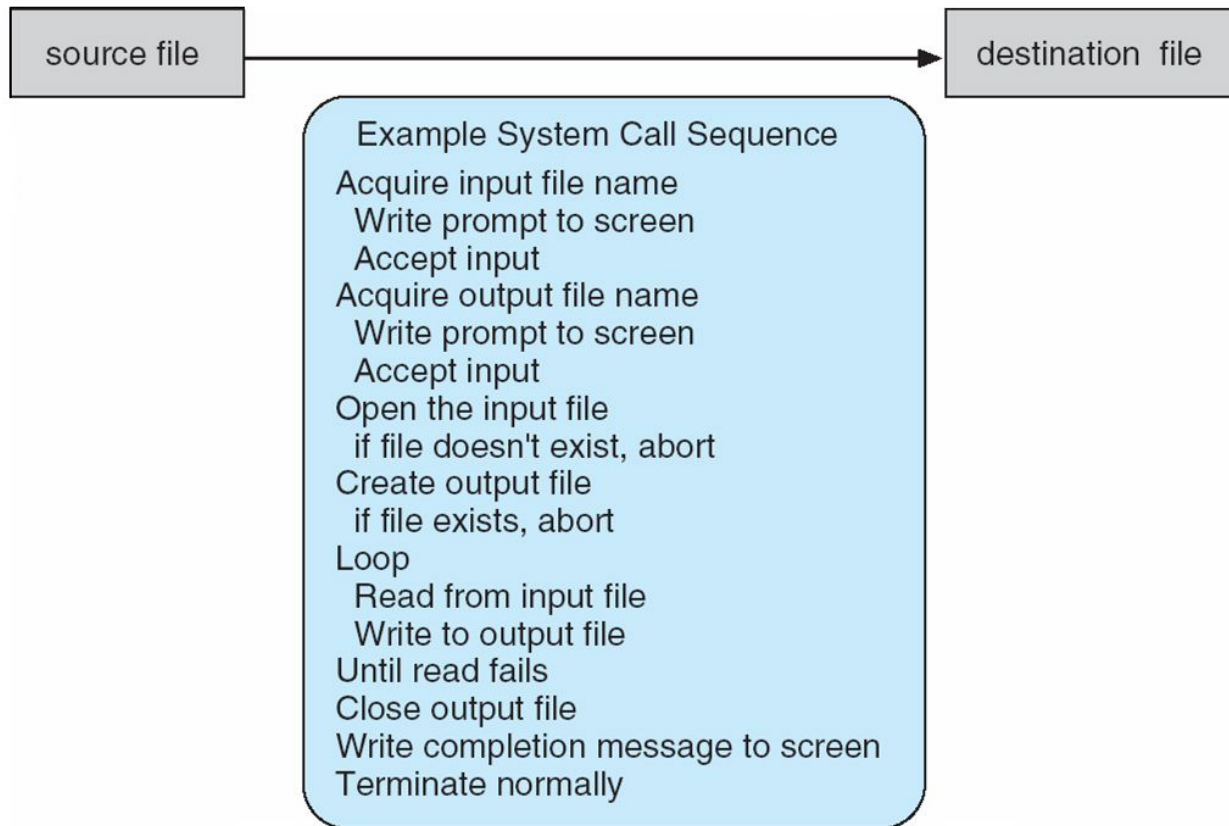
Multiple Sys Calls are needed ..

- Consider following Unix command
`cp in.txt out.txt`
- This command copies the input file `in.txt` to the output file `out.txt`
- This seemingly simple command requires a large number of system calls
 - Write prompting message on terminal
 - To read keyboard input
 - To write text onto the terminal
 - Open input file and open output file
 - Copy text from input file to output file – in loop
 - Possible error message if files don't exist
 - Closing both files on completion

• And so on

Example of System Calls

- System call sequence to copy the contents of one file to another file



Application APIs

Application Programming Interface (API)

- A simple program uses **hundreds of system calls** in a short time
- Application Programmers use APIs
 - API specifies
 - a set of functions that are available to an application programmer
 - Parameters passed to each function and return values it can expect
- Three most common APIs are **Win32 API** for Windows, **POSIX API** for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and **Java API** for the Java virtual machine (JVM)
- A programmer accesses an API via a library of code provided by the operating system
- **libc**: For Linux and Unix, library for programs written in C

APIs vs System Call

- **System Call** is an explicit request to the kernel made via a software interrupt
- API is a function definition that specifies how to obtain a given service
 - An OS has several libraries of functions that provide APIs to programmers:
 - Some APIs provide services to users in User Mode
 - Other APIs use **system calls** (in turn use kernel mode)
 - Wrapper over system call with added functionality
- Example:
 - In Linux, the `malloc()`, `calloc()`, and `free()` APIs are implemented in the `libc` library
 - The code in this library keeps track of the allocation and deallocation requests and uses the `brk()` system call to enlarge or shrink the process heap

Examples of API calls and System Calls

C Library Call:

fopen
fclose
getc/putc
fread/fwrite
scanf/printf
fprintf
fseek

rand

Linux System Call:

open
close
read/write

lseek

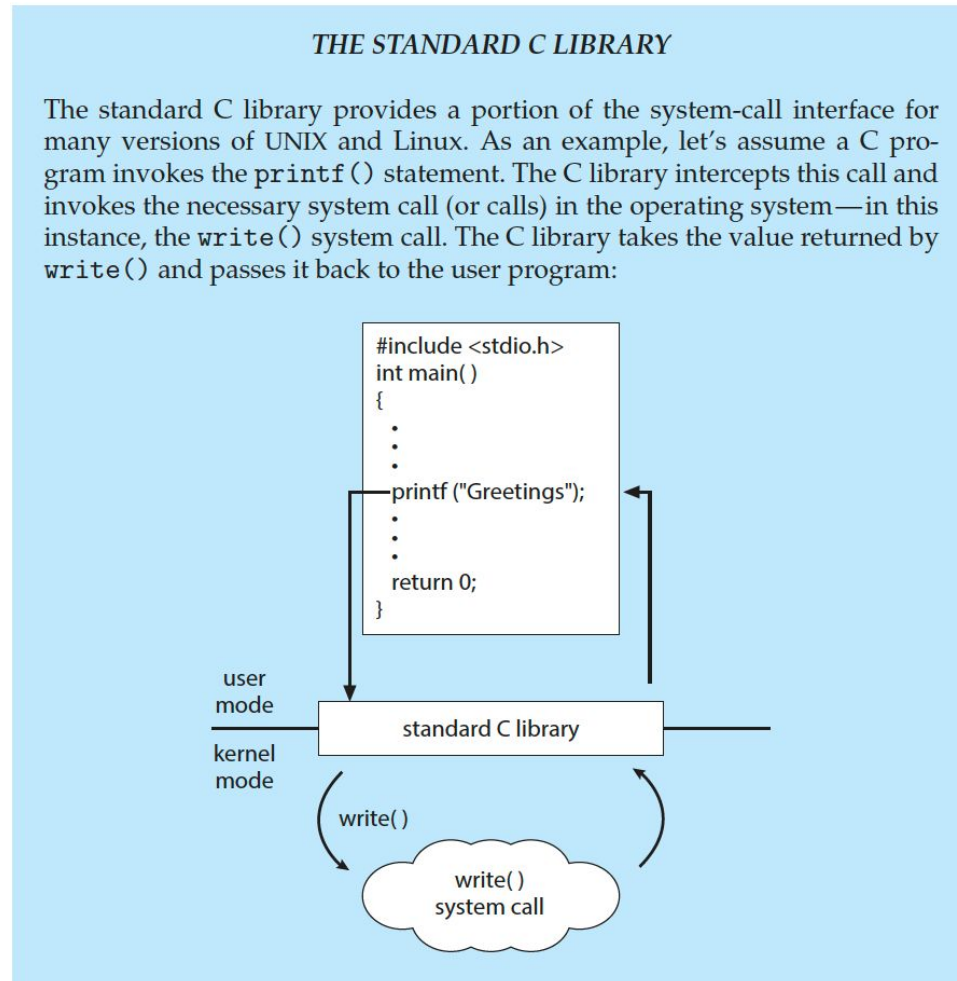
Win32 System Call:

CreateFileA
CloseHandle
ReadFile / WriteFile

SetFilePointer

Standard C Library Example

- C program invoking `printf()` library call, which calls `write()` system call



Advantages of using APIs instead of System Calls

- Program Portability
 - Application programmer need not worry about processor and OS?
- Simple to use compared to system calls
- Standard `read` API is described as

```
#include <unistd.h>
```

<code>ssize_t</code>	<code>read</code>	<code>(int fd, void *buf, size_t count)</code>
return value	function name	parameters

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer into which the data will be read
- `Size_t count`—the maximum number of bytes to be read into the buffer

Python Code

Python Code:

Python: **open(...)**



Python is written in C ("CPython"), making a call to the C library calls...

C++: **fopen(...)**

When compiled on
Win32 system...



SysCall: **CreateFileA(...)**

When compiled on a
Linux system....



SysCall: **open(...)**

System Call Implementation

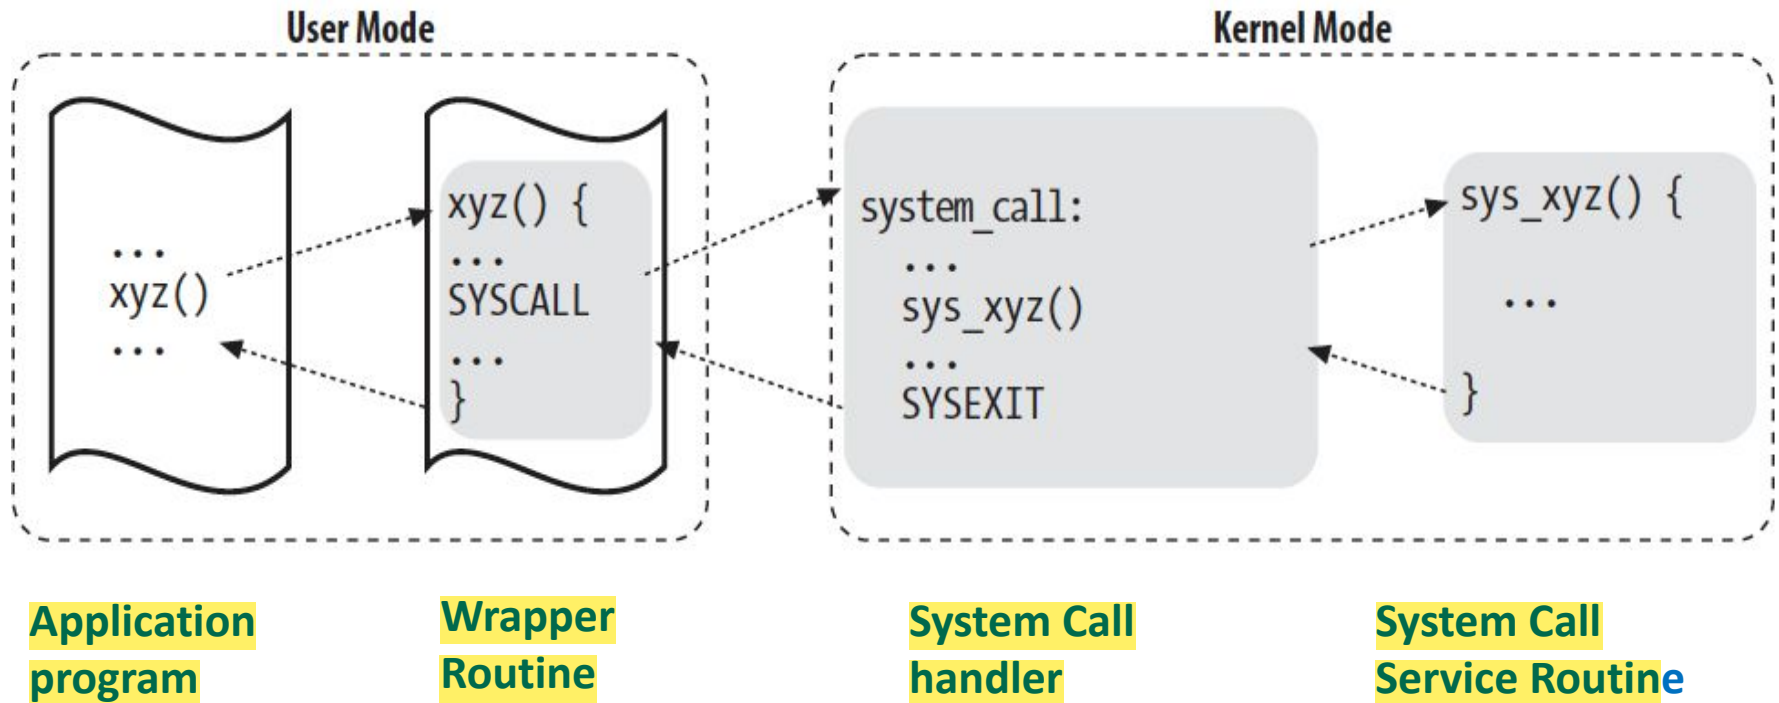
System Call Implementation

- Typically, a number is associated with each system call
 - **System-call interface** maintains a table indexed according to these numbers
- The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call
 - Most details of OS interface hidden from programmer by API
 - Managed by run-time support library (set of functions built into libraries included with compiler)

System Call Implementation

- When a process makes a system call, the processor switches to kernel mode
- In Kernel mode it runs System Call handler
- User process passes a parameter called system call number and other parameters
- System call handler is similar to other exception handler (will study it later). It performs the following actions:
 - Saves contents of registers in Kernel Mode Stack
 - Handles the system call by invoking a function called system call service routine
 - Exit from handler: Regs are loaded with values saved in Kernel stack ; CPU is switched back to user mode

Invoking a system call from Applications



Types of System Calls

Types of System Calls

- Process control
- File Management
- Device Management
- Information Management
- Communications
- Protection

System Calls in Windows and Unix

EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

The following illustrates various equivalent system calls for Windows and UNIX operating systems.

	Windows	Unix
Process control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File management	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communications	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Process Control - System Calls

- create process, terminate process
- end, abort
- load, execute
- get process attributes, set process attributes
- wait for time
- wait event, signal event
- allocate and free memory
- Dump memory if error
 - **Debugger** for determining **bugs, single step** execution
- **Locks** for managing access to shared data between processes – `acquire_lock()` and `release_lock()`

File and Device Management - System Calls

- File management

- create file, delete file
- open, close file
- read, write, reposition
- get and set file attributes

- Device management

- request device, release device
- read, write, reposition
- get device attributes, set device attributes
- logically attach or detach devices

Information and Communication - System Calls

- Information maintenance
 - get time or date, set time or date
 - get system data, set system data
 - get and set process, file, or device attributes
- Communications
 - create, delete communication connection
 - **message passing model** - send, receive messages to **host name** or **process name**
 - From **client** to **server**
 - **Shared-memory model** create and gain access to memory regions

Protection - System Calls

- Protection
 - Control access to resources
 - Get and set permissions
 - Allow and deny user access

Linux System Call

http://www.cheat-sheets.org/saved-copy/Linux_Syscall_quickref.pdf

- Linux Kernel 3.7 has 393 system calls

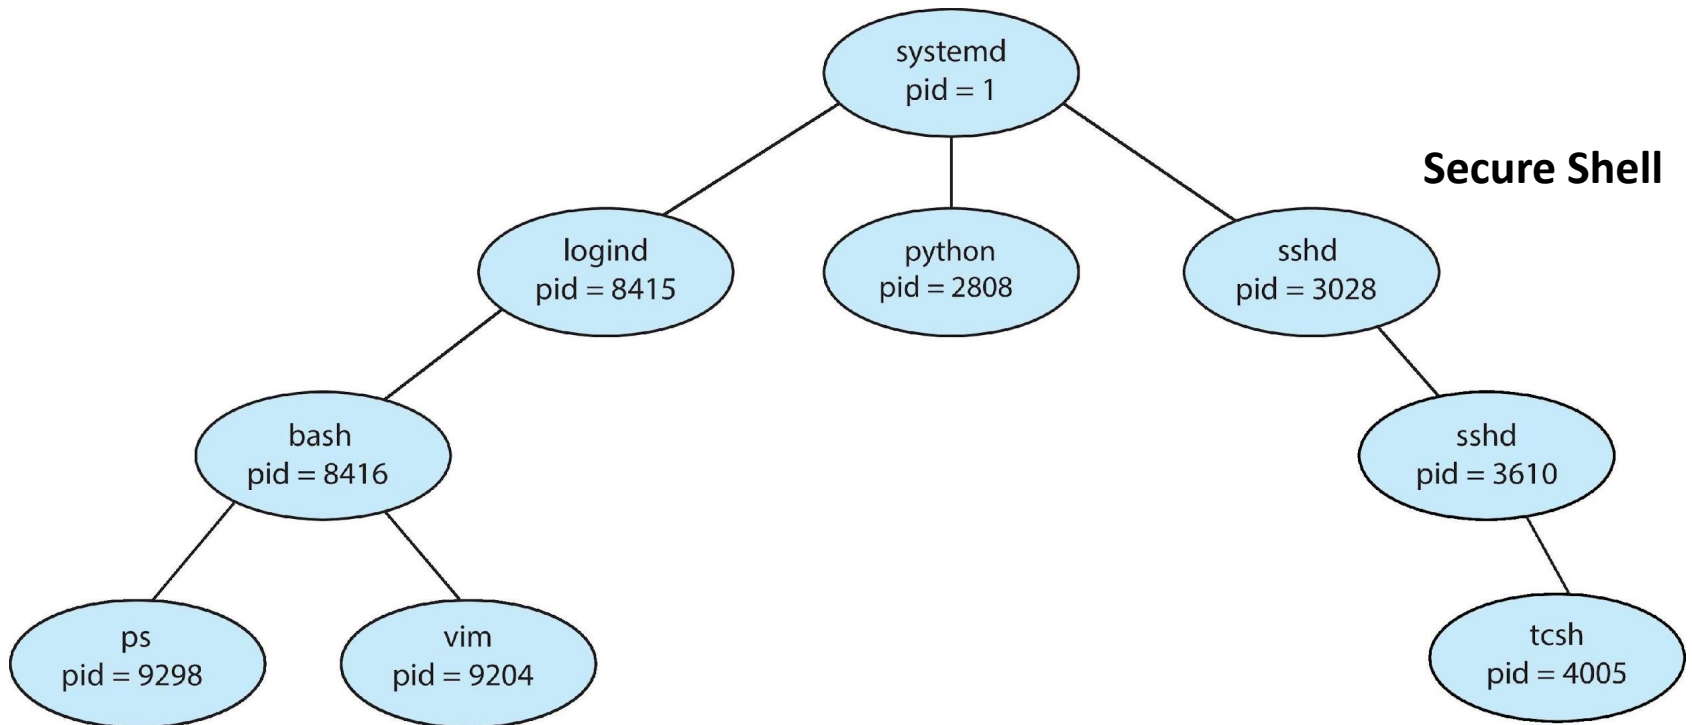
No	Func Name	Description	Source
1	<u>exit</u>	terminate the current process	<i>kernel/exit.c</i>
2	<u>fork</u>	create a child process	<i>arch/i386/kernel/process.c</i>
3	<u>read</u>	read from a file descriptor	<i>fs/read_write.c</i>
4	<u>write</u>	write to a file descriptor	<i>fs/read_write.c</i>
5	<u>open</u>	open a file or device	<i>fs/open.c</i>
6	<u>close</u>	close a file descriptor	<i>fs/open.c</i>
7	<u>waitpid</u>	wait for process termination	<i>kernel/exit.c</i>
8	<u>creat</u>	create a file or device ("man 2 open" for information)	<i>fs/open.c</i>
9	<u>link</u>	make a new name for a file	<i>fs/namei.c</i>
10	<u>unlink</u>	delete a name and possibly the file it refers to	<i>fs/namei.c</i>
11	<u>execve</u>	execute program	<i>arch/i386/kernel/process.c</i>
12	<u>chdir</u>	change working directory	<i>fs/open.c</i>
13	<u>time</u>	get time in seconds	<i>kernel/time.c</i>
14	<u>mknod</u>	create a special or ordinary file	<i>fs/namei.c</i>
15	<u>chmod</u>	change permissions of a file	<i>fs/open.c</i>
16	<u>lchown</u>	change ownership of a file	<i>fs/open.c</i>
18	<u>stat</u>	get file status	<i>fs/stat.c</i>
19	<u>lseek</u>	reposition read/write file offset	<i>fs/read_write.c</i>
20	<u>getpid</u>	get process identification	<i>kernel/sched.c</i>

Process Creation

Process Creation

- During execution, a process may create several new processes
 - Creating Process: Parent Process
 - New Processes: Children Processes
- Each of the children process may create other processes
 - Forming Tree of processes
- The **systemd** process (which always has a pid of 1) serves as the root parent process for all user processes
 - Systemd process in Unix is called init process

A Tree of Processes in Linux



ps -e command gives details of all active processes in the system

Class Summary

- System Calls

- System calls are interface for OS for manipulation of

- Processes

- Memory

- File and folders

- Application APIs

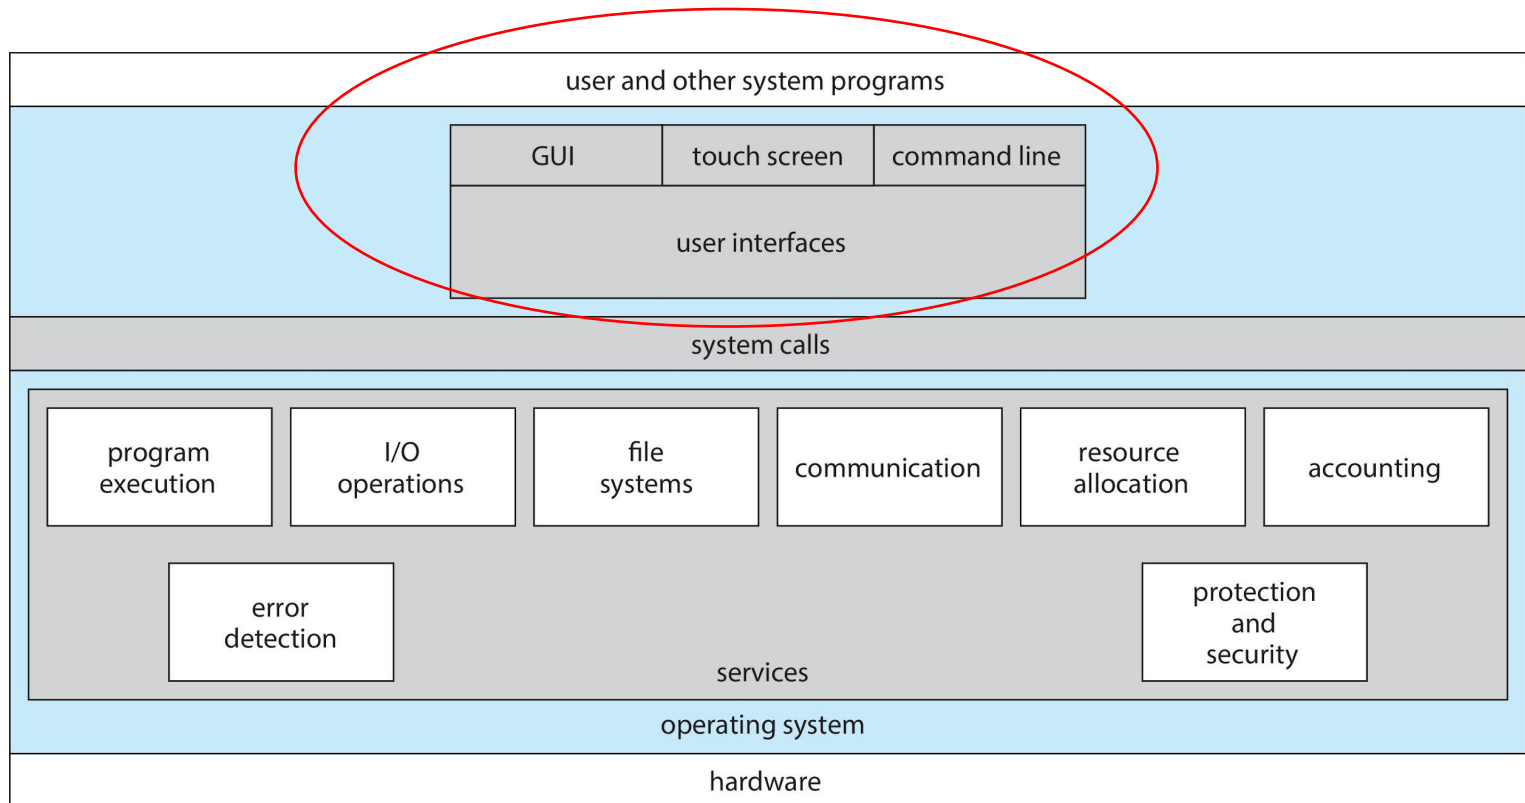
- A programmer usually uses higher level APIs instead of system calls

- User Interfaces: GUI, Command line, Touch Screen

- A tree of processes

User Interfaces (For self reading)

User Interfaces



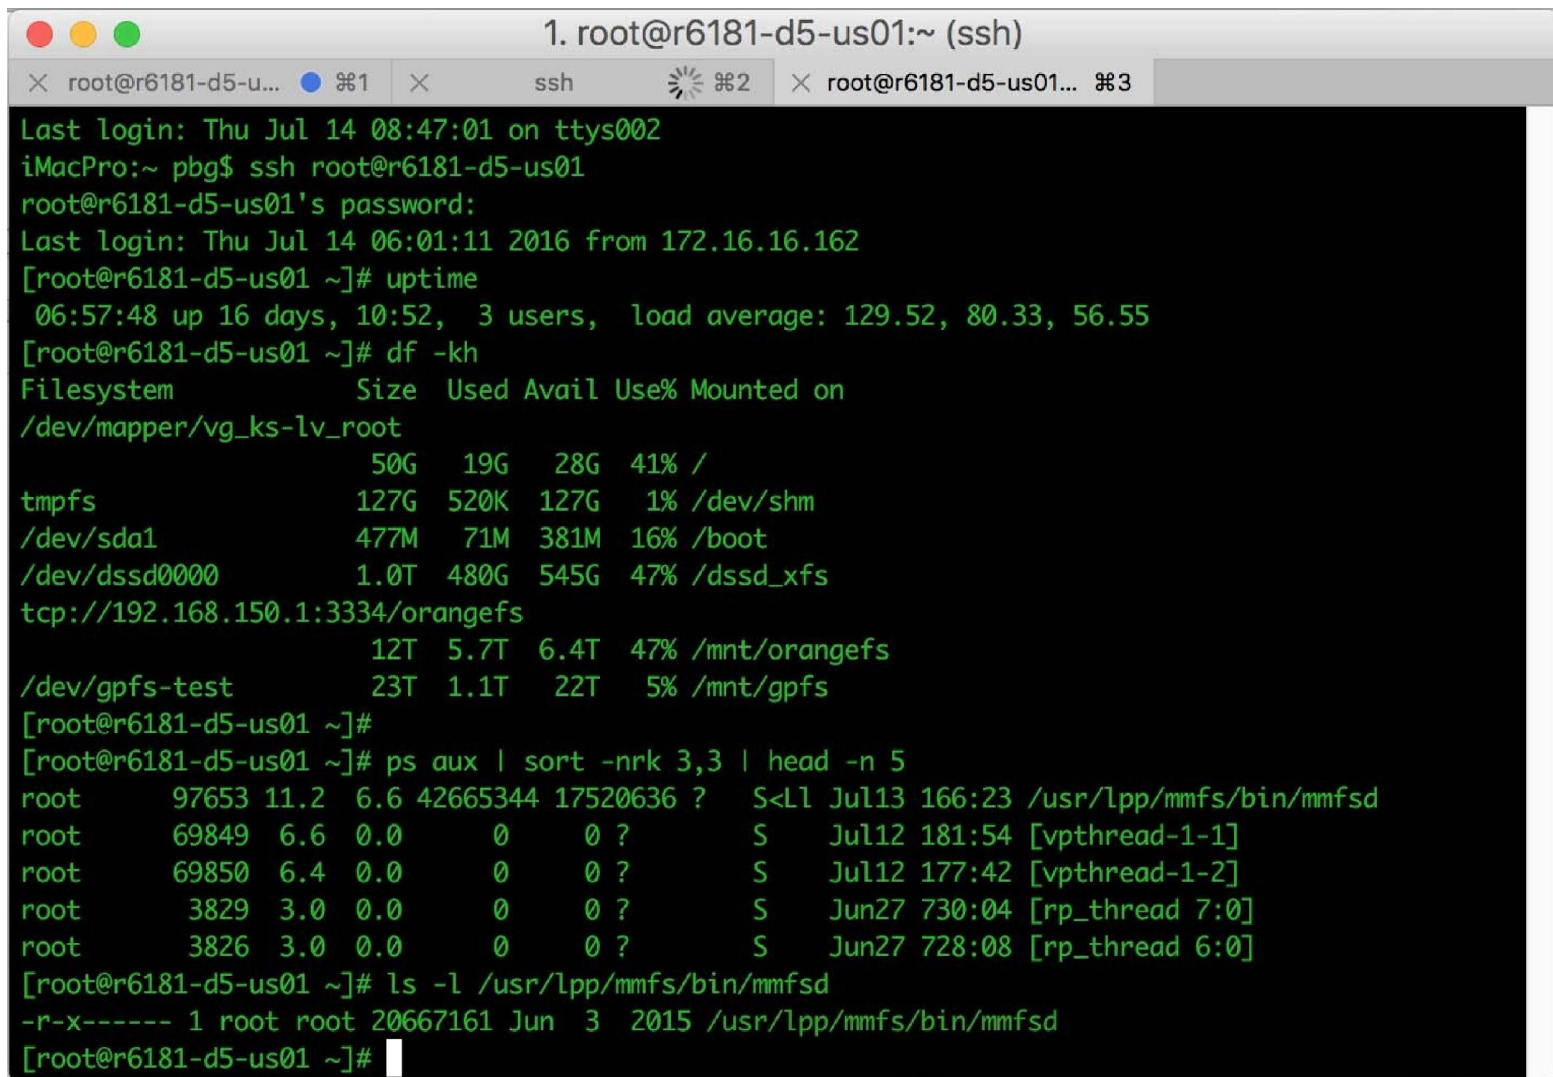
User Interfaces

- GUI
- Touch Screen
- Command Line Interface (CLI)
 - Uses text command
 - Allows users to directly enter commands executed by the OS
- Shells: Command Line Interpreters
 - C-Shell
 - Bourne-Again Shell (bash)
 - Korn Shell
- Example: Shell commands to
 - Create, delete, list, print, copy, execute - files

Interfaces

- GUIs are convenient for common users
 - Not all functionalities are available with GUIs
- CLIs : More powerful .. Used by advanced users and system administrators
 - More efficient, faster access, makes repetitive tasks easier
 - If a frequent task requires a set of command- line steps, those steps can be recorded into a file, and that file can be run just like a program
 - This program is interpreted by CLI
 - Shell Scripts
- Windows OS provides GUI
- MacOS provides both GUI and CLI


Bourne Shell Command Interpreter



```
1. root@r6181-d5-us01:~ (ssh)
root@r6181-d5-us01:~ (ssh)
Last login: Thu Jul 14 08:47:01 on ttys002
iMacPro:~ pbg$ ssh root@r6181-d5-us01
root@r6181-d5-us01's password:
Last login: Thu Jul 14 06:01:11 2016 from 172.16.16.162
[root@r6181-d5-us01 ~]# uptime
 06:57:48 up 16 days, 10:52,  3 users,  load average: 129.52, 80.33, 56.55
[root@r6181-d5-us01 ~]# df -kh
Filesystem                Size      Used Avail Use% Mounted on
/dev/mapper/vg_ks-lv_root   50G       19G   28G  41% /
tmpfs                      127G      520K   127G   1% /dev/shm
/dev/sda1                   477M       71M   381M  16% /boot
/dev/dssd0000               1.0T     480G   545G  47% /dssd_xfs
tcp://192.168.150.1:3334/orangefs 12T     5.7T   6.4T  47% /mnt/orangefs
/dev/gpfs-test              23T      1.1T    22T   5% /mnt/gpfs
[root@r6181-d5-us01 ~]#
[root@r6181-d5-us01 ~]# ps aux | sort -nrk 3,3 | head -n 5
root      97653 11.2  6.6 42665344 17520636 ?    S<Ll  Jul13 166:23 /usr/lpp/mmfs/bin/mmfsd
root      69849  6.6  0.0      0      0 ?        S    Jul12 181:54 [vpthread-1-1]
root      69850  6.4  0.0      0      0 ?        S    Jul12 177:42 [vpthread-1-2]
root       3829  3.0  0.0      0      0 ?        S    Jun27 730:04 [rp_thread 7:0]
root       3826  3.0  0.0      0      0 ?        S    Jun27 728:08 [rp_thread 6:0]
[root@r6181-d5-us01 ~]# ls -l /usr/lpp/mmfs/bin/mmfsd
-r-x----- 1 root root 20667161 Jun  3  2015 /usr/lpp/mmfs/bin/mmfsd
[root@r6181-d5-us01 ~]#
```

Touchscreen Interfaces

- Touchscreen devices require new interfaces
 - Mouse not possible or not desired
 - Actions and selection based on gestures
 - Virtual keyboard for text entry
- Voice commands



System Calls - types