

CS310 Operating Systems

Lecture 8: Process - System Calls – exit, wait, exec

Ravi Mittal
IIT Goa

Acknowledgements !


- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - Class presentation: University of California, Berkeley: David Culler, Anthony D. Joseph, John Kubiawicz, AJ Shankar, George Necula, Alex Aiken, Eric Brewer, Ras Bodik, Ion Stoica, Doug Tygar, and David Wagner
 - Operating Systems: Three Easy Pieces, by Remzi and Andrea Arpaci-Dusseau,
 - Chapter 5: Process APIs
 - Programs are taken from this chapter
 - CS 423 Operating System Design, Univ of Illinois, Prof Fagen-Ullmschneider
 - CS351 University of Washington

Read the following:

- Operating Systems: Principles and Practice (2nd Edition)
Anderson and Dahlin
 - Volume 1, Kernel and Processes
 - Chapter 4
- Operating Systems: Three Easy Pieces, by Remzi and Andrea
Arpaci-Dusseau,
 - Chapter 5: Process APIs

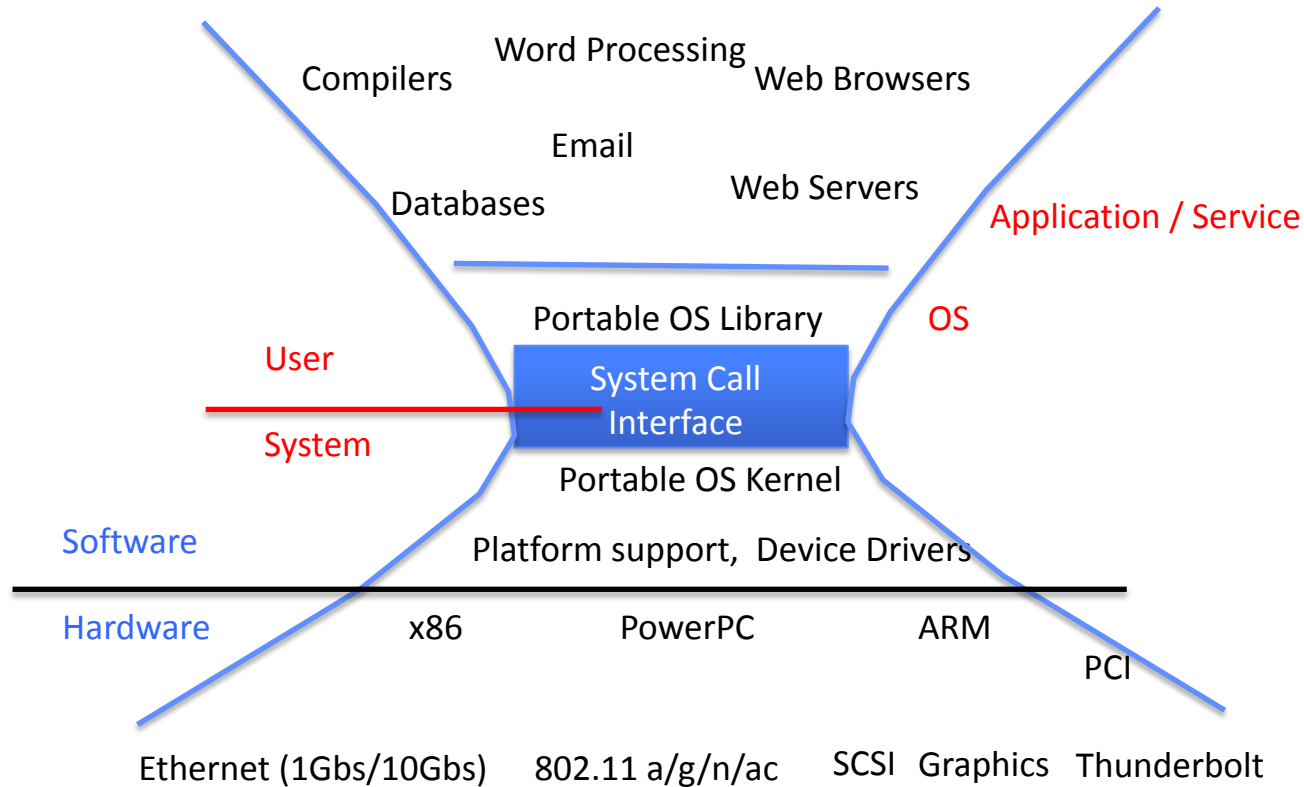
We will study..

- `exit()` system call
- `wait()` system call
- `exec*` system calls

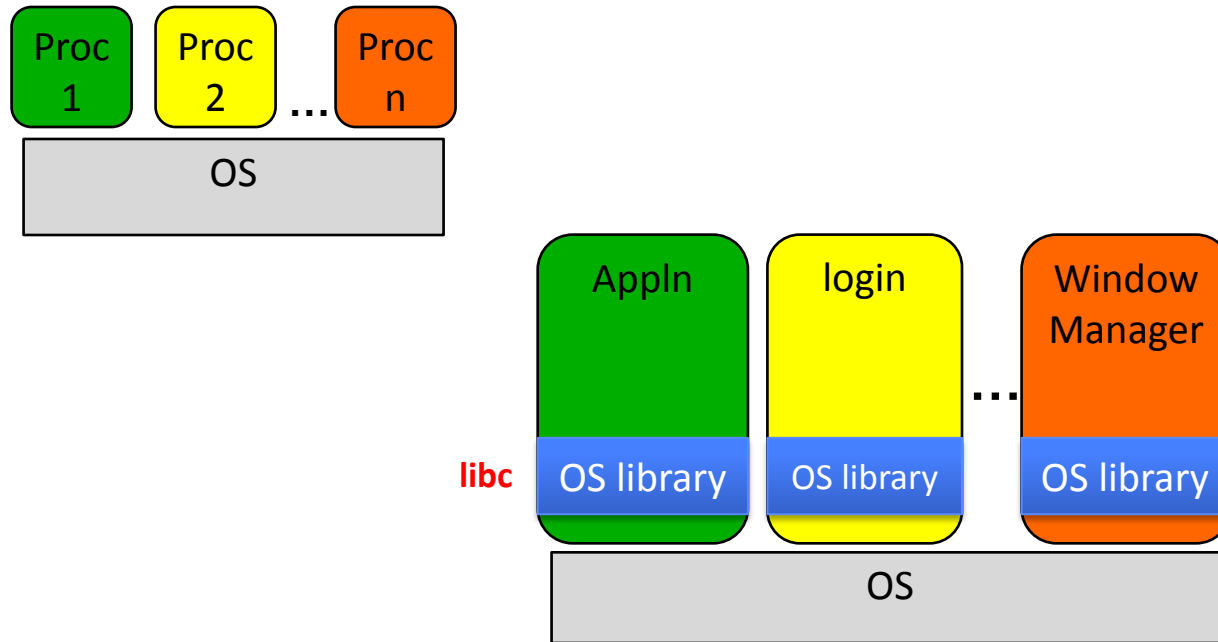


We have studied so far ...

System Calls (“Syscalls”)



OS Library Issues Syscalls



System Calls

Process Related System Calls (in Unix)


Last class

- `fork()` creates a new child process
 - All processes are created by forking from a parent
 - The *init* process is ancestor of all processes
- `exec()` makes a process execute a given executable
- `exit()` terminates a process
- `wait()` causes a parent to block until child terminates
- There are many variants of the above system calls with different arguments


Creating a process

- Just one process – Parent process
- Initially there is one process – `init` with `id = 1`

id = 1



id = 3042




System Calls – **exit()** and **wait()**

Process termination

- Multiple ways for a process to get destroyed
 - Process issues and *exit()* call – Voluntary
 - The parent process issues a *kill()* call - Involuntary
 - Process receives a *terminate signal* - Involuntary
 - When it did something illegal !
- On death
 - Reclaim all of process's memory regions
 - Make process un-runnable
 - Put the process in the *zombie state* (will discuss it later)
 - However, do not remove its *process descriptor* from the list of processes

Waiting for children to die with **wait()**

- The parent can wait for the child to die by executing the **wait** system call
- It is quite useful for a parent to wait for a child process to finish what it has been doing
 - on success, **returns** the process ID of the terminated child; on error, -1 is **returned**.



wait and waitpid syscalls

- A terminated process's information is collected via a call of a wait operation by its parent
- Operations:
 - wait
 - blocks the calling process until a child process terminates, returning the child's pid
 - If caller has no children, wait immediately returns -1 (error)
 - the termination status (return value) of the child may be obtained via the argument
 - waitpid
 - permits a caller to wait for a particular child, identified by its pid

wait() system call

p2.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int
7  main(int argc, char *argv[])
8  {
9      printf("hello world (pid:%d)\n", (int) getpid());
10     int rc = fork();
11     if (rc < 0) {
12         // fork failed; exit
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) {
16         // child (new process)
17         printf("hello, I am child (pid:%d)\n", (int) getpid());
18         sleep(1);
19     } else {
20         // parent goes down this path (original process)
21         int wc = wait(NULL);
22         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
23                rc, wc, (int) getpid());
24     }
25     return 0;
26 }
```

wait() system call

Parent process waits for the child to finish


```
(base) Ravis-MacBook-Pro-2:cp ravimittal$ ./p2  
hello world (pid:13858)  
hello, I am child (pid:13859)  
hello, I am parent of 13859 (wc:13859) (pid:13858)
```


System Calls – **exec family**

Executing a new program

- After **fork**, parent and child are running same code
 - Not too useful!
- A common use of fork is to launch a new executable program
- A process can run **exec()** to load another executable to its memory image
 - So, a child can run a different program from parent
 - It can also be shell script
- The **exec** system call replaces the current process image with a new image
 - *If **exec** succeeds, it never returns*
- **exec** requires you to specify the file you program to run

exec() system call



- The `exec` family of system calls replaces the program executed by a process
- When a process calls `exec`, all code (text) and data in the process is lost and replaced with the executable of the new program
- All open file descriptors remains open after calling `exec`
 - unless explicitly set to close-on-exec

The `exec()` System Call

- There's no a syscall under the name `exec()`. By `exec()` we usually refer to a family of calls:
 - `int execl(char *path, char *arg, ...);`
 - `int execv(char *path, char *argv[]);`
 - `int execlp(char *path, char *arg, ..., char *envp[]);`
 - `int execve(char *path, char *argv[], char *envp[]);`
 - `int execlp(char *file, char *arg, ...);`
 - `int execvp(char *file, char *argv[]);`
- Here's what *l*, *v*, *e*, and *p* mean:
 - *l* means an argument list,
 - *v* means an argument vector,
 - *e* means an environment vector, and
 - *p* means a search path.

execvp() system call

```
#include <unistd.h>
```

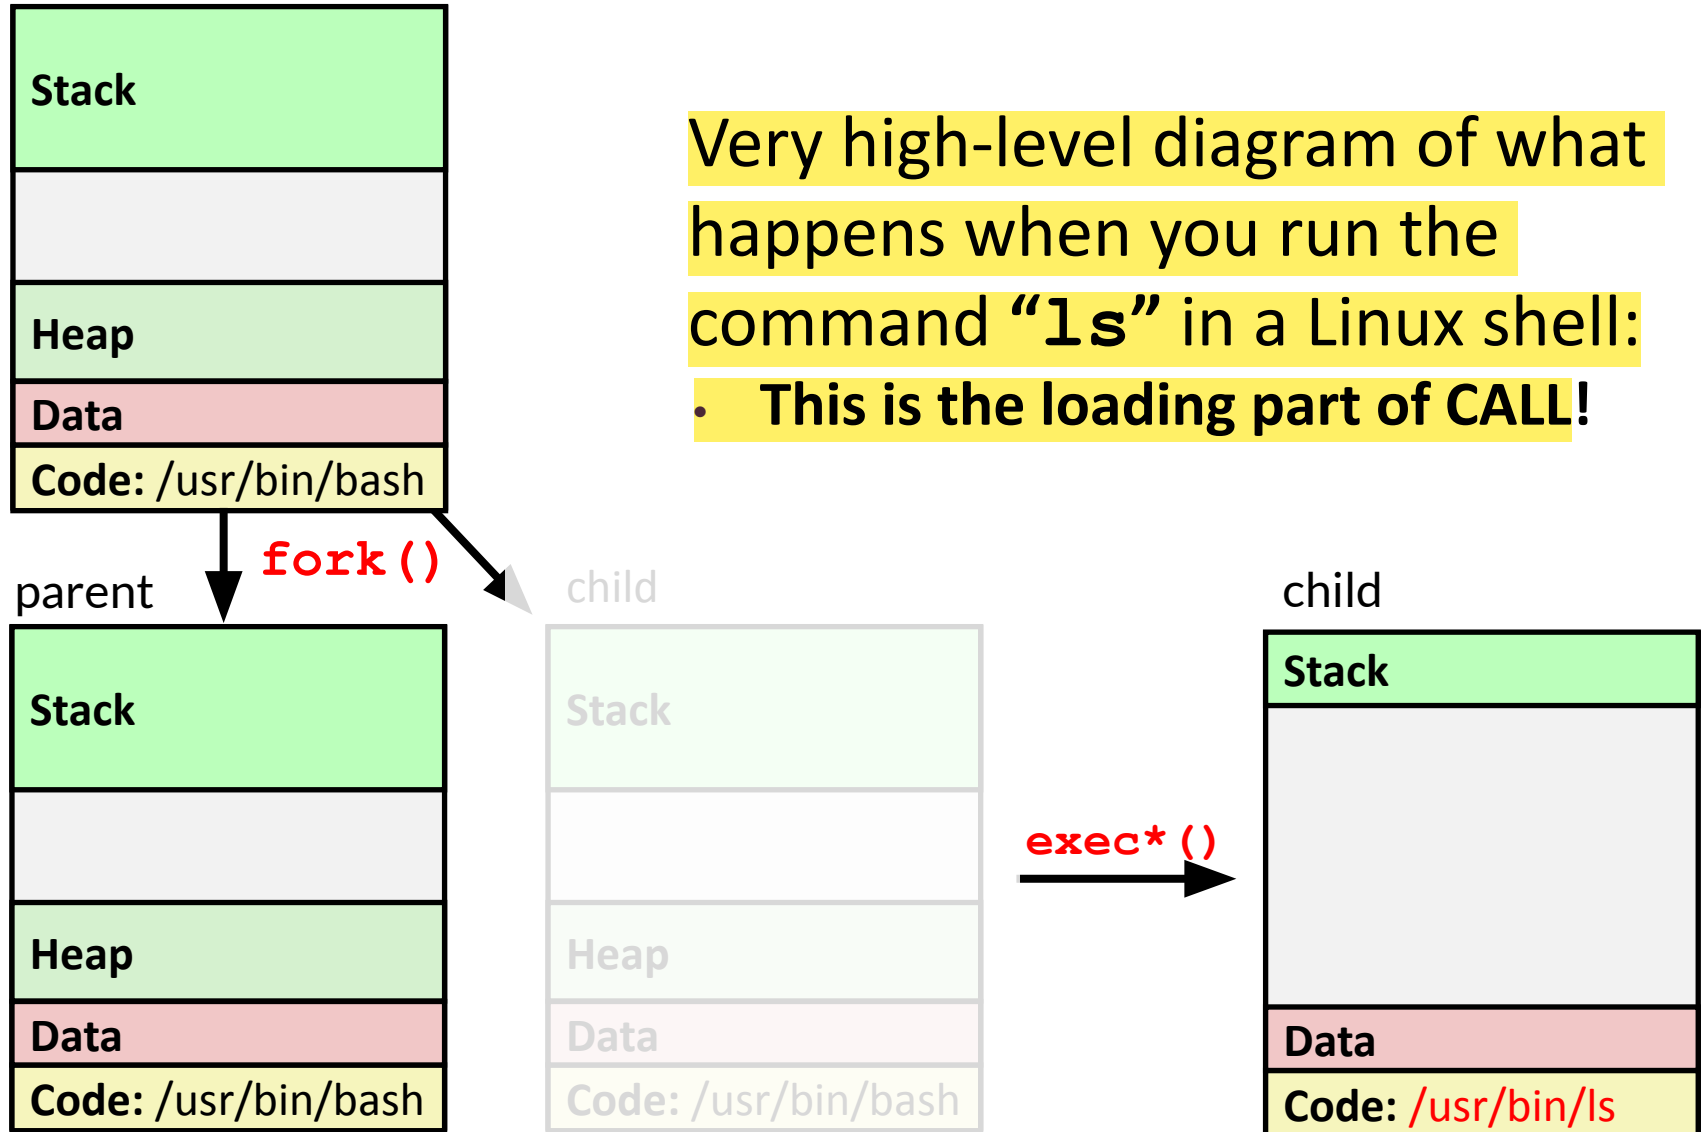
```
int execvp(const char *file, char *const argv[]);
```

- *file*: Used to construct a pathname that identifies the new process image file
 - If the *file* argument contains a slash character, the *file* argument is used as the pathname for the file
- *argv*: An array of character pointers to NULL-terminated strings
 - These strings constitute the argument list available to the new process image
 - *argv[0]* must point to a filename that's associated with the process being started
 - the last member of this array is a NULL pointer

Exec-ing a new program

Very high-level diagram of what happens when you run the command “ls” in a Linux shell:

- This is the loading part of CALL!



Example: Using `execvp()`

```
1
2 #include <unistd.h> // execvp()
3 #include <stdio.h>  // perror()
4 #include <stdlib.h> // EXIT_SUCCESS, EXIT_FAILURE
5
6 int main(void) {
7     char *const cmd[] = {"ls", "-l", NULL};
8     execvp(cmd[0], cmd);
9     perror("Return from execvp() not expected");
10    exit(EXIT_FAILURE);
11 }
12 |
```

```
(base) Ravis-MacBook-Pro-2:cp ravimittal$ ./p5
total 536
-rw-r--r--@ 1 ravimittal  staff    42  6 Sep 10:50 clients.txt
-rw-r--r--@ 1 ravimittal  staff 142187  5 Sep 21:33 fig11_02.c
-rwxr-xr-x  1 ravimittal  staff 12784  3 Sep 09:12 p1
-rw-r--r--@ 1 ravimittal  staff   486  3 Sep 08:24 p1.c
-rwxr-xr-x  1 ravimittal  staff 12872  2 Sep 21:02 p2
-rw-r--r--  1 ravimittal  staff   655 30 Aug 12:11 p2.c
-rwxr-xr-x  1 ravimittal  staff 13044 30 Aug 13:21 p3
-rw-r--r--@ 1 ravimittal  staff   968 31 Aug 09:26 p3.c
-rw-r--r--@ 1 ravimittal  staff   262 10 Aug 16:00 p4.c
-rwxr-xr-x  1 ravimittal  staff 12644  6 Sep 22:07 p5
-rw-r--r--@ 1 ravimittal  staff   268 30 Aug 13:14 p5.c
-rwxr-xr-x  1 ravimittal  staff 12644 30 Aug 13:14 pexe
-rwxr-xr-x  1 ravimittal  staff 13036  5 Sep 22:17 pf
-rw-r--r--@ 1 ravimittal  staff   281  5 Sep 21:50 pfile1.c
-rw-r--r--@ 1 ravimittal  staff   933  5 Sep 22:17 pfile2.c
```

```
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>
```

```
int
main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```


Program Output

```
hello world (pid:25155)
hello, I am child (pid:25156)
    32      123      966 p3.c
hello, I am parent of 25156 (wc:25156) (pid:25155)
```


exec() – more info

- Upon success, exec() never returns to the caller
- A successful exec replaces the current process image, so it cannot return anything to the program that made the call
- If it does return, it means the call failed. Typical reasons are: non-existent file (bad path) or bad permissions
- As a new process is not created, the process identifier (PID) does not change
 - However, machine code, data, heap, and stack of the process are replaced by those of the new program



fork() and exec() combined

- Often after doing fork() we want to load a new program into the child



Zombies

- A terminated process still consumes system resources
 - Various tables maintained by OS
 - Called a “**zombie**” (a living corpse, half alive and half dead)
- *Reaping* is performed by parent on terminated child
 - Parent is given exit status information and kernel then deletes zombie child process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (pid of 1)
 - **Note:** on recent Linux systems, `init` has been renamed `systemd`
 - In long-running processes (*e.g.* shells, servers) we need *explicit* reaping

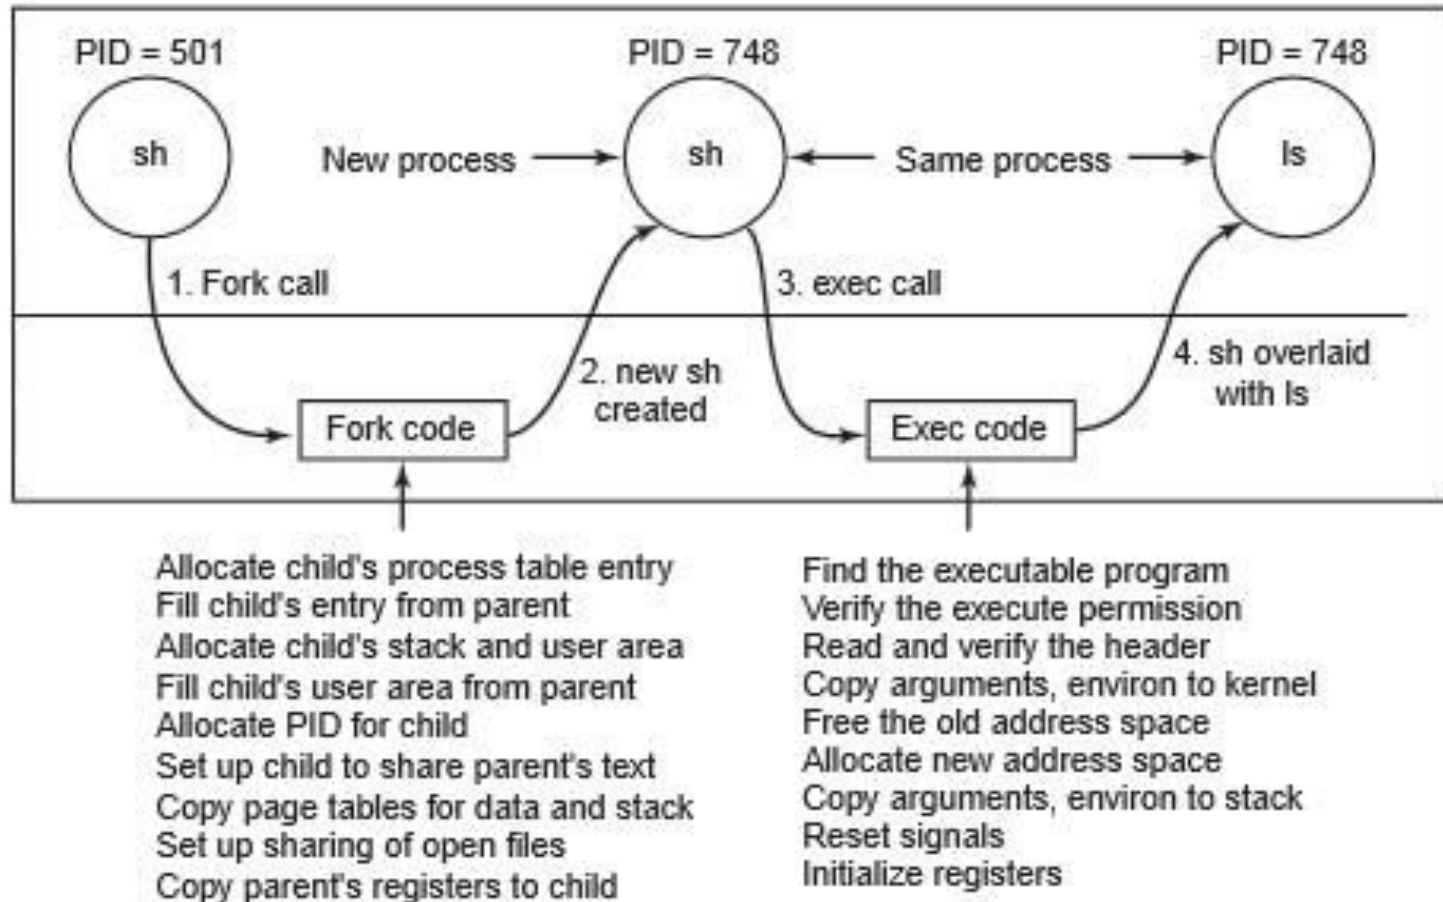
Zombie State

- Why keep process descriptor around?
 - Parent may be waiting for child to terminate
 - via the *wait()* system call
 - Parent needs to get the exit code of the child
 - this information is stored in the descriptor (PCB)
 - If descriptor was destroyed immediately, this information could not be gotten
 - After getting this information, the process descriptor (PCB) can be removed
 - no more remnants of the process

Class Summary

- `fork` makes two copies of the same process (parent & child)
 - Returns different values to the two processes
- `exec*` replaces current process from file (new program)
 - Two-process program:
 - First `fork()`
 - `if (pid == 0) { /* child code */ } else { /* parent code */ }`
 - Two different programs:
 - First `fork()`
 - `if (pid == 0) { execv(...) } else { /* parent code */ }`
- `wait` or `waitpid` used to synchronize parent/child execution and to reap child process

fork() AND exec() combined (1/2)



exec() – More info

- The exec() call replaces a current process' image with a new one (i.e. loads a new program within current process)
- The new image is either regular executable **binary file** or a **shell script**
- There is not a syscall under the name exec(). By exec() we usually refer to a family of calls:
 - `int execl(char *path, char *arg, ...);`
 - `int execv(char *path, char *argv[]);`
 - `int execlp(char *path, char *arg, ..., char *envp[]);`
 - `int execve(char *path, char *argv[], char *envp[]);`
 - `int execlp(char *file, char *arg, ...);`
 - `int execvp(char *file, char *argv[]);`
- Here's what l, v, e, and p mean:
 - **l** means an argument list,
 - **v** means an argument vector,
 - **e** means an environment vector, and
 - **p** means a search path.