

# **CS310 Operating Systems**

## **Lecture 14: Threads – Examples**

Ravi Mittal  
IIT Goa

# Acknowledgements !

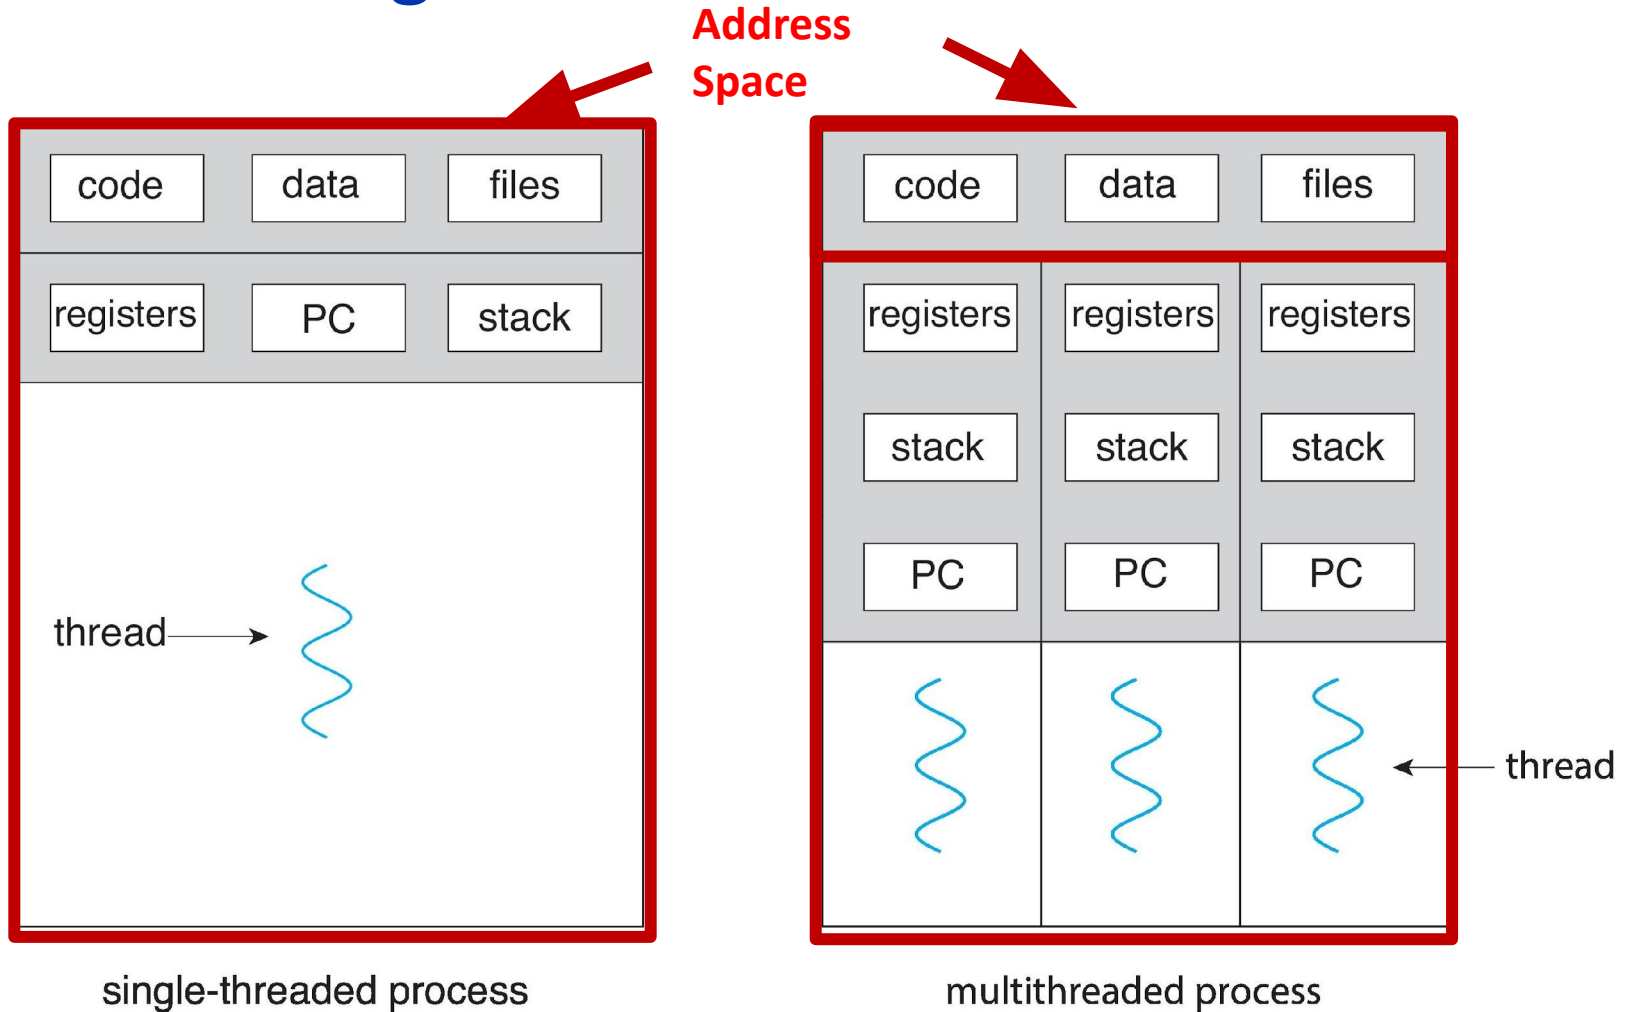
- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
  - CS162, Operating System and Systems Programming, Profs. Natacha Crooks and Anthony D. Joseph, University of California, Berkeley
  - CS240 Computer Systems, Univ of Illinois, Prof. Wade Fagen-Ulmschneider
  - Operating Systems: Three Easy Pieces, by Remzi and Andrea Arpaci-Dusseau, available for free online
  - Book: Modern Operating Systems, Andrew Tenenbaum, and Herbert Bos, 4<sup>th</sup> Edition, Pearson

# Reading

- Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Part 2, Chapter 4
- CS162, Operating Systems and Systems Programming, University of California, Berkeley
- CS4410, Operating Systems, Course, Cornell University, Spring 2019, Lecture on Threads
- Operating Systems: Three Easy Pieces, by Remzi and Andrea Arpaci-Dusseau, available for free online

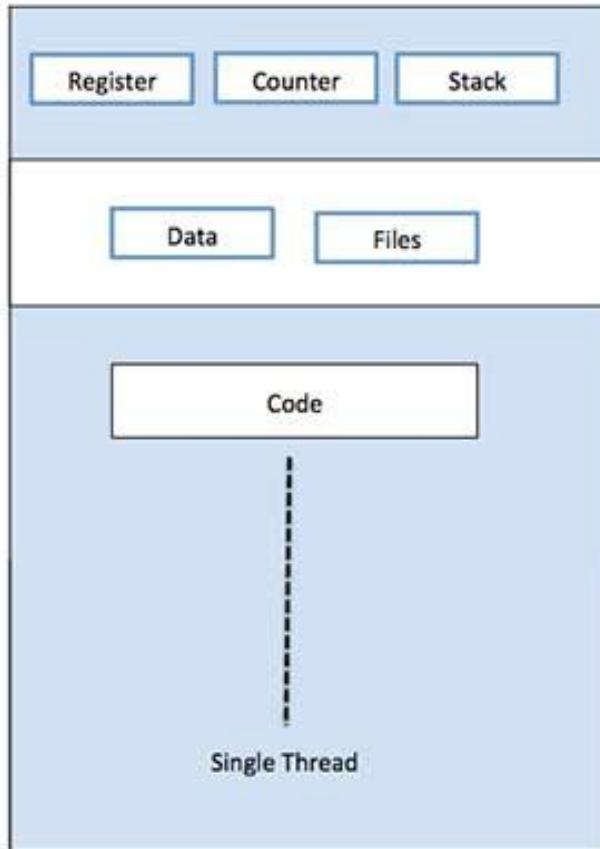
**We have learnt so far ..**

# Recall: Single and Multithreaded Processes

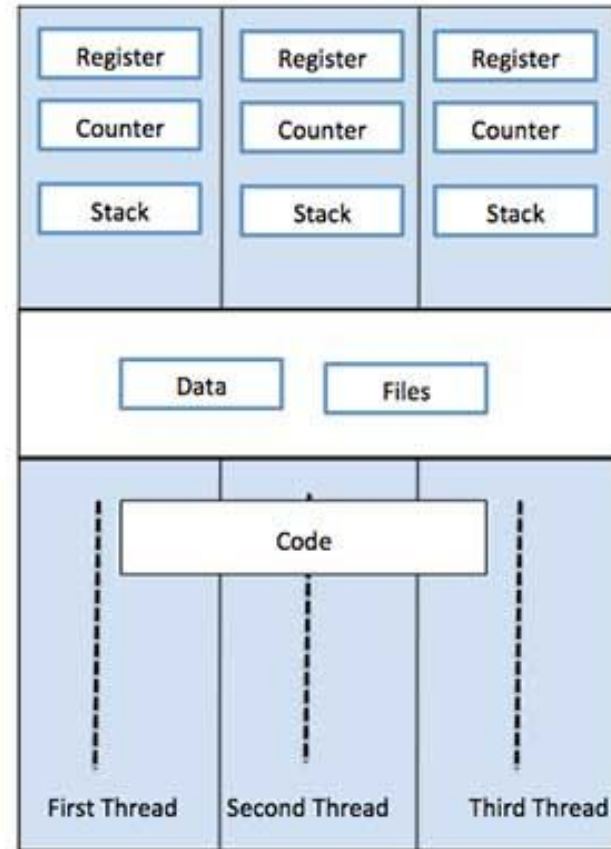


- Address spaces encapsulate protection: Passive Part
- Threads share code, data, files, heap


# Thread Abstraction



Single Process P with single thread



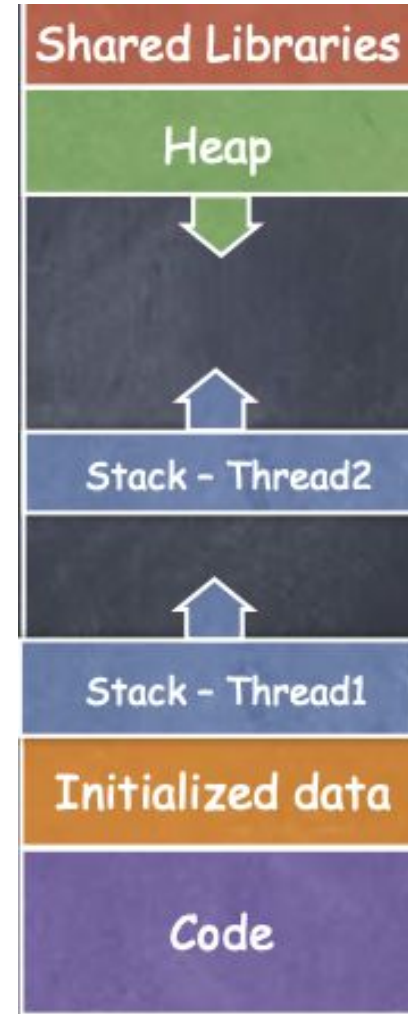
Single Process P with three threads



# Threads

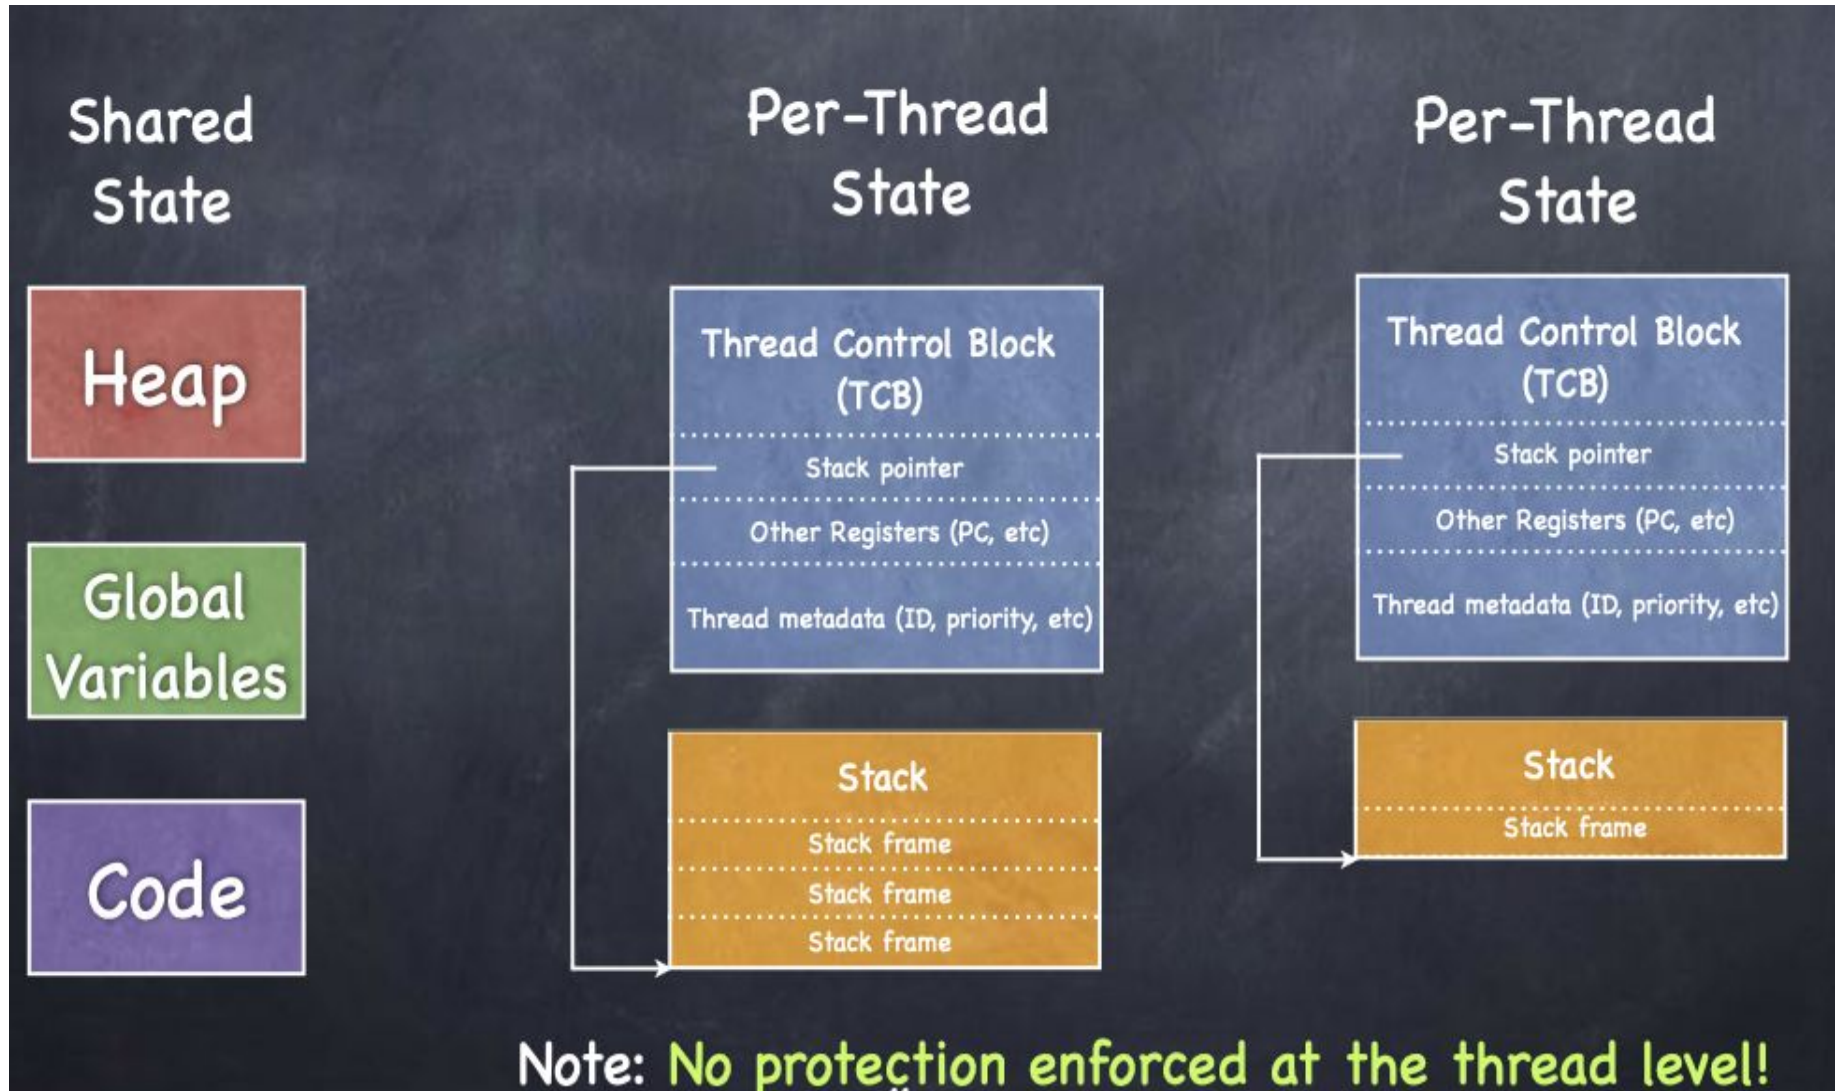
- All threads within a process share
  - Heap
  - Global/static data
  - Libraries
- Each thread has a separate
  - Program Counter
  - Stack
  - registers

Process Address Space






# Implementing Thread Abstraction




# Implementing Thread in User Space (1)



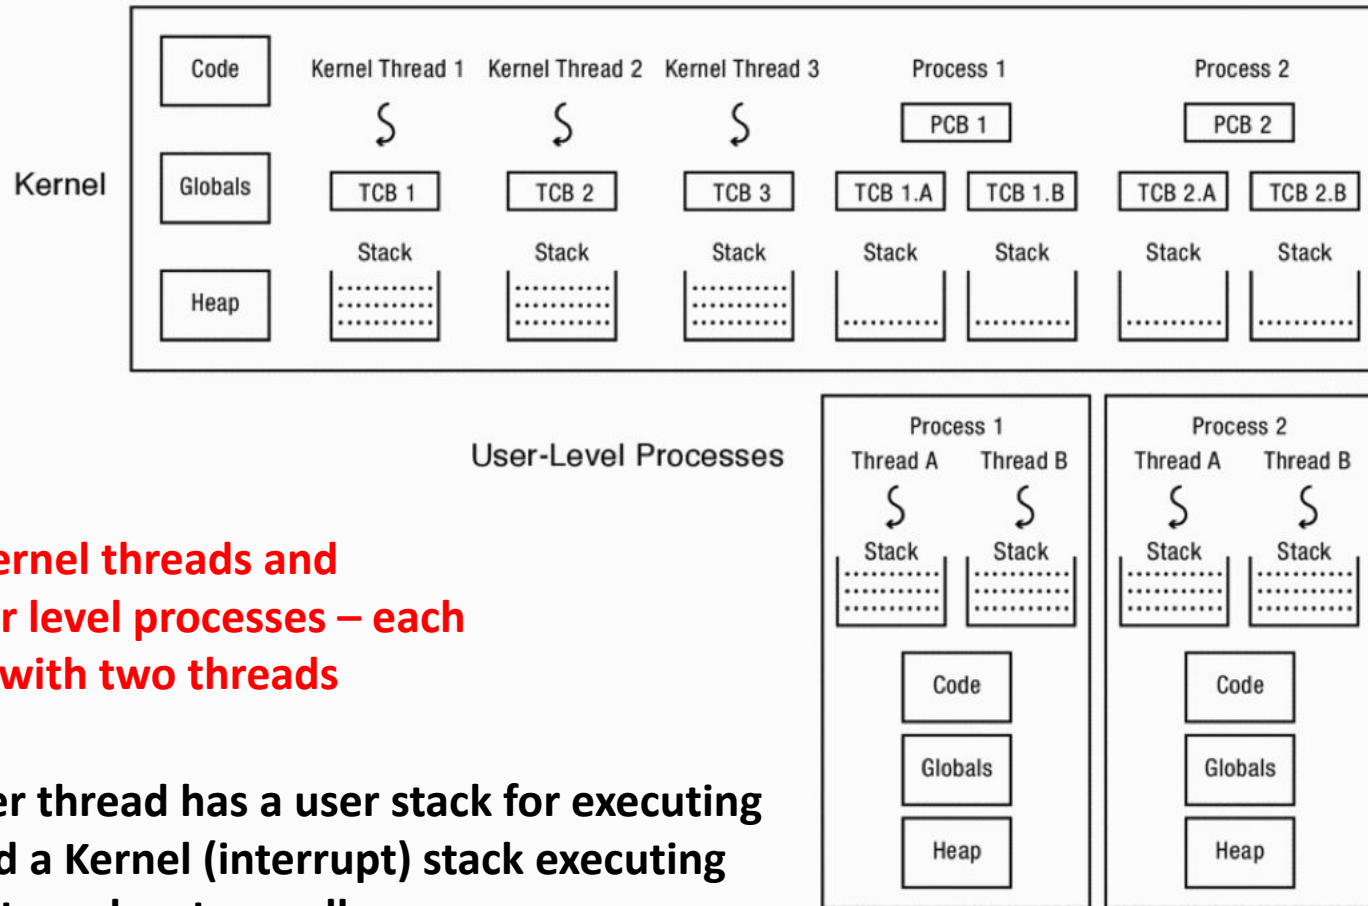
- Threads run on top of run-time system
  - Collection of procedures to manage threads
- Each process needs own private thread table (Thread control Block)

# Thread Management in Kernel (1)

- Kernel knows how to manage threads in a user process
  - No run-time system required
  - No thread table in each process
- When a thread wants to create a new thread or destroy an existing thread, makes a system call
  - Kernel creates or destroys thread as per request.
  - Updates kernel thread table



# Thread Management in Kernel



**Three Kernel threads and  
Two user level processes – each  
process with two threads**


**Each user thread has a user stack for executing  
code and a Kernel (interrupt) stack executing  
interrupts and system calls**

**Thread management is done in Kernel**


# Logical View of Threads

- Threads associated with a process form a pool of peers
  - Unlike processes, which form a tree hierarchy

Threads associated with process foo



Process hierarchy



# Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
    - All code and data structures for the library exist in user space with no kernel support
  - Kernel-level library supported by the OS
    - Code and data structures exist in Kernel Space
- Invoking a function in the API for the library typically results in a system call to the kernel

# Today we will study ..

- Posix Thread library
- Thread Example 1

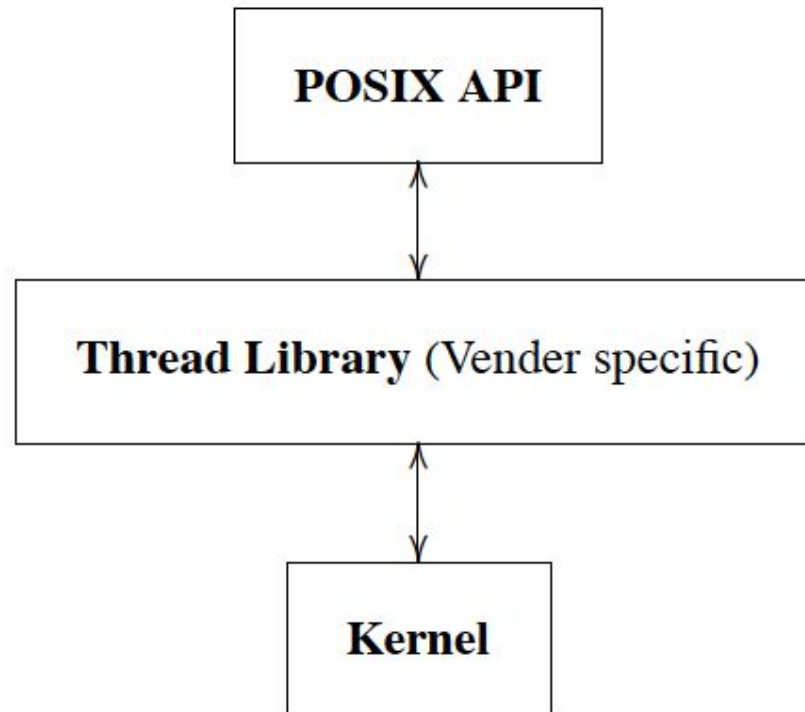
# Thread Libraries

- Three main Thread Libraries in use today
  - POSIX Pthread
    - User or Kernel level
  - Windows
    - Kernel-level
  - Java
    - Using Windows APIs
- Linux, Unix, MacOS
  - Pthreads



**Thread Library: pthread**

# Thread Implementation: Layers of Abstraction



# What are **pthread**s?

- POSIX standard IEEE 1003.1c defines a thread interface
  - `pthread`
- Unix/Linux provides `pthread` library
  - APIs to create and manage threads
- Implementation is up to development of the library
- Simply a collection of C functions
- Standard interface for ~60 functions that manipulate threads from C programs
- Primary way of doing threading in Linux is `pthread`
- Since 2003 (Kernel 2.6), Linux implements POSIX threads as `kernel-scheduled threads`

# pthread APIs

- Creating and reaping threads
  - `pthread_create`, `pthread_join`
- Determining your thread ID
  - `pthread_self`
- Terminating threads
  - `pthread_cancel`, `pthread_exit`
- Synchronizing access to shared variables
  - `pthread_mutex_init`, `pthread_mutex_[un]lock`
  - `pthread_cond_init`, `pthread_cond_[timed]wait`

## Aside: Pointers to Functions in C (self reading)

- A function can take many types of arguments including the address of another function
- Instructions of a program are stored in MM locations
- Instruction too have addresses
  - Like data
- So, we can point to an individual instruction
  - Pointing to an individual instruction is not useful
- Point to the first instruction of a function □ useful
- **Function Pointer :**
  - a variable or parameter that points at a function
- Note that a function name is a pointer to it
  - Helps to initialize such variables or parameters

## Aside: Pointers to Functions in C (self reading)

- Function name is actually starting address of the code in the memory (think Computer Architecture – Machine instructions)
- Start of the function code or the address of a function is a function pointer
- Function pointer is different from other pointers since you do not allocate or deallocate memory with them
- Function pointers can be passed as arguments to other functions or return from functions
- Unfortunately function pointers have complicated syntax and therefore are not widely used

## Aside: Pointers to Functions in C (self reading)

- Provides late binding
  - deciding the proper function during runtime instead of compile time
- Ex: Determine sorting function based on type of data at run time
- Technically speaking, the name of any function is a pointer to that function
- However, when we refer to a function pointer, we mean a variable or parameter that points to a function

## Aside: Pointers to Functions in C (self reading)

```
int (*fn)(int,int) ;
```

- Here we define a function pointer `fn`, that can be initialized to any function that takes two integer arguments and return an integer
- Consider a function:

```
int sum(int x, int y) {  
    return (x+y);  
}
```

Now to initialize `fn` to the address of the `sum`, we can do the following:

```
fn = &sum    /* make fn points to the address of sum */
```

```
int x = (*fn)(12,10); /* call to the function through a pointer */
```



# Pthread\_create

```
int pthread_create(  
    pthread_t *thread,      /* thread struct */  
    const pthread_attr_t *attr, /* usually NULL */  
    void *(*start_routine) (void *), /* start func */  
    void *arg                /* thread start arg */  
);
```

- This function takes 4 parameters
- The first parameter is a pointer to structure of type `pthread_t` which describes a thread
  - `Pthread_create` function fills in this structure with information about the newly created thread
  - We can use type `pthread_t` to refer to the thread in other library calls
  - The thread parameter points to the thread ID of the newly created thread
- The second argument lets you specify particular attribute of the thread
  - Usually Null. The new thread has default attributes.

# Pthread\_create

```
int pthread_create(  
    pthread_t  *thread,      /* thread struct */  
    const pthread_attr_t *attr, /* usually NULL */  
    void *(*start_routine) (void *), /* start func */  
    void *arg                /* thread start arg */  
);
```

- The third argument specifies the entry point to the new thread
  - Unlike the entire program, which starts at main, there is no pre-defined entry point for other threads
  - The thread that creates the new thread must specify the function in which the new thread must start
  - This function is specified by passing a function pointer which must have type `void* (*)(void *)`
    - A pointer to a function that takes a `void*` and returns a `void*`

# Pthread\_create

```
int pthread_create(  
    pthread_t  *thread,      /* thread struct */  
    const pthread_attr_t *attr, /* usually NULL */  
    void *(*start_routine) (void *), /* start func */  
    void *arg                /* thread start arg */  
);
```

- **arg** parameter specifies what to pass into this function as its arguments when the function is called on the newly spawned thread
- It returns 0 on success; non-zero error code on error
- The thread created by **pthread\_create** is runnable without requiring a separate start code. It begins execution on the first line of the function routine provided in the third argument.

# Pthread\_create

- When `pthread_create` is called a few things happen ( that are different from we have seen before)
- A new stack is created that is independent of caller's stack
  - Frames can be created and destroyed on this stack – which is independent of frame creation and destruction on other stacks (for other threads)
  - The stack with a frame is created for the function requested (as the third argument) passing in the arguments
- Now second execution arrow (corresponding to the current value of PC) is created at the start of the function that we specify as the entry point of the new thread
  - You can imagine execution arrow as a point in execution of the thread
- Concurrent execution of both threads begins

# pthread – Completion

- A Thread exits
  - By returning from the function it started in, or
  - By calling `pthread_exit`
  - It's return value ( either what the function returned or argument passed to `pthread_exit` must be kept available for another thread to retrieve
- `pthread_join`
  - Thread calling `pthread_join` waits for the specified thread to exit and obtaining the return value
  - When one thread calls `pthread_join` - it blocks until the thread it is joining terminates
    - Resources associated with the thread that is ended

# Pthread - Completion


```
int pthread_join(pthread_t thread, void **value_ptr);
```

- The first argument is of type `pthread_t` and it specifies the thread to wait for
- The second argument is a pointer to the return value you expect to get back
  - It ignores the return value by passing in NULL for the second argument
- When one thread calls `pthread_join` it stops advancing – until the thread it is joining terminates
- Joining a thread also allows for the **pthread library** to release resources associated with the thread
- A **zombie thread** is a joinable thread which has terminated, but which hasn't been joined.

# pthread\_detach; Main terminating

- Pthread-detach
  - A thread tells the pthread library that it will never join another thread with pthread\_detach
  - This allows library to release resources as soon as the thread exits
- The Main thread (in main function) behaves differently
  - If main returns, the entire process exits – thus terminating all threads inside it
  - Note that main was not called by pthread library. It was called by C library

# Fork-Join Parallelism



- Main thread *creates* (forks) collection of sub-threads passing them **args** to work on...
- ... and then *joins* with them, collecting results
- Data may be safely shared between parent and child
  - It is written by parent before child starts
  - It is written by the child and read by the parent after the join



# Thread - Example

## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## th1.c

## Thread 0

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

main

## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23     pthread_join(thread, NULL);
24     printf("join completed \n");
25     return EXIT_SUCCESS;
26 }
27
28
29
```

## Thread 0

main


Thread

?

## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```


## Thread 0




## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15
16 int main(void) {
17     pthread_t thread;
18     int x = 3;
19     pthread_create(&thread, NULL, sqr, &x);
20     sleep(1);
21     for (int i = 0; i < 100; i++){
22         printf("main: %d\n", i);
23     }
24     pthread_join(thread, NULL);
25     printf("join completed \n");
26     return EXIT_SUCCESS;
27 }
28
29
```

## Thread 1




## Thread 0




## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```


## Thread 1



## Thread 0




## output




## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 1



## Thread 0




## output




## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("squaare: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 1



## Thread 0




## output


## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 1



## Thread 0




## output


## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 1



## Thread 0




## output


# th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10        int fsq = (i * i);
11        printf("square: %d\n", fsq);
12    }
13    return NULL;
14 }
15
16 int main(void) {
17     pthread_t thread;
18     int x = 3;
19     pthread_create(&thread, NULL, sqr, &x);
20     sleep(1);
21     for (int i = 0; i < 100; i++){
22         printf("main: %d\n", i);
23     }
24     pthread_join(thread, NULL);
25     printf("join completed \n");
26     return EXIT_SUCCESS;
27 }
28
29
```

## Thread 1



## Thread 0




## output


## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10        int fsq = (i * i);
11        printf("square: %d\n", fsq);
12    }
13    return NULL;
14 }
15
16 int main(void) {
17     pthread_t thread;
18     int x = 3;
19     pthread_create(&thread, NULL, sqr, &x);
20     sleep(1);
21     for (int i = 0; i < 100; i++){
22         printf("main: %d\n", i);
23     }
24     pthread_join(thread, NULL);
25     printf("join completed \n");
26     return EXIT_SUCCESS;
27 }
28
29
```

## Thread 1



## Thread 0




## output

square 0  
square 1  
square 4

## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27 }
28
29
```

## Thread 0



thread 1 exited


## output

square 0  
square 1  
square 4

## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23     pthread_join(thread, NULL);
24     printf("join completed \n");
25     return EXIT_SUCCESS;
26 }
27
28
29
```

## Thread 0



## output

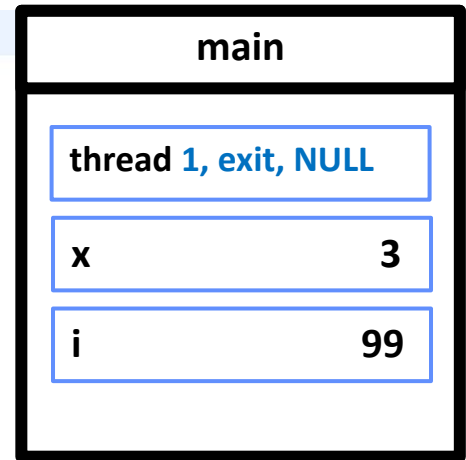
main 0



## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 0



## output


main 0  
main 1  
  
main 99



## th1.c

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 void * sqr(void *arg){
7     int * ptr = arg;
8     int n = *ptr;
9     for (int i = 0; i < n ; i++){
10         int fsq = (i * i);
11         printf("square: %d\n", fsq);
12     }
13     return NULL;
14 }
15 int main(void) {
16     pthread_t thread;
17     int x = 3;
18     pthread_create(&thread, NULL, sqr, &x);
19     sleep(1);
20     for (int i = 0; i < 100; i++){
21         printf("main: %d\n", i);
22     }
23 }
24 pthread_join(thread, NULL);
25 printf("join completed \n");
26 return EXIT_SUCCESS;
27
28 }
29
```

## Thread 0



Thread 0 unblocks

## output

main 0  
main 1

main 99  
join competed

# Lecture Summary

- We have used pthread library's APIs
  - pthread\_create
  - pthread\_join