

# CS310 Operating Systems

## Lecture 20: Need for Synchronization – 2

**Too much Milk! and Lock**

Ravi Mittal  
IIT Goa

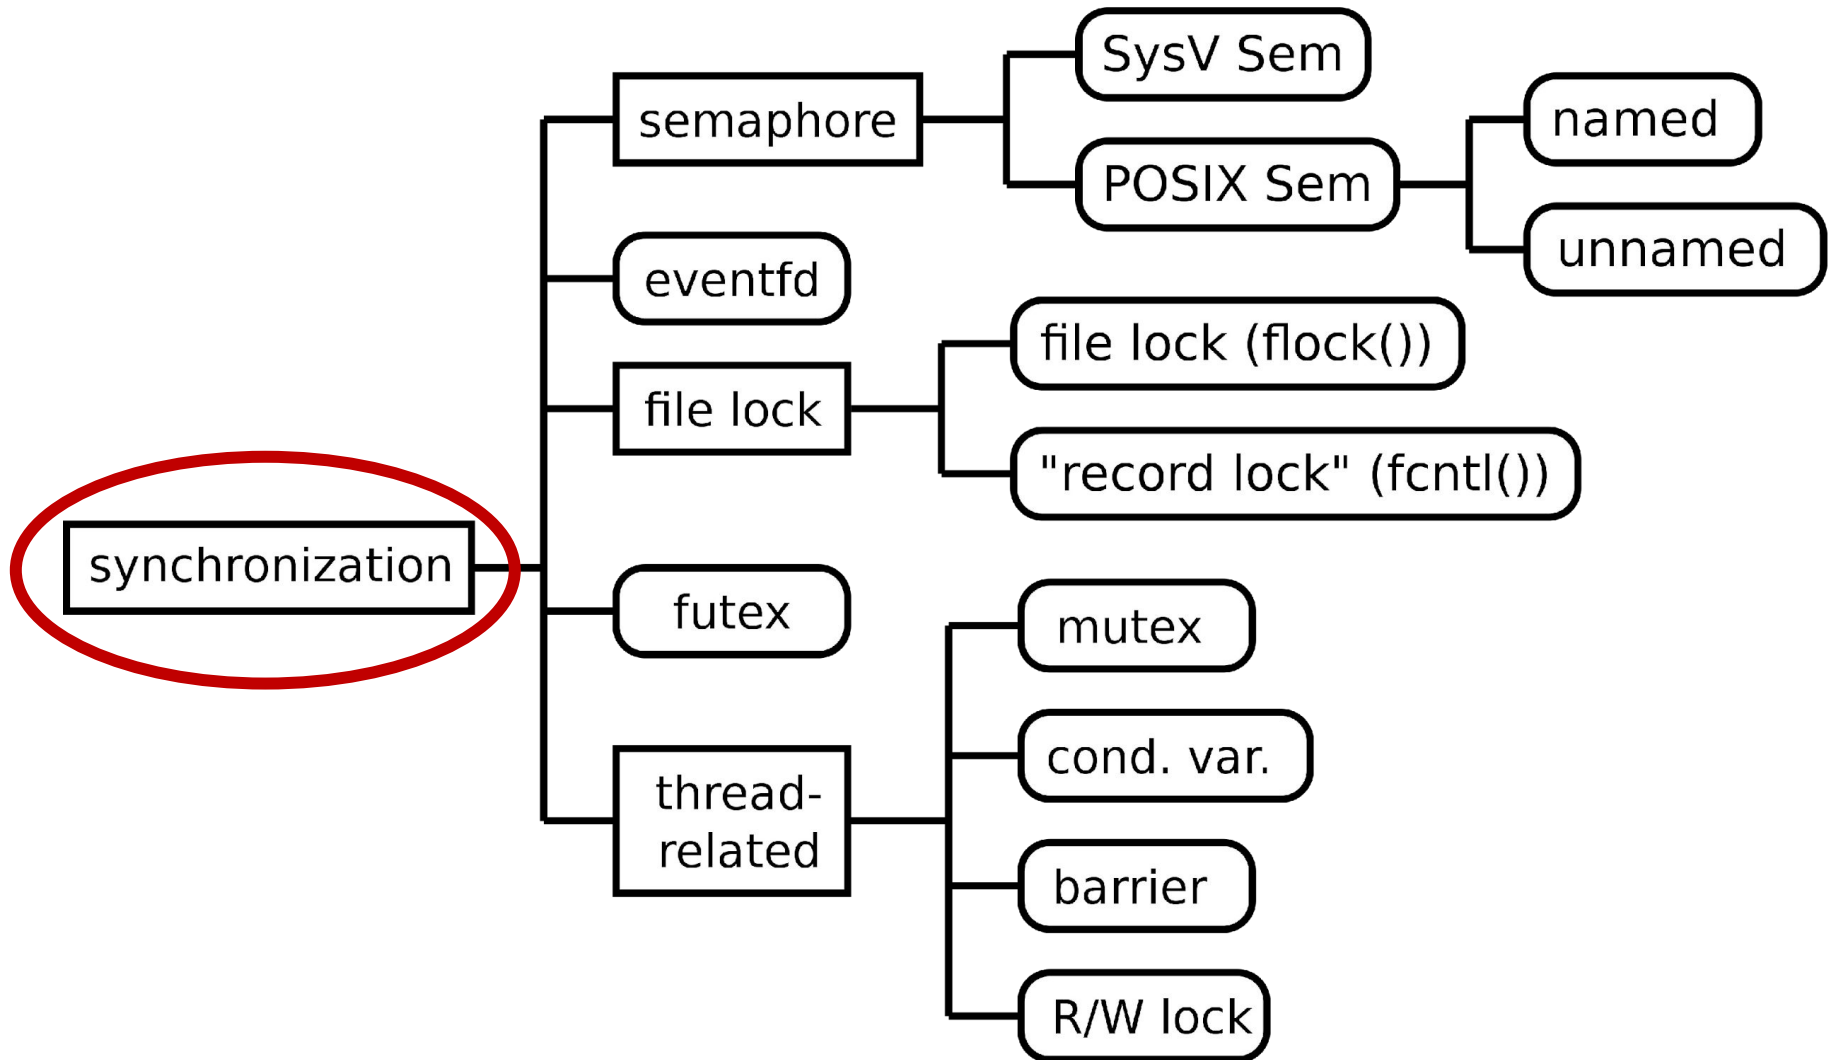
# So far we have studied

- Threads
- Processes
- Concurrent execution of Threads and Processes require
  - Communication
  - Synchronization
- In the last class we looked at two cases of race condition with
  - Concurrent Processes
  - Concurrent Threads

# Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
  - CS162, Operating System and Systems Programming, University of California, Berkeley
  - Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II, Chapter 5
    - Section 5.1

# Needs for Synchronization



# We will start with High level primitives

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Others
Hardware	Disable Ints Test&Set Compare&Swap, others

- Our focus will be on concepts

# Today we will study ..

- Example: Too much Milk
- Lock definition and properties
- Pthread mutex
- Example: Counting with two threads

**Too much Milk !**

# Motivating Example: “Too Much Milk”

- Great thing about OS's – analogy between problems in OS and problems in real life
  - Help you understand real life problems better
  - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away



# Use **lock** to fix Milk problem

- Fix the milk problem by putting a key on the refrigerator
  - Lock it and take key if you are going to go buy milk
  - **Fixes too much:** roommate angry if only wants **Beer**




# Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
- What are the correctness properties for the “Too much milk” problem???
  - Never more than one person buys
  - Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

# Too Much Milk: **Solution #1**

(1/3)

- **Use a note to avoid buying too much milk:**
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- **Suppose a computer tries this**
  - Remember, only memory read/write are atomic



```
if (noMilk) {  
  if (noNote) {  
    leave Note;  
    buy milk;  
    remove note;  
  }  
}
```

**Does it work properly?**

# Too Much Milk: Solution #1

(2/3)

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):


```
Thread A      Thread B
if (noMilk) {
  if (noMilk) {
    if (noNote) {
      if (noNote) {
        leave Note;
        buy Milk;
        remove Note;
      }
    }
  }
}

      leave Note;
      buy Milk;

      remove Note;
```

# Too Much Milk: Solution #1

(3/3)



- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
        if (noMilk) {  
        if (noNote) {  
            leave Note;  
            buy milk;  
            remove note;  
        }  
    }
```

- Result?
  - Still too much milk **but only occasionally!**
  - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
  - Makes it really hard to debug...
  - Must work despite what the dispatcher does!

# Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
  - Let's try to fix this by placing note first
- Another try at previous solution:

```
    leave Note;  
    if (noMilk) {  
        if (noNote) {  
            buy milk;  
        }  
    }  
    remove Note;
```

- What happens here?
  - Well, with human, probably nothing bad
  - With computer: no one ever buys milk




## Too Much Milk: **Solution #2**

- How about labeled notes?
  - Now we can leave note before checking
- Algorithm looks like this:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
if (noNote B) {	if (noNoteA) {
if (noMilk) {	if (noMilk) {
buy Milk;	buy Milk;
}	}
}	}
remove note A;	remove note B;

- Does this work?
- **Possible for neither thread to buy milk**
  - Context switches at exactly the wrong times can lead each to think that the other is going to buy

# Too Much Milk: **Solution #2: problem!**



- *I'm not getting milk, You're getting milk*
- This kind of lockup is called starvation!



# Too Much Milk Solution #3 (self reading)

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {	if (noNote A) {
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	
}	remove note B;
remove note A;	

- Does this work? **Yes**. Both can guarantee that:
  - It is safe to buy, or
  - Other will buy, ok to quit
- At X:
  - If no note B, safe for A to buy,
  - Otherwise wait to find out what will happen
- At Y:
  - If no note A, safe for B to buy

# Solution #3 discussion (self reading)

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
        if (noMilk) {  
            buy milk;  
        }
```

- Solution #3 works, but it’s really unsatisfactory
  - Really complex
    - Hard to convince yourself that this really works
  - A’s code is different from B’s – what if lots of threads?
    - Code would have to be slightly different for each thread
  - While A is waiting, it is consuming CPU time
    - This is called “busy-waiting”
- There’s a better way
  - Have hardware provide higher-level primitives than atomic load & store
  - Build even higher-level programming abstractions on this hardware support

# Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock
  - `lock.Acquire()` – wait until lock is free, then grab
  - `lock.Release()` – Unlock, waking up anyone waiting
    - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:


```
milklock.Acquire();  
if (nomilk)  
    buy milk;  
milklock.Release();
```

**Critical section**

# **Lock Definition and use**

# Lock Definition

- Best mainstream solution: **Locks**
  - Implements **mutual exclusion**
    - You can't have it if I have it, I can't have it if you have it
- Two states of a lock: **BUSY or FREE**
- A lock is initially in the **FREE** state
- **lock.Acquire()**
  - Waits until the lock is **FREE** and then atomically makes the lock **BUSY**
  - Checking the state to see if it is **FREE** and setting the state to **BUSY** are together an atomic operation
  - Multiple threads try to acquire the lock, at most one thread will succeed
  - One thread observes that the lock is **FREE** and sets it to **BUSY**; the other threads just see that the lock is **BUSY** and wait



# Lock Definition

- `lock.Release()`
  - makes the lock FREE
  - If there are pending acquire operations, this state change causes one of them to proceed
- *Only the lock “owner” should release the lock*
- Different OSs have different names for these operations of acquiring and releasing locks
  - `lock()`
  - `Unlock()`

# Lock Properties

- Mutual Exclusion (safety property)
  - At most one thread holds the lock
    - Only one thread in the critical section
- Progress (liveness property)
  - If no thread holds the lock and any thread attempts to acquire the lock, then eventually some thread succeeds in acquiring the lock
- Bounded Waiting (liveness property)
  - If thread T attempts to acquire a lock, then there exists a bound on the number of times other threads can successfully acquire the lock before T does

# Lock Non-Property: Thread Ordering

- The **bounded waiting property** guarantees that a thread will eventually get a chance to acquire the lock
- However, there is no guarantee that waiting threads acquire the lock in FIFO order



# Lock Use

- Consider an update of shared variable
  - `balance = balance + 1;`
- This statement forms critical section. We need to protect it with a special lock variable (`mutex` in the example below)

```
1  lock_t mutex; // some globally-allocated lock 'mutex'
2  ...
3  lock(&mutex);
4  balance = balance + 1;
5  unlock(&mutex);
```

- Note that lock is a variable that holds the state of the lock at any instant in time
- All threads accessing a critical section share a lock
- Only one thread becomes successful in holding the lock – current owner of the lock
  - Thus the thread is in critical section

# Lock Use

- `lock()` and `unlock()` are routines
- The thread that acquires the lock enters the critical section
- Then, if another thread calls the lock on the **same lock variable**, it will **not return**
  - So the thread is prevented from entering critical section
  - Function `lock()` doesn't return means it doesn't come out of `lock()` function
- Note that the owner of lock calls `unlock`
- Thread entities are created by the programmer but scheduled by the OS
  - Locks give some control to the programmer
- `pthread` library in Linux provides such locks

# Pthread lock calls

# Pthread lock calls

- **pthread\_mutex\_init**
  - Creates a new lock in the **unlocked** state
- **pthread\_mutex\_lock**
  - When the lock is **unlocked**, change the lock to the **locked** state and advance to the next line of code.
  - When the lock is **locked**, this function **blocks** execution of other threads until the lock can be acquired
- **pthread\_mutex\_unlock**
  - Moves the lock to the **unlocked** state
- Let's use these APIs in our example.

# OS Library Locks: *pthread*s

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                        const pthread_mutexattr_t *attr)
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

# Mutex Variables

- A typical sequence in the use of a mutex
  1. Create and initialize **mutex**
  2. Several threads attempt to lock **mutex**
  3. Only one succeeds and now owns **mutex**
  4. The owner performs some set of actions
  5. The owner unlocks **mutex**
  6. Another thread acquires **mutex** and repeats the process
  7. Finally **mutex** is destroyed

**Example**

# thread3.c

```
5  int ct = 0;
6
7  void *thread_start(void *ptr) {
8      int countTo = *((int *)ptr);
9
10     int i;
11     for (i = 0; i < countTo; i++) {
12         ct = ct + 1;
13     }
14
15     return NULL;
16 }
17
18 int main(int argc, char *argv[]) {
19     // Parse Command Line:
20     if (argc != 3) {
21         printf("Usage: %s <countTo> <thread count>\n",
22 argv[0]);
23         return 1;
24     }
25
26     const int countTo = atoi(argv[1]);
27     if (countTo == 0) { printf("Valid `countTo` is
28 required.\n"); return 1; }
29
30     const int thread_ct = atoi(argv[2]);
31     if (thread_ct == 0) { printf("Valid thread count is
32 required.\n"); return 1; }
33
34     // Create threads:
35     int i;
36     pthread_t tid[thread_ct];
37     for (i = 0; i < thread_ct; i++) {
38         pthread_create(&tid[i], NULL,
39 thread_start, (void *)&countTo);
40     }
41
42     // Join threads:
43     for (i = 0; i < thread_ct; i++) {
44         pthread_join(tid[i], NULL);
45     }
46
47     // Display result:
48     printf("Final Result: %d\n", ct);
49     return 0;
50 }
```



# thread3.c - Observations

- Command line inputs
  - countTo : Each thread counts upto CountTo
  - thread\_conunt: number of threads
- For example:
  - countTo = 1000
  - thread\_count = 2
  - ☐ final result = 2000
- However, we see that we don't get exact result for all values of countTo and thread\_count
  - At architecture level  $ct = ct + 1$  is not atomic
  - For example for the first thread – after a few counts – it is context switched after a few steps:
    - Lw s1, ct (assuming ct is in a memory location)
    - Addi s1, 1
    - Context switched.... Other threads now change ct variable
    - After some time the first thread completes operation

# Synchronization using mutex

- Mutex : mutual exclusive (to each thread)
- The simplest way to protect a region of code from being accessed is through the use of a **mutex** lock
  - **mutex** stands for *mutual exclusive*

# thread3.c

```
5  int ct = 0;
6
7  void *thread_start(void *ptr) {
8      int countTo = *((int *)ptr);
9
10     int i;
11     for (i = 0; i < countTo; i++) {
12         ct = ct + 1;
13     }
14
15     return NULL;
16 }
17
18 int main(int argc, char *argv[]) {
19     // Parse Command Line:
20     if (argc != 3) {
21         printf("Usage: %s <countTo> <thread count>\n",
22             argv[0]);
23         return 1;
24     }
25
26     const int countTo = atoi(argv[1]);
27     if (countTo == 0) { printf("Valid `countTo` is
28         required.\n"); return 1; }
29
30     const int thread_ct = atoi(argv[2]);
31     if (thread_ct == 0) { printf("Valid thread count is
32         required.\n"); return 1; }
33
34     // Create threads:
35     int i;
36     pthread_t tid[thread_ct];
37     for (i = 0; i < thread_ct; i++) {
38         pthread_create(&tid[i], NULL,
39             thread_start, (void *)&countTo);
40     }
41
42     // Join threads:
43     for (i = 0; i < thread_ct; i++) {
44         pthread_join(tid[i], NULL);
45     }
46
47     // Display result:
48     printf("Final Result: %d\n", ct);
49     return 0;
50 }
```

Non-atomic operation

Critical section – that should be  
accessed by only one thread at a  
time

# thread3d.c


```
5 pthread_mutex_t lock;
6 int ct = 0;
7
8 void *thread_start(void *ptr) {
9     int countTo = *((int *)ptr);
10
11     int i;
12     for (i = 0; i < countTo; i++) {
13         pthread_mutex_lock(&lock);
14         ct = ct + 1;
15         pthread_mutex_unlock(&lock);
16     }
17
18     return NULL;
19 }
20
21 int main(int argc, char *argv[]) {
22     // Parse Command Line:
23     if (argc != 3) {
24         printf("Usage: %s <countTo> <thread count>\n",
25             argv[0]);
26         return 1;
27     }
28
29     const int countTo = atoi(argv[1]);
30     if (countTo == 0) { printf("Valid `countTo` is
31         required.\n"); return 1; }
32
33     const int thread_ct = atoi(argv[2]);
34     if (thread_ct == 0) { printf("Valid thread count is
35         required.\n"); return 1; }
36
37     // Create Lock:
38     pthread_mutex_init(&lock, NULL);
39 }
```

1 Declare lock variable

3 Locking and unlocking  
Critical section

2 initialize lock to null (open)

# Summary

- Concurrent threads are a very useful abstraction
  - Concurrent threads introduce problems when accessing shared data
    - Programs must be insensitive to arbitrary interleavings
    - Without careful design, shared variables can become completely inconsistent
  - Important concept: Atomic Operations
    - An operation that runs to completion or not at all
    - These are the primitives on which to construct various synchronization primitives
  - We will study several mechanisms of implementing locks
- 

# Backup Slides

# pthread lock calls

- The name that the POSIX library uses for a lock is a **mutex**
  - `pthread_mutex_init()` - get a mutex
  - `pthread_mutex_lock()` - lock a mutex (acquire it), block until available
  - `pthread_mutex_trylock()` - try to lock a mutex (acquire it), do not block if unavailable
  - `pthread_mutex_unlock()` - unlock a mutex (release it)
  - `pthread_mutex_destroy()` - destroy a mutex (remove it)

The name that the POSIX library uses  
for a lock is a mutex,

# Pthread Synchronization

- Two primitives
  - Mutex
    - Semaphore with maximum value 1
  - Condition variable
    - Provides a shared signal
    - Combined with a mutex for synchronization



# Pthread Mutex

- States
  - Locked
    - Some thread holds the mutex
  - Unlocked
    - No thread holds the mutex
- When several threads compete
  - One wins
  - The rest block
    - Queue of blocked threads

# Mutex Variables

- A typical sequence in the use of a mutex
  1. Create and initialize **mutex**
  2. Several threads attempt to lock **mutex**
  3. Only one succeeds and now owns **mutex**
  4. The owner performs some set of actions
  5. The owner unlocks **mutex**
  6. Another thread acquires **mutex** and repeats the process
  7. Finally **mutex** is destroyed

# Creating a mutex

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
    const pthread_mutexattr_t *attr);
```

- Initialize a pthread mutex: the mutex is initially unlocked
- Returns
  - 0 on success
  - Error number on failure
    - **EAGAIN**: The system lacked the necessary resources; **ENOMEM**: Insufficient memory ;
    - **EPERM**: Caller does not have privileges; **EBUSY**: An attempt to re-initialise a mutex;
    - **EINVAL**: The value specified by attr is invalid
- Parameters
  - **mutex**: Target mutex
  - **attr**:
    - NULL: the default mutex attributes are used
    - Non-NULL: initializes with specified attributes

# Creating a mutex

- Default attributes
  - Use **PTHREAD\_MUTEX\_INITIALIZER**
    - Statically allocated
    - Equivalent to dynamic initialization by a call to **pthread\_mutex\_init()** with parameter **attr** specified as NULL
    - No error checks are performed

# Destroying a mutex

```
#include <pthread.h>
```

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

- Destroy a pthread mutex
- Returns
  - 0 on success
  - Error number on failure
    - **EBUSY**: An attempt to re-initialise a mutex; **EINVAL**: The value specified by attr is invalid
- Parameters
  - **mutex**: Target mutex

# Locking/unlocking a mutex

```
#include <pthread.h>
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Returns
  - 0 on success
  - Error number on failure
    - **EBUSY**: already locked; **EINVAL**: Not an initialised mutex; **EDEADLK**: The current thread already owns the mutex; **EPERM**: The current thread does not own the mutex

# How to implement Locks?

- How can we build multi-instruction atomic operations?
  - Recall: dispatcher gets control in two ways
    - Internal: Thread does something to relinquish the CPU
    - External: Interrupts cause dispatcher to take CPU
  - On a uniprocessor, can avoid context-switching by:
    - Avoiding internal events
    - Preventing external events by disabling interrupts

# Naïve use of Interrupt Enable/Disable

- Naïve Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

- Problems with this approach:

- Can't let user do this!

- Consider following:

```
    LockAcquire();  
    while(TRUE) {;}    // infinite loop
```

- Critical Sections might be arbitrarily long
- What happens with I/O or other important events?
  - Flight control



**backup**

# Too Much Milk Solution #3 (self reading)

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {	if (noNote A) {
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk; }	
}	remove note B;
remove note A;	

- Does this work? **Yes**. Both can guarantee that:
  - It is safe to buy, or
  - Other will buy, ok to quit
- At X:
  - If no note B, safe for A to buy,
  - Otherwise wait to find out what will happen
- At Y:
  - If no note A, safe for B to buy

# Solution #3 discussion (self reading)

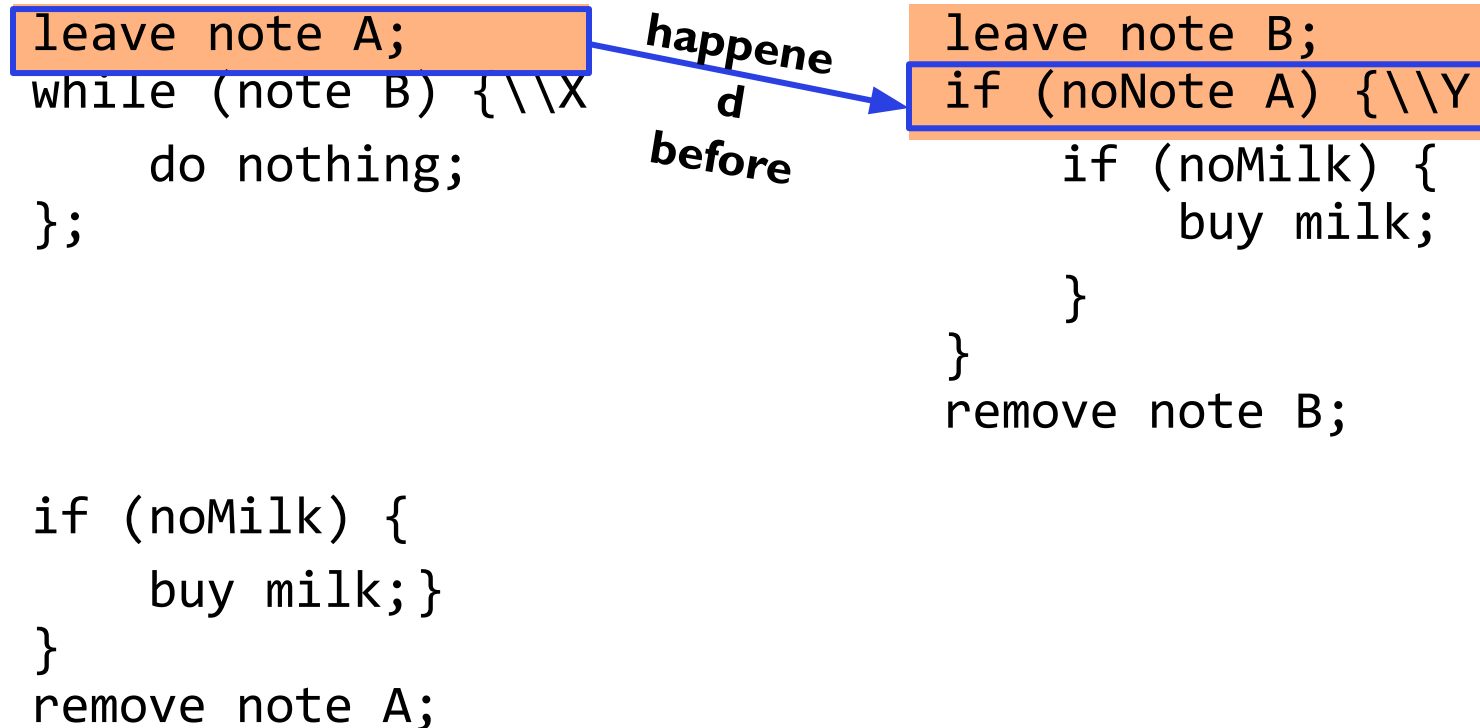
- Our solution protects a single “Critical-Section” piece of code for each thread:

```
        if (noMilk) {  
            buy milk;  
        }
```

- Solution #3 works, but it’s really unsatisfactory
  - Really complex
    - Hard to convince yourself that this really works
  - A’s code is different from B’s – what if lots of threads?
    - Code would have to be slightly different for each thread
  - While A is waiting, it is consuming CPU time
    - This is called “busy-waiting”
- There’s a better way
  - Have hardware provide higher-level primitives than atomic load & store
  - Build even higher-level programming abstractions on this hardware support

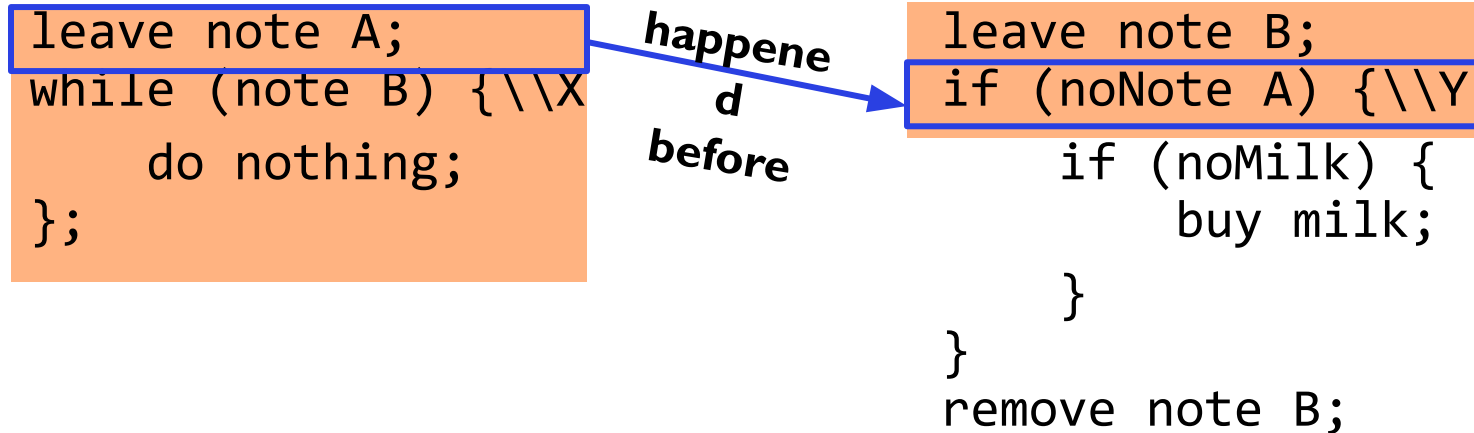
# Case 1

- “leave note A” happens before “if (noNote A)”



# Case 1

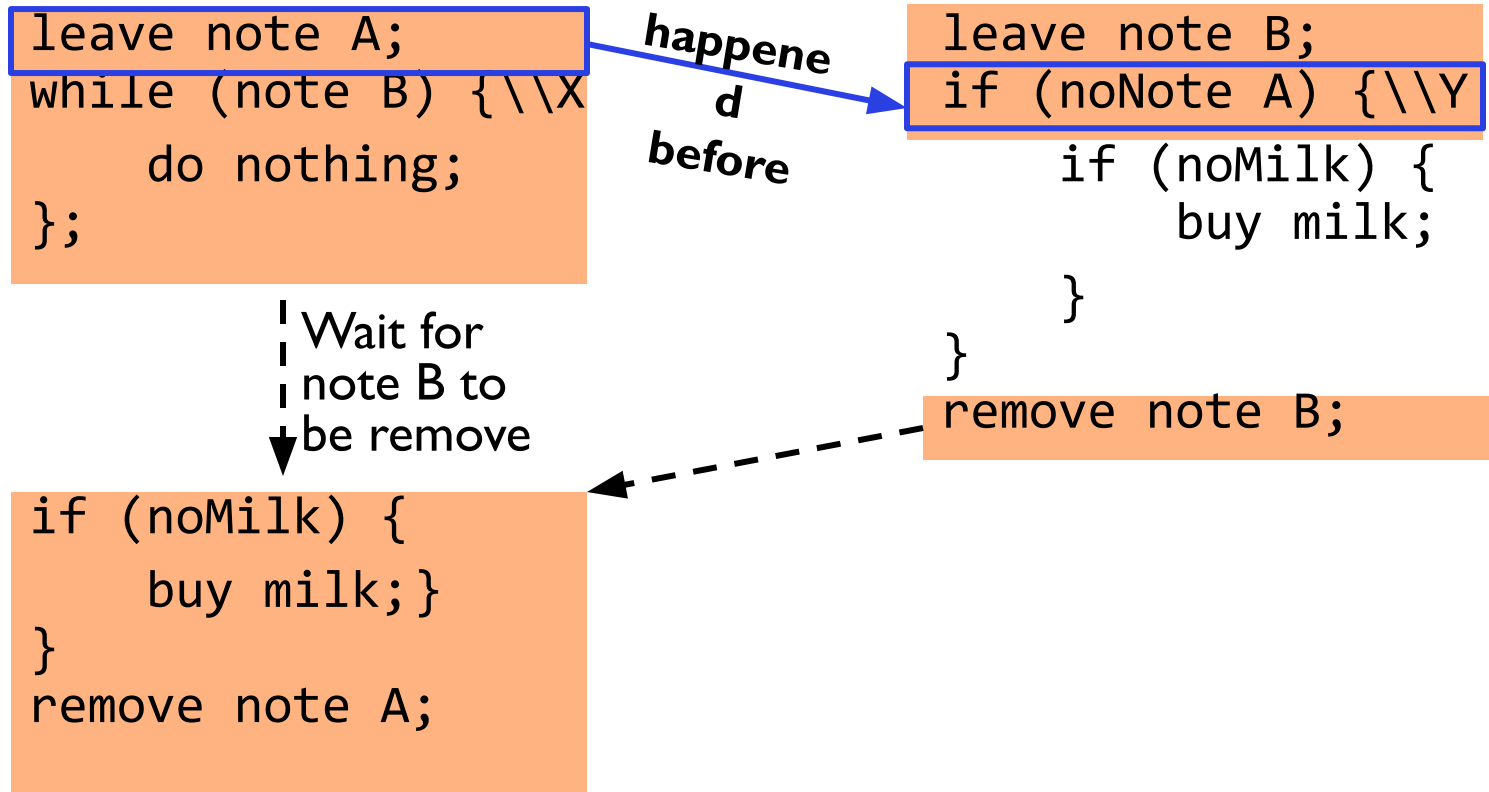
- “leave note A” happens before “if (noNote A)”



```
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

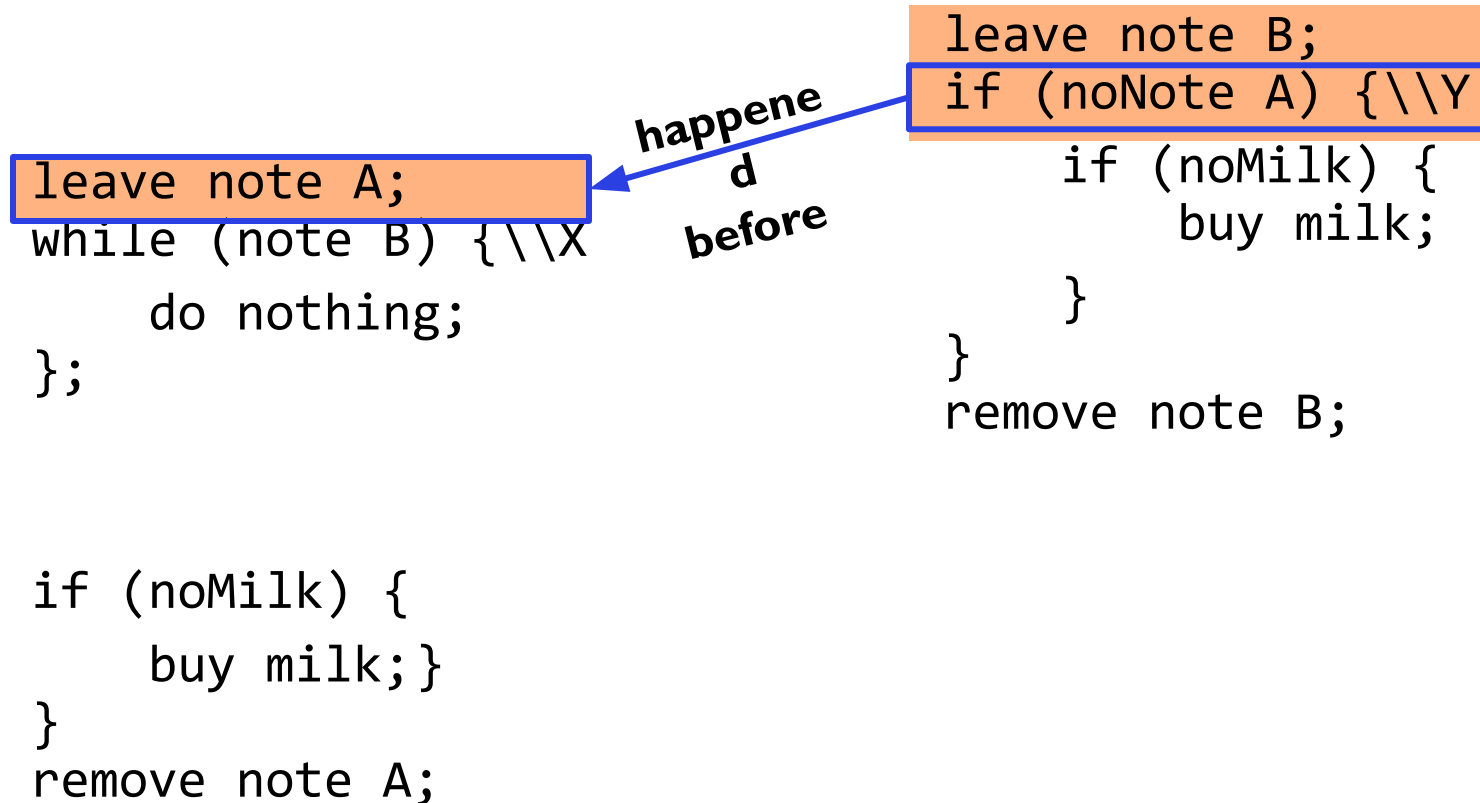
# Case 1

- “leave note A” happens before “if (noNote A)”



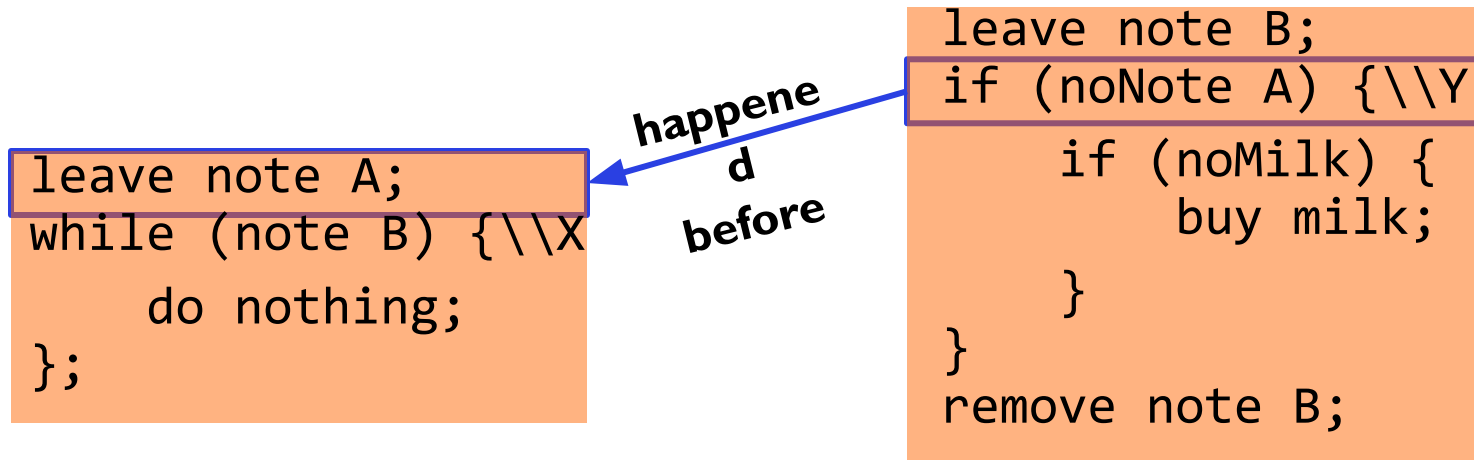
## Case 2

- “if (noNote A)” happens before “leave note A”



## Case 2

- “if (noNote A)” happens before “leave note A”

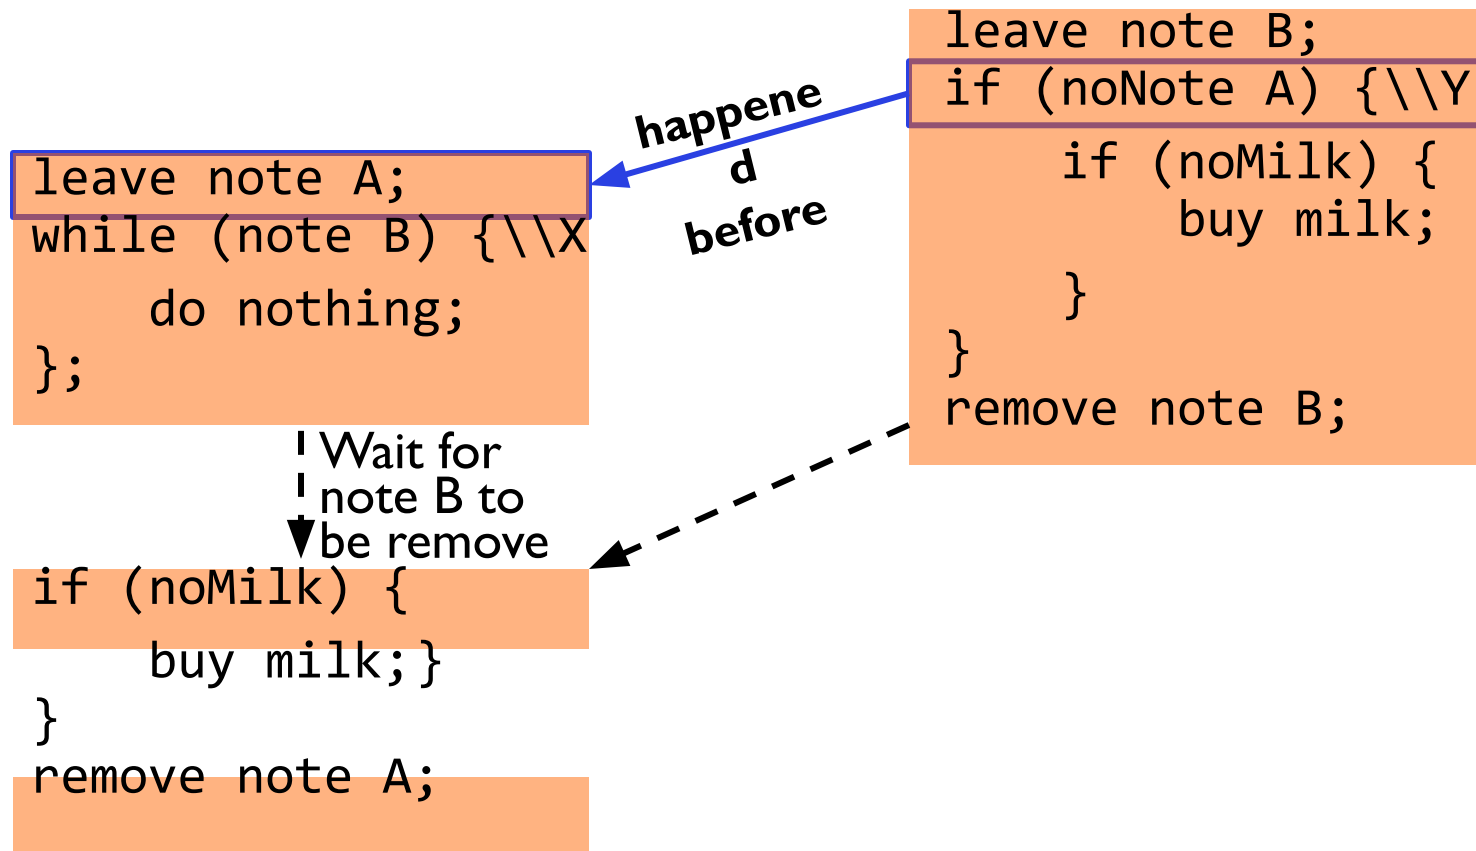


```
if (noMilk) {  
    buy milk;}  
}  
remove note A;
```



## Case 2

- “if (noNote A)” happens before “leave note A”



# Atomic Operations

- Indivisible operations that cannot be interleaved with or split by other operations
- What is the need for atomic operations?