

CS310 Operating Systems

Lecture 26 Scheduling – Multilevel Feedback Queue, Lottery Scheduling

Ravi Mittal
IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, University of California, Berkeley
 - Book: [Operating System: Three Easy Pieces](#)

Reading

- CS162, Operating System and Systems Programming, University of California, Berkeley
- Book: Operating System: Three Easy Pieces

Previous lectures ..

First-Come, First-Served (FCFS) Scheduling

- First-Come, First-Served (FCFS)
 - Also **First In First Out (FIFO)** or **Run until done**
 - In early systems, FCFS meant one program scheduled until done (including I/O)
 - Now, means keep CPU until thread blocks
- Simple Algorithm, Easy to implement (+)
- FCFS Scheme: Potentially bad for short jobs!
 - Depends on submit order
 - If you are first in line at supermarket with milk, you don't care who is behind you, on the other hand...
 - **Convoy effect:** short process stuck behind long process (-)



Shortest Job First (SJF)

- Non-preemptive
- Run whatever job has least amount of computation to do
- Provably optimal
- Need to know run times in advance

Shortest Remaining Time First (SRTF)


- Preemptive version of SJF
- If job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
- Sometimes called Shortest Remaining Time to Completion First (SRTCF)
- Both SJF and SRTF:
 - These can be applied to whole program or current CPU burst
 - Idea is to get short jobs out of the system
 - Big effect on short jobs, only small effect on long ones
 - Result is better average response time

Round Robin (RR) Scheduling



- Uses **Preemption!**
- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds
- After quantum expires, the process is preempted and added to the end of the ready queue
- n processes in ready queue and time quantum is q
 - Each process gets $1/n$ of the CPU time
 - In chunks of at most q time units
 - No process waits more than $(n-1)q$ time units

Multilevel Queue Scheduling – Strict Priority



- Execution Plan

- Always execute highest-priority runnable jobs to completion
- Each queue can be processed in RR with some time-quantum
- A priority is assigned statically to each process, and a process remains in the same queue for the duration of the run time

- Problems

- Starvation
 - Lower priority jobs don't get to run because higher priority jobs
- Deadlock: Priority Inversion
 - Happens when low priority task holds a lock needed by high-priority task

Today we will study

- Multilevel Feedback Queue
- Changing landscape of scheduling
- Proportionate share scheduling
- Lottery Scheduling

Multilevel Feedback Queue

Dealing with Starvation

- Strict priority scheduling may lead to situation where low priority processes may starve
- One solution is
 - Don't assign priorities that are static – fixed for lifetime of the process
 - But, change priority of a process – dynamically
 - Scheduler can decrease the priority of the current running process
 - Penalizing it for taking too much CPU time
 - Alternatively, Scheduler keeps track of low priority processes that are not getting a chance to run
 - Increase their priority

Multilevel Feedback Queue

- Multilevel queues with no strict priority
- Scheduler is allowed to change priority of a process
 - Dynamic priority
- A multilevel feedback queue uses two basic rules:
 - A new process gets placed in the highest priority queue
 - If a process does not finish its quantum (due to I/O)
 - Then it will stay at the same priority level (round robin)
 - Else it moves to the next lower priority level
- A process with long CPU bursts will use its entire time slice
 - Get preempted and get placed in a lower-priority queue
 - A highly interactive process will not use up its quantum and will remain at a high priority level

Example of Multilevel Feedback Queue

- Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS

- Scheduling

- A new process enters queue Q_0 which is served in RR


- When it gains CPU, the process receives 8 milliseconds
- If it does not finish in 8 milliseconds, the process is moved to queue Q_1

- At Q_1 job is again served in RR and receives 16 additional milliseconds

- If it still does not complete, it is preempted and moved to queue Q_2

- Starvation problem still exists

- When new processes are frequently created or there are too many interactive processes □ CPU bound processes are never scheduled



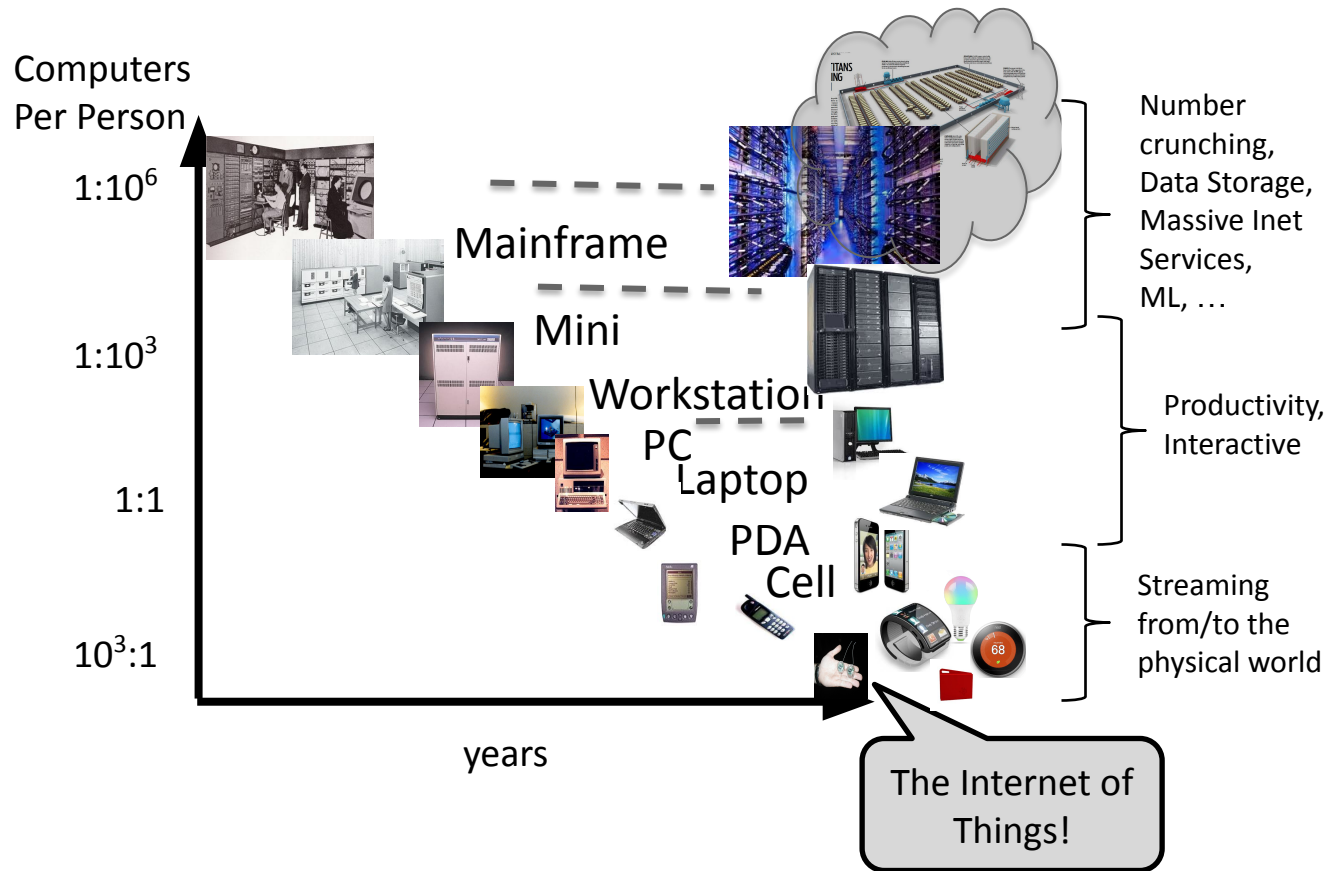
Cause for Starvation: **Priorities?**

- The policies we've studied so far:
 - Always prefer to give the CPU to a prioritized job
 - Non-prioritized jobs may never get to run
- But priorities were a means, not an end
- Our end goal was to serve a mix of CPU-bound, I/O bound, and Interactive jobs effectively on common hardware
 - Give the I/O bound ones enough CPU to issue their next file operation and wait (on those slow discs)
 - Give the interactive ones enough CPU to respond to an input and wait (on those slow humans)

Changing landscape ..

Changing Landscape...

Bell's Law: New computer class every 10 years



Changing Landscape of Scheduling

- Priority-based scheduling rooted in “time-sharing”
 - Allocating precious, limited resources across a diverse workload
 - CPU bound, vs interactive, vs I/O bound
- 80’s brought about personal computers, workstations, and servers on networks
 - Different machines of different types for different purposes
 - Shift to **fairness and avoiding extremes** (starvation)
- 90’s emergence of the web, rise of internet-based services, the data-center-is-the-computer
 - Server consolidation, massive clustered services, huge flash crowds
 - It’s about **predictability, 95th percentile performance guarantees**


Does prioritizing some jobs *necessarily* starve those that aren't prioritized?

Key Idea: Proportional-Share Scheduling

- The policies we've studied so far:
 - Always prefer to give the CPU to a prioritized job
 - Non-prioritized jobs may never get to run
- Instead, we can share the CPU *proportionally*
 - Give each job a share of the CPU according to its priority
 - Low-priority jobs get to run less often
 - But all jobs can at least make progress (no starvation)

Lottery Scheduling (Fair Share Scheduling)

Lottery Scheduling



- Given a set of jobs (the mix), provide each with a share of a resource
 - e.g., 50% of the CPU for Job A, 30% for Job B, and 20% for Job C
- Idea: Give out tickets according to the proportion each should receive
- Every quantum: draw one at random, schedule that job (thread) to run

Lottery Scheduling



- Proportional share scheduling
 - Give each job some number of lottery tickets
 - On each time slice, randomly pick a winning ticket
 - On average, CPU time is proportional to number of tickets given to each job
- How to assign tickets? **Users can define policy**
 - To approximate SRTF, short running jobs get more, long running jobs get fewer
 - To avoid starvation, every job gets at least one ticket (everyone makes progress)
- Advantage over strict priority scheduling: behaves gracefully as load changes
 - Adding or deleting a job affects all jobs proportionally, independent of how many tickets each job possesses

Use Randomness

Requires a good random number generator

Lottery Scheduling Example (Cont.)

- Lottery Scheduling Example
 - Assume short jobs get 10 tickets, long jobs get 1 ticket (total 11 tickets)

# short jobs/ # long jobs	% of CPU each short jobs gets	% of CPU each long jobs gets
1/1	91%	9%
0/2	N/A	50%
2/0	50%	N/A
10/1	9.9%	0.99%
1/10	50%	5%

- How to do ticket distribution when processes come and go and get blocked?
- Not a useful algorithm for general-purpose scheduling
- Useful for environments with long-running processes that may need to be allocated shares of CPUs
 - Example: running multiple virtual machines on a server

Lottery Fairness

- Two jobs competing against each other
- Both have the same number of tickets say 100
- Both have the same run time R
- Ideally each job must finish roughly at the same time
 - Due to randomness, sometimes one job finishes before the other
- Fairness Metric: F
 - Time the first job completes / time the second job completes
 - Ex: $R = 10$; First job completes at time 10 and second job completes at time 20; $F = 10/20 = 0.5$
 - When both finish at nearly the same time, $F = 1$

A simulation

- R varies from 1 to 1000 over 30 trials




Figure 9.2: Lottery Fairness Study

How to Assign Tickets?

- How many to each user?
- User knows the best
 - Each user is handed over a certain number of tickets
 - User can allocate them to his/her jobs
- Ticket Assignment problem remains open.

Lecture Summary

- Multilevel Feedback Queue
 - It has multiple levels of queues
 - It uses feedback to determine priority of each job
 - Based on how jobs behave over time
 - It is not based on concept of prior knowledge of the nature of the job ... but observes the execution of a job and prioritize accordingly
- We have studied the concept of proportional – share scheduling
 - Lottery Scheduling
- Lottery scheduling uses randomness in a clever way to achieve proportional share