

CS310 Operating Systems

Lecture 28: Introduction to Semaphores and Deadlocks

Ravi Mittal
IIT Goa

In the class we will study

- Semaphore Introduction
- Deadlock Introduction

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, Sam Kumar, University of California, Berkeley,
 - Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II, Chapter 5
 - Book: Modern Operating Systems, Fourth Edition, Andrew Tenenbaum, Herbert Bos, Pearson Publication
 - Chapter 2.3

Reading

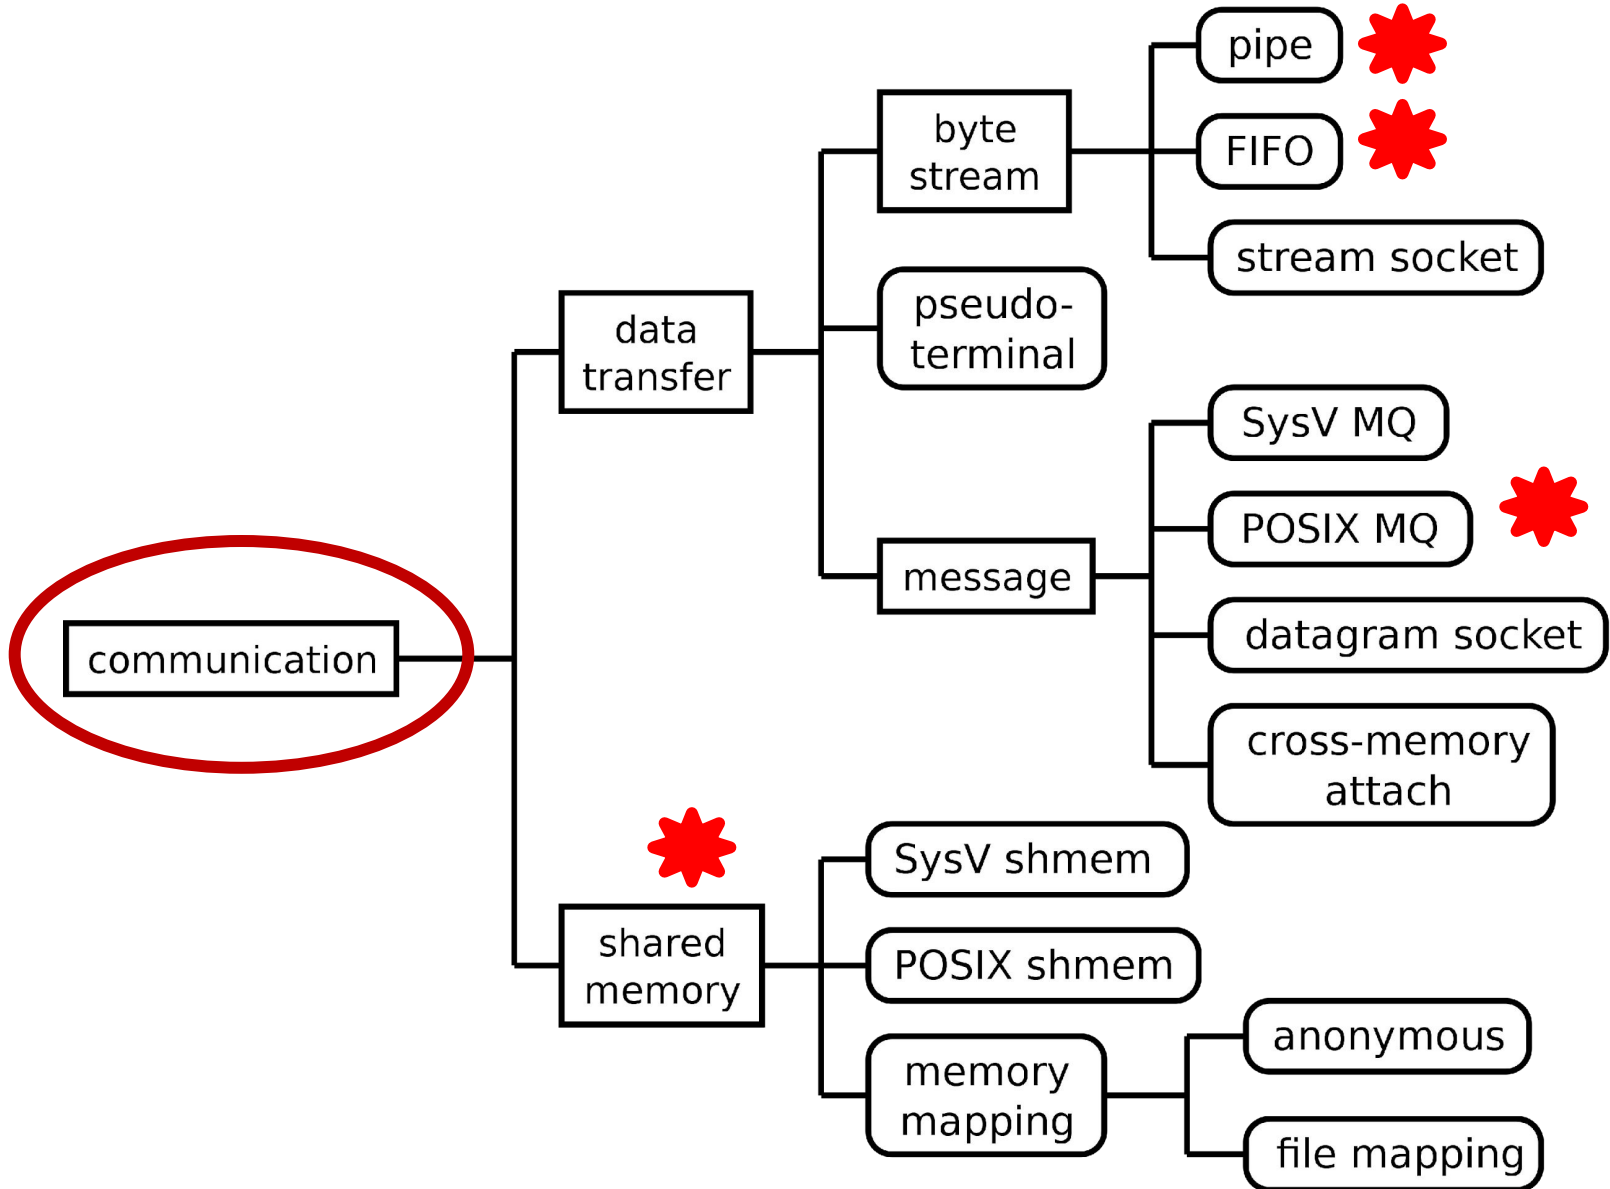
- Book: Operating System: Three Easy Pieces, by Remzi H Arpaci-Dusseau, Andrea C Arpaci-Dusseau
 - Chapter 31,
<https://pages.cs.wisc.edu/~remzi/OSTEP/threads-sema.pdf>
- Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II, Chapter 5.8

Previous Classes on IPC and Synchronization

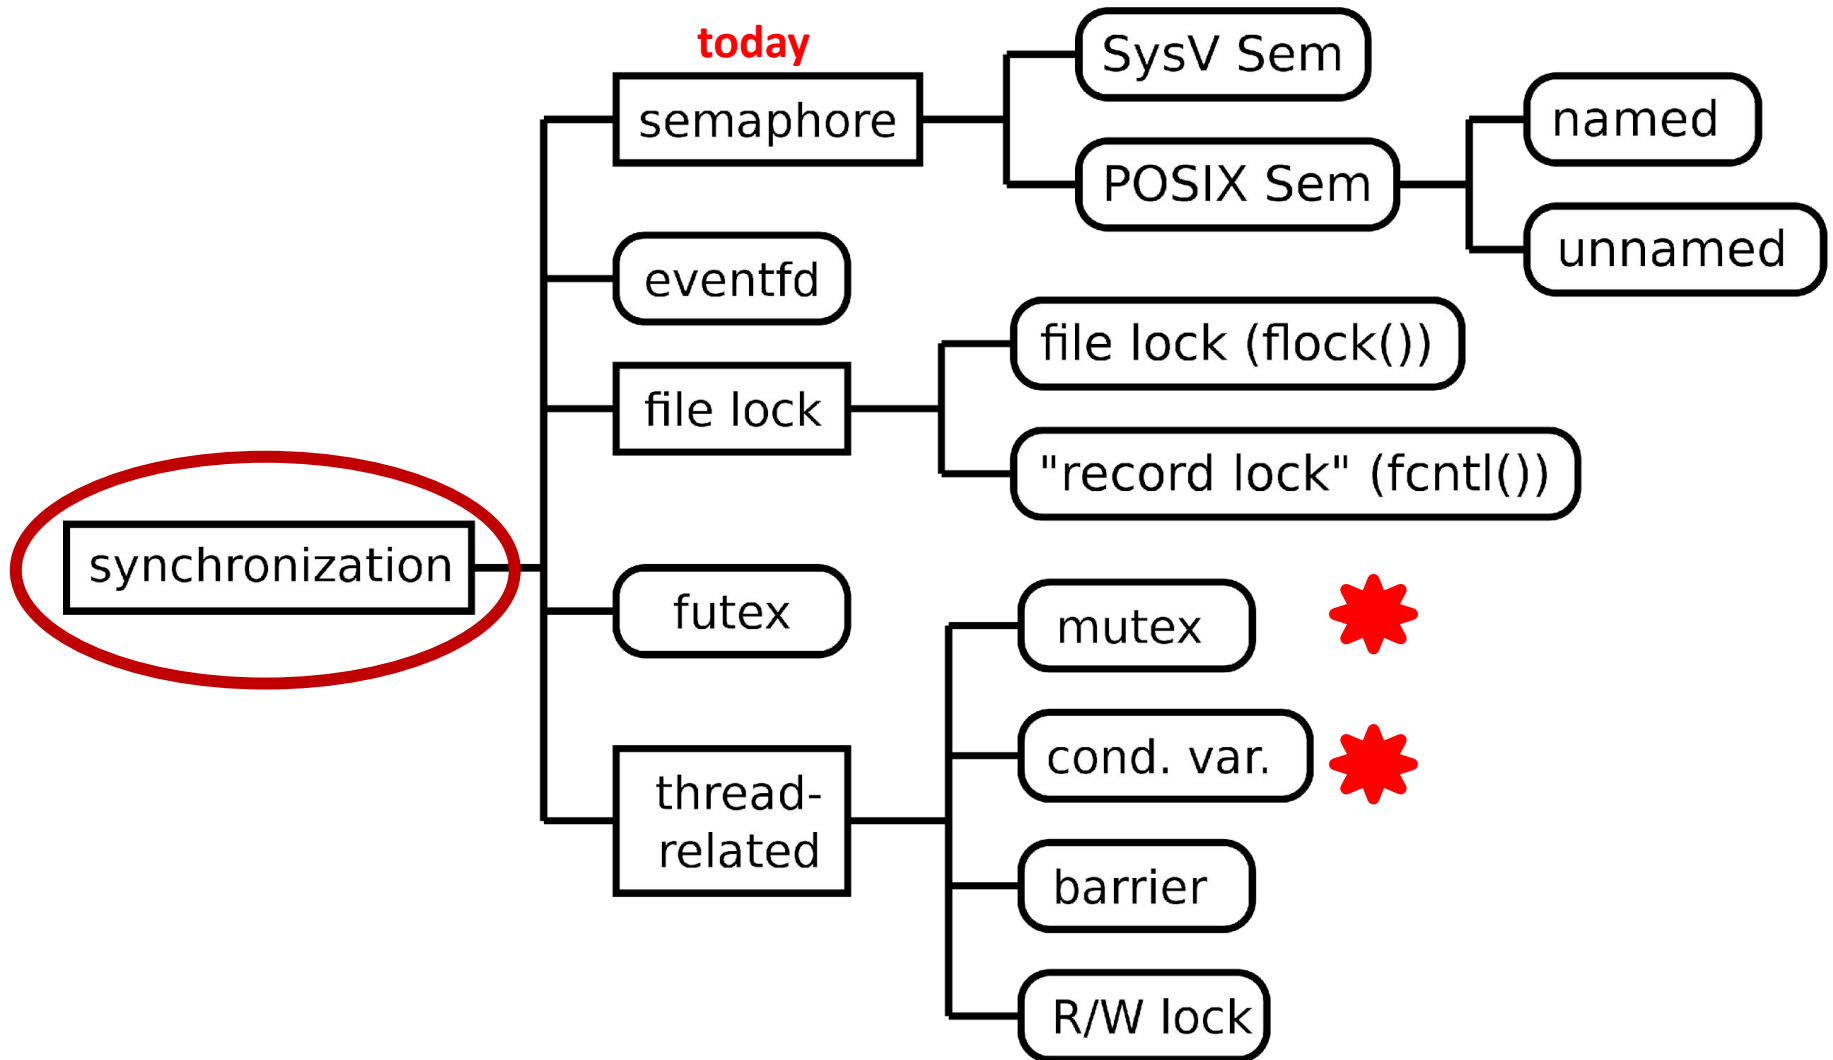
So far we have studied

- Threads
- Processes
- Concurrent execution of Threads and Processes require
 - Communication
 - Synchronization
- Inter-process Communication methods
 - Message Passing
 - Message Queues
 - Pipes
 - Named Pipes or FIFO
 - Shared Memory
- Synchronization
 - Lock, Lock implantation: Test-and-Set, Compare-and-swap
 - Condition Variables

Communication



Needs for Synchronization



We will start with High level primitives

Programs	Shared Programs		
Higher-level API	Locks	Semaphores	Monitors Others
Hardware	Disable Ints	Test&Set	Compare&Swap, others

- Today we will study Semaphores

Semaphore Introduction



Semaphore

- Semaphore = a synchronization primitive
 - Higher level of abstraction than locks
 - invented by Dijkstra in 1968, as part of the THE operating system
 - Synchronization tool that does not require busy waiting
 - So it doesn't waste CPU time
 - An integer value used for signaling among processes
- Fundamental Principle:
 - Two or more processes want to cooperate by means of simple signals
- Semaphores can be used both as locks and condition variables

Semaphore Definition

- A semaphore is an object with an integer value
 - We can manipulate semaphore with two routines:
 - `sem_wait()`
 - `sem_post()`
- The initial value to the semaphore determines its behavior
 - So, we must initialize it to some value
- Initializing Semaphore:
 - `s`: semaphore
 - Initial value: 1
 - The third argument
- Second argument to `sem_init()` is set to 0
 - This indicates that the semaphore is shared between threads in a process
 - Other options: see man pages

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1);
```

Semaphore Definition

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1);
```

- After initialization we can call `sem_wait()` or `sem_post()`
- `sem_wait()`
 - Returns right away – if the value of semaphore is 1 or more when `sem_wait()` was called, Enter critical section
 - If semaphore value was **zero or negative** --> caller to suspend execution
 - Waits for subsequent post – puts itself to sleep
 - If multiple calling threads may call `sem_wait()`
 - Gets queued waiting to be woken
- `sem_post()`
 - It simply increments the value of the semaphore
 - If there is any thread waiting to be woken, one is woken up
- The negative value of the semaphore = number of waiting threads

Semaphore Definition

- The semaphore value is not seen by the users of semaphores
- `wait()` and `post()`

```
1  int sem_wait(sem_t *s) {  
2      decrement the value of semaphore s by one  
3      wait if value of semaphore s is negative  
4  }  
5  
6  int sem_post(sem_t *s) {  
7      increment the value of semaphore s by one  
8      if there are one or more threads waiting, wake one  
9  }
```

Binary Semaphore (Locks)

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize to X; what should X be?  
3  
4  sem_wait(&m);  
5  // critical section here  
6  sem_post(&m);
```

- What should be the value of X ?

Binary Semaphore (Locks)

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize to X; what should X be?  
3  
4  sem_wait(&m);  
5  // critical section here  
6  sem_post(&m);
```

- What should be the value of X ?
- **Answer: 1**

First Case: Single Thread using semaphore

- Consider two threads
 - Only one thread is using Semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sem_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

Second Case: Two threads using a Semaphore

- Initial Value of semaphore is 1
- Thread 0 holds the lock using `sem_wait()`
- Thread 1 tries to enter critical section by calling `sem_wait()`
 - Thread 1 decrements the value of semaphore to -1 and waits by putting itself to sleep
 - It **relinquishes** the processor
- Thread 0 runs again
 - Eventually calls `sem_post()`
 - Increments semaphore to 0
 - Wakes up sleeping thread 1
- Thread 1 now enters critical section
 - Note that value of semaphore remains 0 and thread 1 has already decremented the semaphore .. It's not checked again

Binary Semaphore (Locks) – Another case

Val	Thread 0	State	Thread 1	State
1		Run		Ready
1	call sem_wait()	Run		Ready
0	sem_wait() returns	Run		Ready
0	(crit sect begin)	Run		Ready
0	<i>Interrupt; Switch→T1</i>	Ready		Run
0		Ready	call sem_wait()	Run
-1		Ready	decr sem	Run
-1		Ready	(sem<0)→sleep	Sleep
-1		Run	<i>Switch→T0</i>	Sleep
-1	(crit sect end)	Run		Sleep
-1	call sem_post()	Run		Sleep
0	incr sem	Run		Sleep
0	wake(T1)	Run		Ready
0	sem_post() returns	Run		Ready
0	<i>Interrupt; Switch→T1</i>	Ready		Run
0		Ready	sem_wait() returns	Run
0		Ready	(crit sect)	Run
0		Ready	call sem_post()	Run
1		Ready	sem_post() returns	Run

Semaphore for Ordering

- We can use semaphores to order events in a concurrent program
- For example
 - A thread may wait for a list to become non-empty so that it can delete an item from it
- In general, we can use semaphores to develop programs where
 - One thread is *waiting* for something to happen
 - Another thread making that something happen and then *signaling* that it has happened
 - Thus, waking the sleeping thread
 - Similar to use of condition variable
- Example: Imagine a thread creates another thread and then wants to wait for it to complete its execution

Example: A parent waiting for it's child

```
1  sem_t s;
2
3  void *child(void *arg) {
4      printf("child\n");
5      sem_post(&s); // signal here: child is done
6      return NULL;
7  }
8
9  int main(int argc, char *argv[]) {
10     sem_init(&s, 0, X); // what should X be?
11     printf("parent: begin\n");
12     pthread_t c;
13     Pthread_create(&c, NULL, child, NULL);
14     sem_wait(&s); // wait here for child
15     printf("parent: end\n");
16     return 0;
17 }
```

Example: A parent waiting for it's child

- Execution of the program results prints:
Parent: begin
child
Parent: end
- Value of the semaphore should be 0

Example: A parent waiting for it's child

- Case 1
 - Parent has created child but child has not run, yet (in ready queue)

Val	Parent	State	Child	State
0	create(Child)	Run	(Child exists, can run)	Ready
0	call sem_wait()	Run		Ready
-1	decr sem	Run		Ready
-1	(sem<0) → sleep	Sleep		Ready
-1	Switch → Child	Sleep	child runs	Run
-1		Sleep	call sem_post()	Run
0		Sleep	inc sem	Run
0		Ready	wake(Parent)	Run
0		Ready	sem_post() returns	Run
0		Ready	Interrupt → Parent	Ready
0	sem_wait() returns	Run		Ready

Example: A parent waiting for it's child

- Case 2
 - Child gets to completion before parent gets chance to call `wait()`

Val	Parent	State	Child	State
0	<code>create (Child)</code>	Run	<i>(Child exists; can run)</i>	Ready
0	<i>Interrupt</i> → <i>Child</i>	Ready	child runs	Run
0		Ready	<code>call sem_post ()</code>	Run
1		Ready	<code>inc sem</code>	Run
1		Ready	<code>wake (nobody)</code>	Run
1		Ready	<code>sem_post () returns</code>	Run
1	parent runs	Run	<i>Interrupt</i> → <i>Parent</i>	Ready
1	<code>call sem_wait ()</code>	Run		Ready
0	<code>decrement sem</code>	Run		Ready
0	<code>(sem ≥ 0) → awake</code>	Run		Ready
0	<code>sem_wait () returns</code>	Run		Ready


Producer Consumer Problem

- We can easily solve producer consumer problem using semaphores
- Note that solution must avoid deadlock
 - Improper use of semaphores may cause both consumer and producer to block
 - Example: The consumer holds the mutex (semaphore) and is waiting for someone to signal. Producer could signal but it is waiting for the mutex.
 - Producer and consumer are stuck – waiting for each other
 - Deadlock
- Read details from Section 31.5 OSTEP book

The Dining Philosophers

- Assume there are five “philosophers” sitting around a table
- Between each pair of philosophers is a single fork (and thus, five total)
- The philosophers each have
 - times where they think, and don’t need any forks
 - Times when the eat – They need both forks
 - One on their right and one on their left
- How do they solve it for contention?
- How can they synchronize?

Dining Philosopher Problem



Dining Philosopher problem

```
while (1) {  
    think();  
    get_forks(p);  
    eat();  
    put_forks(p);  
}
```


- How to write `get_fork()` and `put_fork()` such that
 - There is no deadlock
 - No philosopher starves
 - Concurrency is high (as many philosophers can eat at the same time as possible)

Deadlock

Example: Single-Lane Bridge Crossing



Bridge Crossing Example



- Each segment of road can be viewed as a resource
 - Car must own the segment under them
 - Must acquire segment that they are moving into
- **Deadlock:** Two cars in opposite directions meet in middle
- **Starvation** (not deadlock): Eastbound traffic doesn't stop for westbound traffic

Deadlock Definition

- A deadlock is a cycle of waiting among a set of threads, where each thread waits for some other thread in the cycle to take some action
- Deadlocks can happen due to many conditions

Lecture Summary

- Semaphores are a powerful and flexible primitive for writing concurrent programs
- One could view semaphores as a generalization of locks and condition variables
- Deadlock can happen where each process/thread waits for some other thread in the cycle to take some action