

CS310 Operating Systems

Lecture 31: Deadlock Prevention

Ravi Mittal
IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, University of California, Berkeley
 - Book: Modern Operating System, Andrew Tanenbaum and Herbert Bos – Chapter 6
 - Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II
 - Chapter 5.6

Reading

- Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II
- Book: Operating System Concepts, 10th Edition, by Silberschatz, Galvin, and Gagne
- Book: Modern Operating System, Andrew Tanenbaum and Herbert Bos – Chapter 6

Previous Classes

Necessary Conditions for Deadlock

- There are 4 necessary conditions for deadlock to occur
- If you can prevent any one of these conditions, you can eliminate the possibility of deadlock

1. Bounded Resources (or Mutual Exclusion)

- There are a finite number of threads that can simultaneously use a resource. Only one thread at a time can use a resource

2. No preemption

- Once a thread acquires a resource, its ownership cannot be revoked until the thread acts to release it

3. Wait while Holding

- A thread holds one resource while waiting for another
 - Also called **multiple independent requests**


4. Circular Waiting

- A set of waiting threads such that each thread is waiting for a resource held by another


Resource-Allocation Graph (RAG) Examples

Model: Directed Graph


- request edge
 - $T_i \rightarrow R_j$
- assignment edge
 - $R_j \rightarrow T_i$



Simple Resource
Allocation Graph



Allocation Graph
With Deadlock



Allocation Graph
With Cycle, but
No Deadlock

Instead of thread T , we can also represent a process with a circle

Banker's Algorithm

- Suppose we know the **worst case** resource needs of processes in advance
 - A bit like knowing the credit limit on your credit cards. (This is why they call it the Banker's Algorithm)
- Observation: Suppose we just give some process ALL the resources it could need...
 - Then it will execute to completion.
 - After which it will give back the resources.
- Like a bank: If Visa just hands you all the money your credit lines permit, at the end of the month, you'll pay your entire bill

Banker's Algorithm

- So...
 - A process pre-declares its worst-case needs
 - Then it asks for what it “really” needs, a little at a time
 - The algorithm decides when to grant requests
- It delays a request unless:
 - It can find a sequence of processes...
 - such that it could grant their outstanding need...
 - so they would terminate...
 - letting it collect their resources...
 - and in this way it can execute everything to completion!

Today we will study

- Deadlock Prevention Approaches
- Deadlock Recovery Approaches

How Should a System Deal With Deadlock?

Three different approaches:

1. Deadlock avoidance
2. Deadlock prevention
3. Deadlock recovery

Deadlock Prevention

Issues with Deadlock Avoidance

- Deadlock avoidance is very hard
 - It requires information about future resource requests
- Future Resource requests are not easily known
- Then, how to prevent deadlocks?
- Relook at 4 conditions that can cause deadlock

Conditions for Deadlock

- If you can prevent any one of these conditions, you can eliminate the possibility of deadlock

1. Mutual Exclusion (Bounded Resources)

- Only one thread at a time can use a resource. Only one thread at a time can use a resource

2. No preemption

- Once a thread acquires a resource, its ownership cannot be revoked until the thread acts to release it

3. Wait while Holding

- A thread holds one resource while waiting for another
 - Also called multiple independent requests

4. Circular Waiting

- A set of waiting threads such that each thread is waiting for a resource held by another

Attack Mutual Exclusion Condition

Deadlock Prevention (remove Mutual Exclusion)

- Remove “Mutual Exclusion”
 - Infinite resources
 - Example: Virtual Memory

(Virtually) Infinite Resources

Thread A

AllocateOrWait(1 MB)

AllocateOrWait(1 MB)

Free(1 MB)

Free(1 MB)

Thread B

AllocateOrWait(1 MB)

AllocateOrWait(1 MB)

Free(1 MB)

Free(1 MB)

With virtual memory we have **infinite** space so everything will just succeed.

Attack Mutual Exclusion – Example

- Don't assign a resource exclusively to a single process
- How about writing to a printer?
 - Printer is dedicated to a process at a time
- Solution: Use printer spooler
 - Now, several processes can generate output at a time
 - The only process that actually requests the physical printer is the printer daemon
 - Daemons are programmed to print only after the complete file is available

Attack Mutual Exclusion – Example

- However, NOT all devices can be spooled
- Principle:
 - Avoid assigning a resource unless it is absolutely necessary and try to make sure that as few processes as possible may actually claim the processor

Attack Hold and Wait Condition

Attacking the **Hold and Wait** Condition

- Require all processes request resources before starting
 - If everything is available, the process will be allocated whatever it needs and can run to completion
- Problem with the above approach:
 - Many processes don't know how many resources they need until they start running
 - Resources are not used optimally – resources get tied up until process completes (may take hours)
- Another Solution: **back-off and retry**
 - A process requesting a resource to first temporarily release all the resources it currently holds
 - Then it tries to get everything it needs all at once

Request Resource Atomically

Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

Thread B

```
y.Acquire();  
x.Acquire();  
...  
x.Release();  
y.Release();
```

Consider instead:

Thread A

```
Acquire_both(x, y);  
...  
y.Release();  
x.Release();
```

Thread B

```
Acquire_both(y, x);  
...  
x.Release();  
y.Release();
```

Or consider this:

Thread A

```
z.Acquire();  
x.Acquire();  
y.Acquire();  
z.Release();  
...  
y.Release();  
x.Release();
```


Thread B

```
z.Acquire();  
y.Acquire();  
x.Acquire();  
z.Release();  
...  
x.Release();  
y.Release();
```

Attacking – No preemption Condition

Attacking the **No Preemption** Condition

- This is not a viable option
- Consider a process given the printer
 - halfway through its job
 - now forcibly take away printer
 - !!??




Attacking – Circular Wait Condition

Attacking **Circular Wait** Condition

- There are several ways to eliminate circular wait
- Process is entitled to a single resource at any moment
 - If it needs the second one, it has to release the first one
 - ☐ Not possible in many use cases – eg copy a huge file from a tape to a printer
- Provide a global numbering of all resources
- Rule:
 - Processes can request resources whenever they want to
 - All requests must be made in numerical order
 - Example: A process may request first a printer and then a tape drive, but it may not request first a plotter and then a printer (see diagram in the next slide)

Attacking the Circular Wait Condition

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD Rom drive



- Resource allocation graph will never have cycles
- Two processes A and B
- Deadlock can happen when when A requests j and B requests i
 - If $i < j$, B is not allowed to request i, and vice versa
- Many variations to this approach exist

Summary of all deadlock prevention approaches

Condition	Approach
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away
Circular wait	Order resources numerically

Deadlock Recovery

How Should a System Deal With Deadlock?

Three different approaches:

1. Deadlock avoidance
2. Deadlock prevention
3. Deadlock recovery

Deadlock Recovery

- Let deadlock happen, and then figure out how to deal with it

How to Deal with Deadlock?

- Terminate thread, force it to give up resources
 - Dining Philosophers Example: Remove a dining philosopher
 - In AllocateOrWait example, OS kills a process to free up some memory
 - Not always possible—killing a thread holding a lock leaves world inconsistent
- Roll back actions of deadlocked threads
 - Common techniques in databases (transactions)
 - Of course, if you restart in exactly the same way, may enter deadlock again
- Preempt resources without killing off thread
 - Temporarily take resources away from a thread
 - Doesn't always fit with semantics of computation

Lecture Summary

- Resource deadlock can be avoided by keeping track of which states are safe and which are unsafe.
- Resource deadlock can be structurally prevented by building the system in such a way that it can never occur by design
 - Allowing processes to hold one resource at a time
 - Resource deadlock can also be prevented by numbering all the resources and making processes request them in strictly increasing order
- Once deadlock happens, system must recover from it by
 - Killing a thread or process
 - Rolling back actions of deadlocked thread or process
 - Preempt resources without killing thread or process