

# CS310 Operating Systems

## Lecture 7: Process – fork() system call

Ravi Mittal  
IIT Goa

# Acknowledgements !


- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
  - Class presentation: University of California, Berkeley: David Culler, Anthony D. Joseph, John Kubiatawicz, AJ Shankar, George Necula, Alex Aiken, Eric Brewer, Ras Bodik, Ion Stoica, Doug Tygar, and David Wagner
- Operating Systems: Three Easy Pieces, by Remzi and Andrea Arpaci-Dusseau,
  - Chapter 5: Process APIs
  - Programs are taken from this chapter

# Read the following:

- Operating Systems: Principles and Practice (2nd Edition)  
Anderson and Dahlin
  - Volume 1, Kernel and Processes
    - Chapter 4
- Operating Systems: Three Easy Pieces, by Remzi and Andrea  
Arpaci-Dusseau,
  - Chapter 5: Process APIs

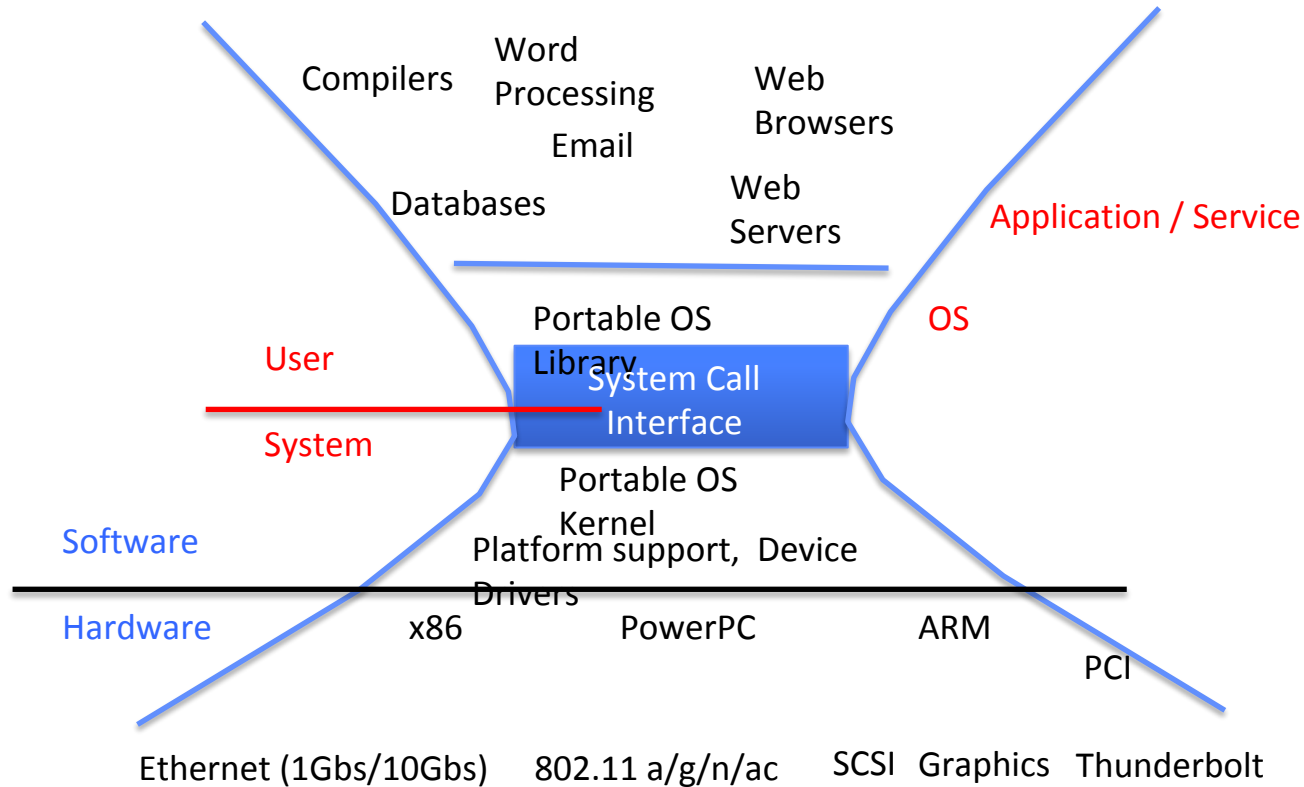
# We will study..

- Fork() system call

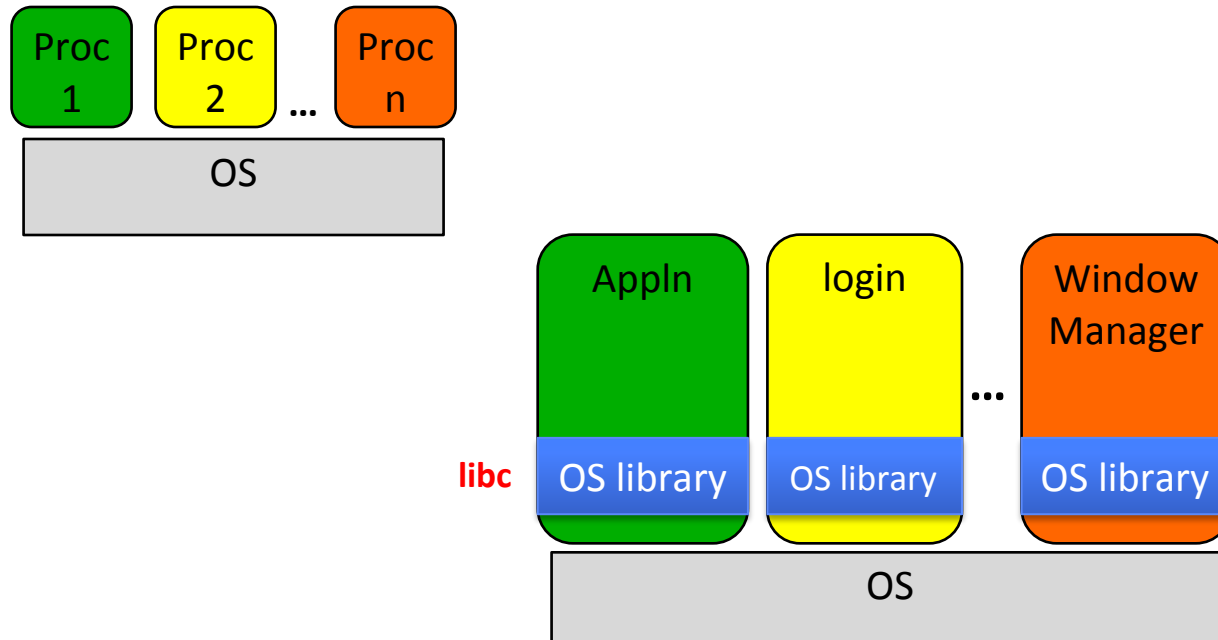


**We have studied so far ...**

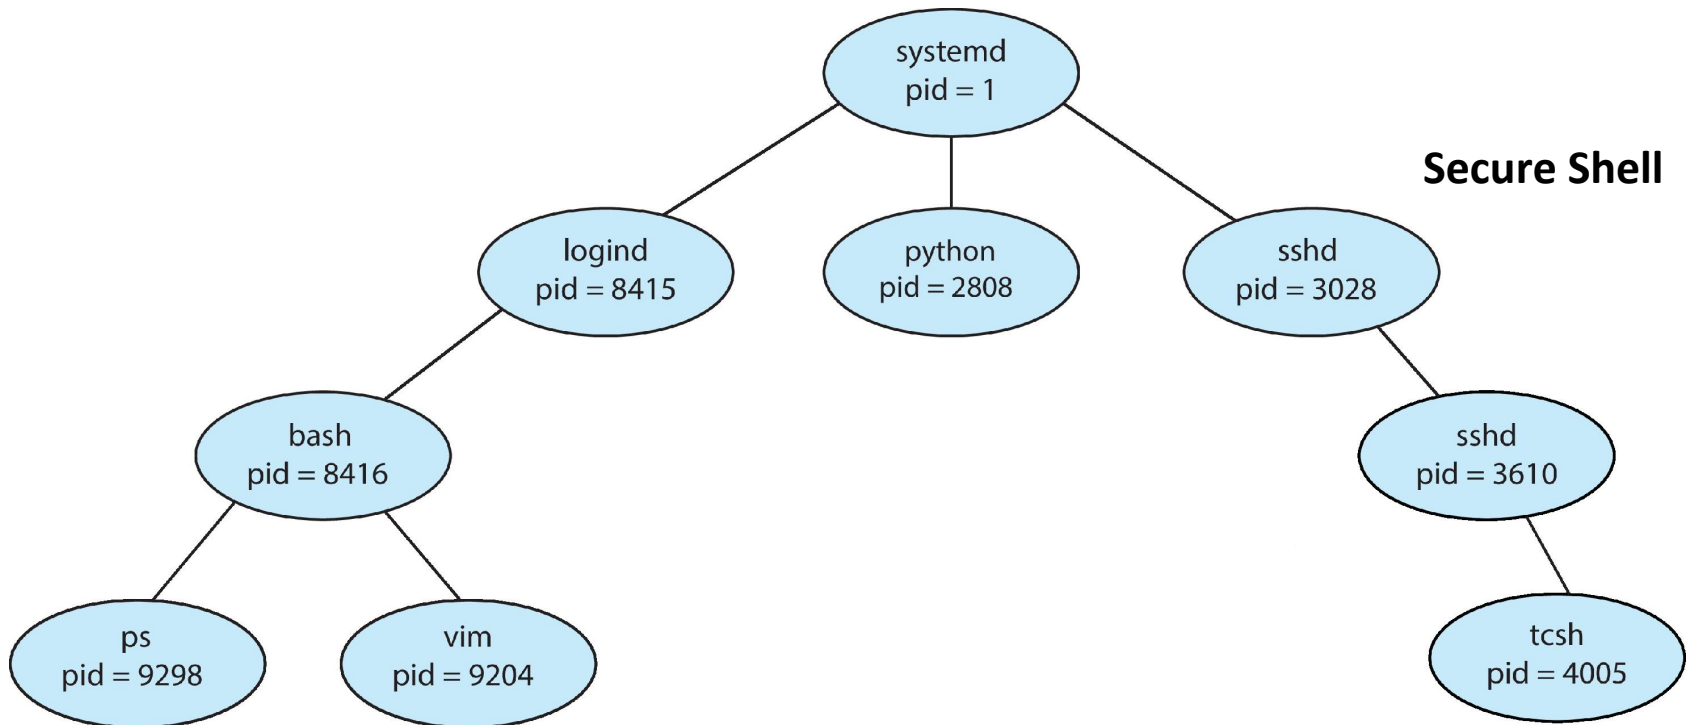
# System Calls (“Syscalls”)



# OS Library Issues Syscalls



# A Tree of Processes in Linux



**ps -e** command gives details of all active processes in the system

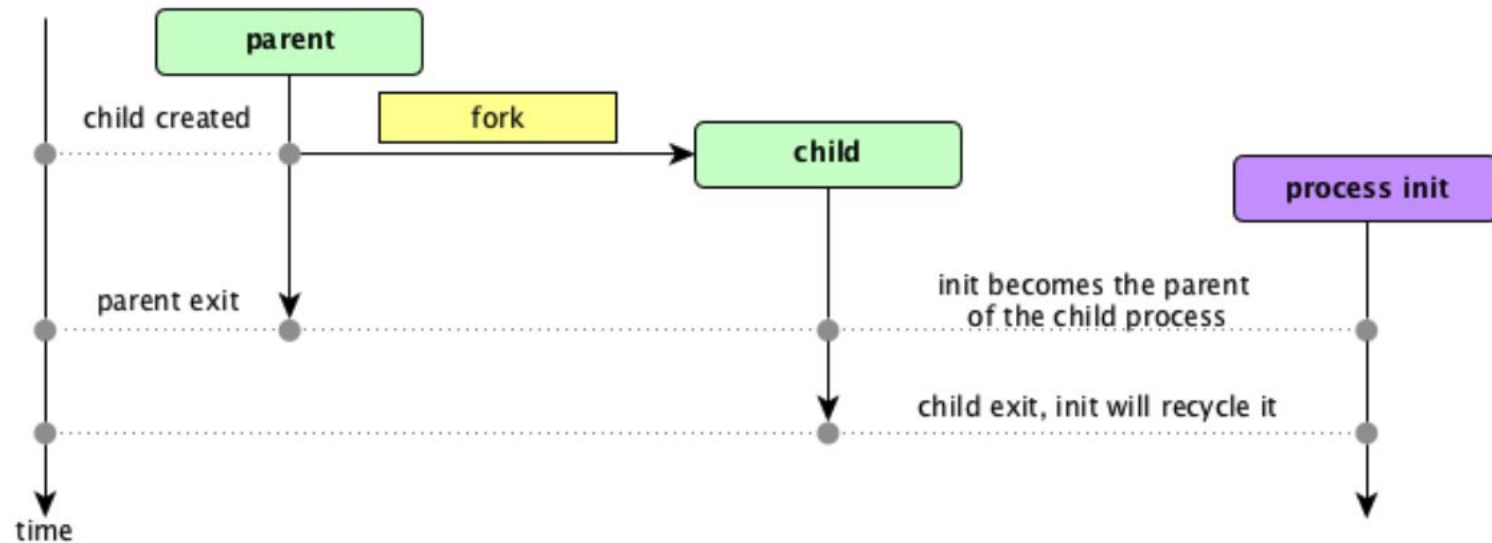


# Process Creation

- During execution, a process may create several new processes
  - Creating Process: Parent Process
  - New Processes: Children Processes
- Each of the children process may create other processes
  - Forming Tree of processes
- The `systemd` process (which always has a pid of 1) serves as the root parent process for all user processes
  - `Systemd` process in Unix is called `init` process

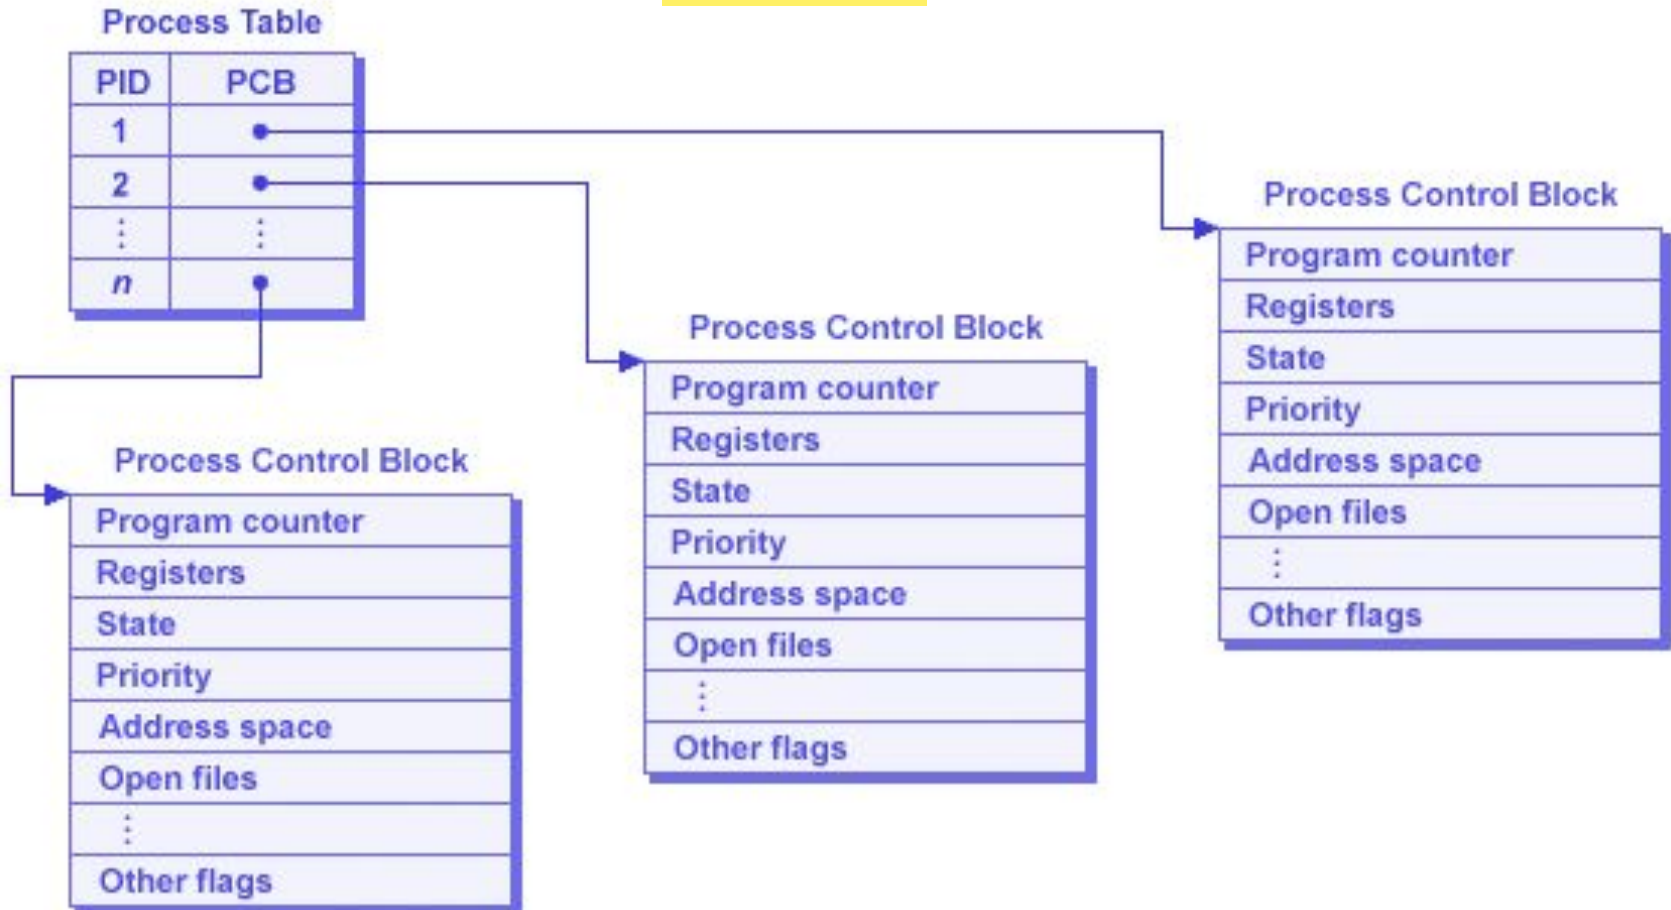
# Process creation and termination – snapshot - 1

In case the parent dies, **init** becomes parent of the child



# Another Data structure – Process Table

Used in Linux



# System Calls

# Process Related System Calls (in Unix)

- `fork()` creates a new child process
  - All processes are created by forking from a parent
  - The *init* process is ancestor of all processes
- `exec()` makes a process execute a given executable
- `exit()` terminates a process
- `wait()` causes a parent to block until child terminates
- There are many variants of the above system calls with different arguments

## System Calls – **fork()**

# Creating a Process (Linux)

- All processes are created using `fork` system call
  - Creates an exact copy of the current process
  - Both processes continue in **parallel from the statement that follows the fork call**
  - The only difference is in the **return value**
    - Parent: Child process ID (“pid” , non-zero)
    - Child: 0
      - Child can get parent ID via `getppid()`
    - Failure: -1

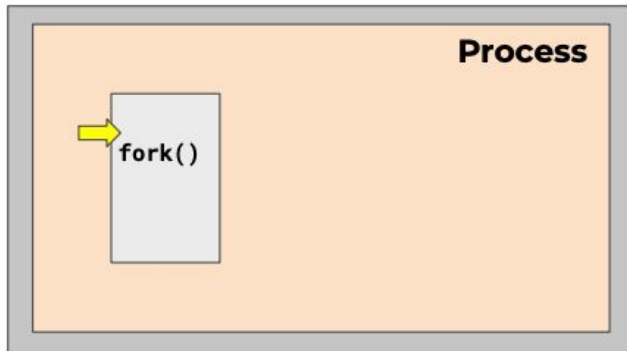
# Actions during a fork

- A new process is created by making a copy of parent's memory image
  - In a separate address space
- The new process is added to the process list and scheduled
- Parent and child start execution just after fork (with different return values)
- Parent and child execute and modify the memory data independently
- Fork copies variables and registers from the parent to the child



# Creating a process


- Just one process – Parent process




# Creating a process

- Just one process – Parent process
- Initially there is one process – `init` with `id = 1`

**id = 1**




**id = 3042**



# Fork() is fun..

- The fork() is one of the those system calls, which is called once, but returns twice



```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0){
        //fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        //child new process
        printf("hello I am child process with (pid = %d) \n", (int) getpid());
    } else {
        //parent
        printf("hello I am parent of %d (pid: %d) \n", rc, (int) getpid());
    }
    return 0;
}
```

```
(base) Ravis-MacBook-Pro-2:Cprograms ravimittal$  
./a.out
```

```
hello world (pid:22305)
```

```
hello I am parent of 22306 (pid: 22305)
```

```
hello I am child process with (pid = 22306)
```

# Notes on Fork()

- Note that the child isn't an exact copy
  - It has its own copy of the address space
  - Its own registers, PC, etc
  - Value it returns is different from its parent
- The output of p1.c is not deterministic
  - After child process is created – there are two active processes
  - On a single CPU, it depends on scheduler which one gets scheduled first. For example, in another run we might get:

```
Ravis-MacBook-Pro-2:Cprograms ravimittal$ ./a.out
hello world (pid:22394)
hello I am child process with (pid = 22395)
hello I am parent of 22395 (pid: 22394)
```

# Class Summary

- `fork()` and `exec()` combination is a powerful way to create and manipulate processes
- In the next class we will study `exec()` system call

**Next Class**

**System Calls – exit, wait and exec**



## System Calls – **exit()** and **wait()**

# Process termination


- Multiple ways for a process to get destroyed
  - Process issues and *exit()* call – Voluntary
  - The parent process issues a *kill()* call - Involuntary
  - Process receives a *terminate signal* - Involuntary
    - When it did something illegal !
- On death
  - Reclaim all of process's memory regions
  - Make process un-runnable
  - Put the process in the *zombie state*
  - However, do not remove its process descriptor from the list of processes

# Zombie State

- Why keep process descriptor around?
  - Parent may be waiting for child to terminate
    - via the *wait()* system call
  - Parent needs to get the exit code of the child
    - this information is stored in the descriptor
  - If descriptor was destroyed immediately, this information could not be gotten
  - After getting this information, the process descriptor can be removed
    - no more remnants of the process

# Waiting for children to die with **wait()**

- The parent can wait for the child to die by executing the **wait** system call
- It is quite useful for a parent to wait for a child process to finish what it has been doing
  - on success, **returns** the process ID of the terminated child; on error, -1 is **returned**.



# wait and waitpid syscalls

- A terminated process's information is **collected** via a call of a wait operation by its parent
- Operations:
  - **wait**
    - blocks the calling process until a child process terminates, returning the child's pid
      - If caller has no children, wait immediately returns -1 (error)
    - the termination status (return value) of the child may be obtained via the argument
  - **waitpid**
    - permits a caller to wait for a particular child, identified by its pid

# wait() system call

p2.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int
7  main(int argc, char *argv[])
8  {
9      printf("hello world (pid:%d)\n", (int) getpid());
10     int rc = fork();
11     if (rc < 0) {
12         // fork failed; exit
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) {
16         // child (new process)
17         printf("hello, I am child (pid:%d)\n", (int) getpid());
18         sleep(1);
19     } else {
20         // parent goes down this path (original process)
21         int wc = wait(NULL);
22         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
23                rc, wc, (int) getpid());
24     }
25     return 0;
26 }
```

# wait() system call

Parent process waits for the child to finish

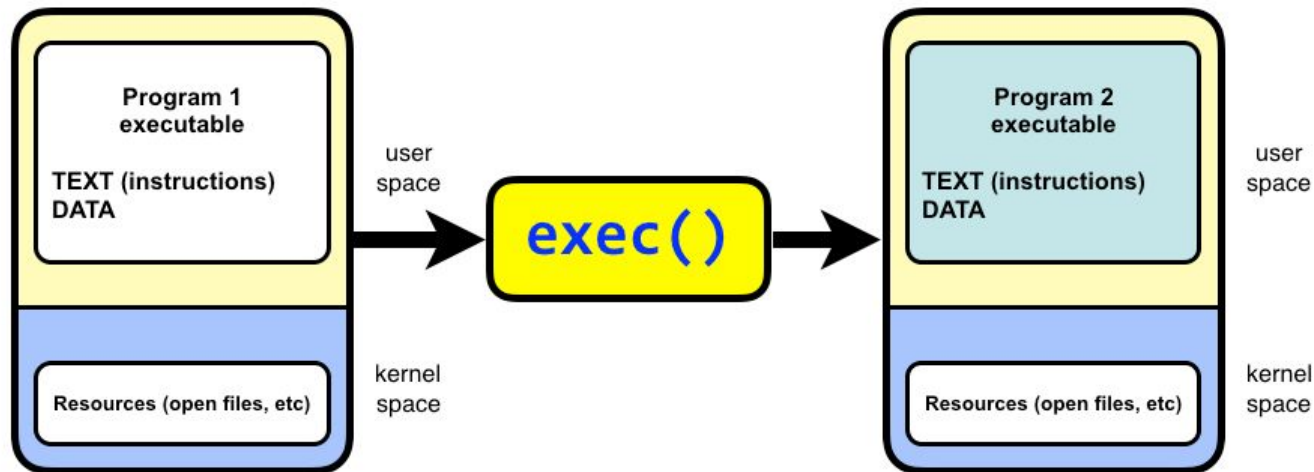
```
(base) Ravis-MacBook-Pro-2:cp ravimittal$ ./p2  
hello world (pid:13858)  
hello, I am child (pid:13859)  
hello, I am parent of 13859 (wc:13859) (pid:13858)
```

# Executing a new program

- After **fork**, parent and child are running same code
  - Not too useful!
- A common use of fork is to launch a new executable program
- A process can run **exec()** to load another executable to its memory image
  - So, a child can run a different program from parent
- The **exec** system call replaces the current process image with a new image
  - *If **exec** succeeds, it never returns*
- **exec** requires you to specify the file you program to run



# exec() system call



- The `exec` family of system calls replaces the program executed by a process
- When a process calls `exec`, all code (text) and data in the process is lost and replaced with the executable of the new program
- All open file descriptors remains open after calling `exec`
  - unless explicitly set to close-on-exec

# `exec()` system call

- We consider system call: `execvp()`
- `execvp` system call requires two arguments
  - The first argument is a character string that contains the name of a file to be executed
  - The second argument is a pointer to an array of character strings. More precisely, its type is **`char **`**

# execvp() system call

- When **execvp()** is executed, the program file given by the first argument will be loaded into the caller's address space
  - duplicate the actions of the shell in searching for an executable file if the specified file name does not contain a slash (/) character
- The second argument will be provided to the program. It's argument vector. The first argument points to the **filename** associated with the file being executed. The array of pointers must be terminated by a NULL pointer.

```
#include <unistd.h>
```

```
int execvp(const char *file, char *const argv[]);
```

# More on Process APIs

- `kill()` system call is used to send **signals** to a process
  - To pause, die, and other imperatives
- In shells, certain keystroke combinations are configured to deliver a specific signal to the currently running process
  - **control-c** sends a SIGINT (interrupt) to the process (normally terminating it)
  - **control-z** sends a SIGTSTP (stop) signal thus pausing the process in mid-execution
    - Can use **bg** command to resume