

CS310 Operating Systems

Lecture 34 : File System - 2

Ravi Mittal

IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - Book: Modern Operating Systems by Andrew Tanenbaum and Herbert Bos,
 - Chapter 4
 - Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love
 - Book: Computer Systems, A programming Perspective, Bryant and O'Hallaron
 - Class presentation: University of California, Berkeley, CS162
 - Book: Linux – The Textbook, by Sarwar, Koretsky

Read the following:

- Book: Modern Operating Systems, by Andrew Tanenbaum and Herbert Bos
 - Chapter 4
- Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love
- Book: Operating Systems: Three Easy Pieces, Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau


Today's Class

- File System – Introduction (previous class)
- Directories
- Hard and Soft (Symbolic) links - Concepts

Previous Classes

File System

- Files are logical unit of information – Created by Processes
- A Disk may contain thousands of files – usually independent of each other
- Processes can read existing files and create new files
- Files remain in existence not affected by process creation and termination
- File must disappear only when the owner explicitly deletes it
- File are managed by Operating system
- File System deals with how files are
 - Structured
 - Named
 - Accessed
 - Used
 - Protected
 - Implemented, and
 - Managed



File Naming

Extension	Meaning
file.bak	Backup file
file.c	C source program
file.gif	Compuserve Graphical Interchange Format image
file.hlp	Help file
file.html	World Wide Web HyperText Markup Language document
file.jpg	Still picture encoded with the JPEG standard
file.mp3	Music encoded in MPEG layer 3 audio format
file.mpg	Movie encoded with the MPEG standard
file.o	Object file (compiler output, not yet linked)
file.pdf	Portable Document Format file
file.ps	PostScript file
file.tex	Input for the TEX formatting program
file.txt	General text file
file.zip	Compressed archive

File Structure – Three categories

- **Byte sequences**

- Maximum flexibility-can put anything in
- Unix and Windows use this approach
- Byte stream

In use


- **Fixed length records (card images in the old days)**

- Read/Write operations are done on one record
 - Example: Based on **Punch cards (80 columns)** – record 80 characters
 - Not in use
- VMS OS provides highly structured file – with records

Obsolete



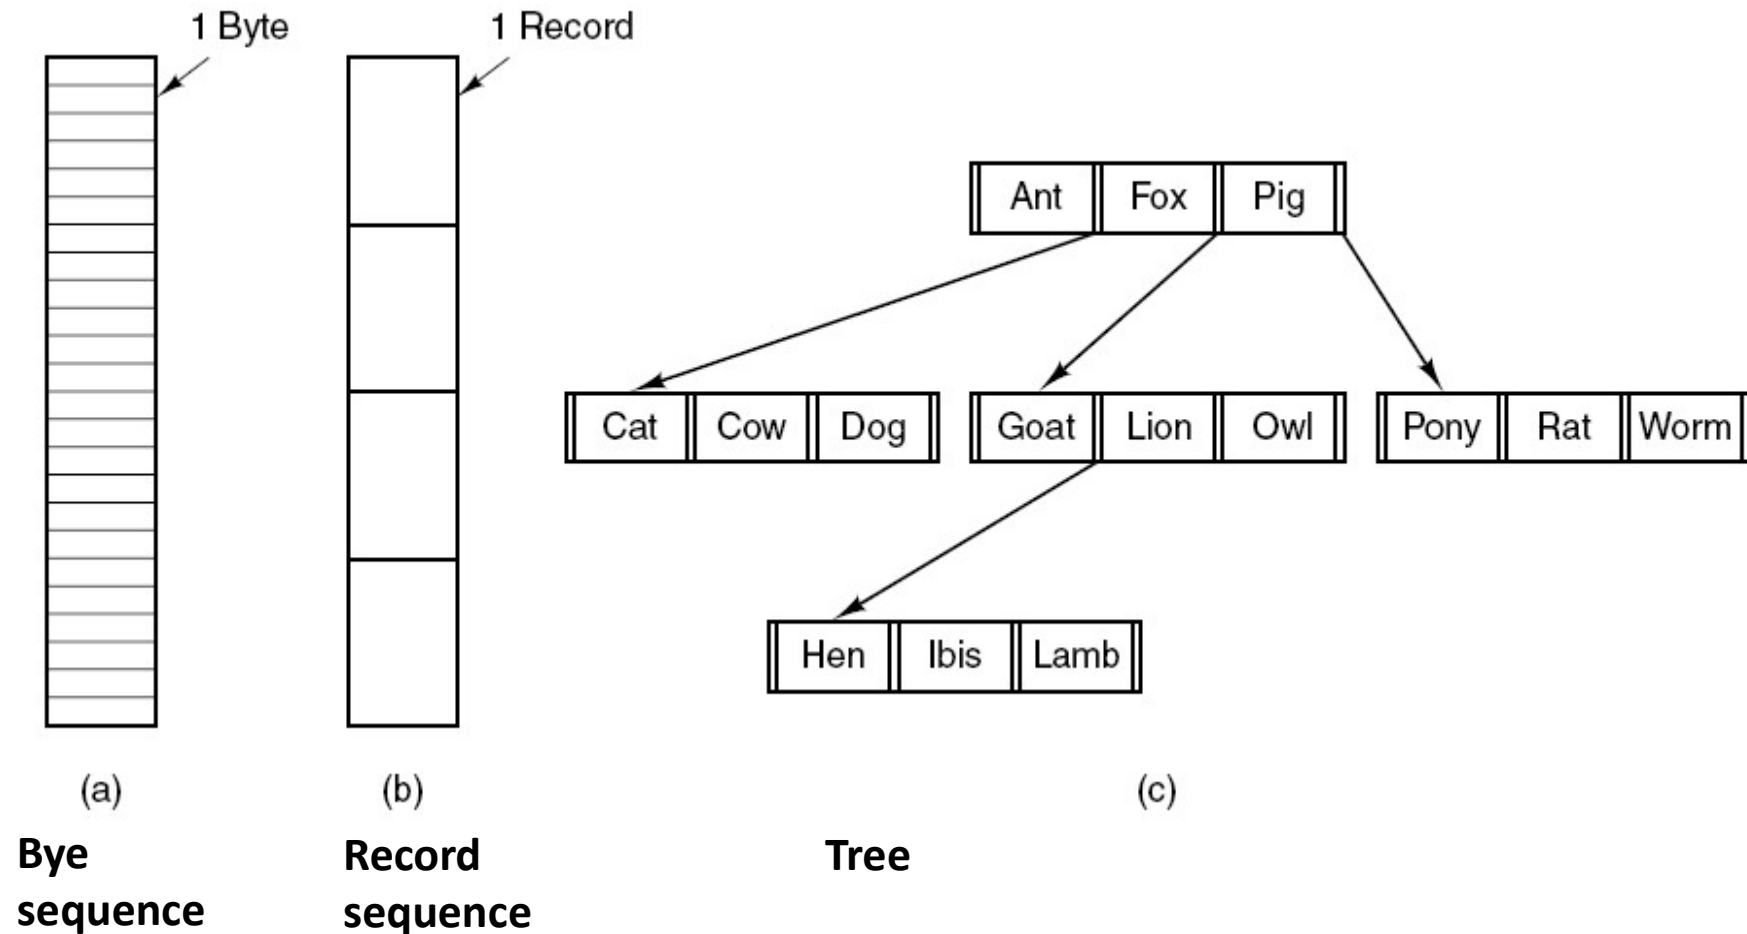
- **Tree of records- uses key field to find records in the tree**

- Tree is sorted on the key field
- Used in some large mainframes

Obsolete



File Structure – Three Types



File Types

- Regular Files (most common)
 - Text
 - Binary
- **Directories**
 - A file that contains the names and locations of other files
- Character Special and Block Special Files
 - Terminals (character special) and disks (block special)
- FIFO (named pipe)
 - A file type used for inter-process communication
- Socket
 - A file type used for network communication between processes

Regular Files

- What we mostly call as files – are *regular files*
- *Regular files* contain bytes of data organized into linear array called a *byte stream*
 - Linux has no further organization or formatting specified for a file
- Any byte within a file can be read or written to
- Read or write operation start at a specific byte
 - Conceptual location
 - *File position* or *file offset*
 - Essential metadata
- At file opening – *file position is zero*

Regular Files (Linux)

- File reading or writing is done byte by byte
 - The **file position** increases
- File position may be set manually



- **File position** starts at zero
- Writing a byte to the middle of a file overwrites over the previous byte
- It is not possible to expand a file by writing into middle of it

Regular Files (Linux)

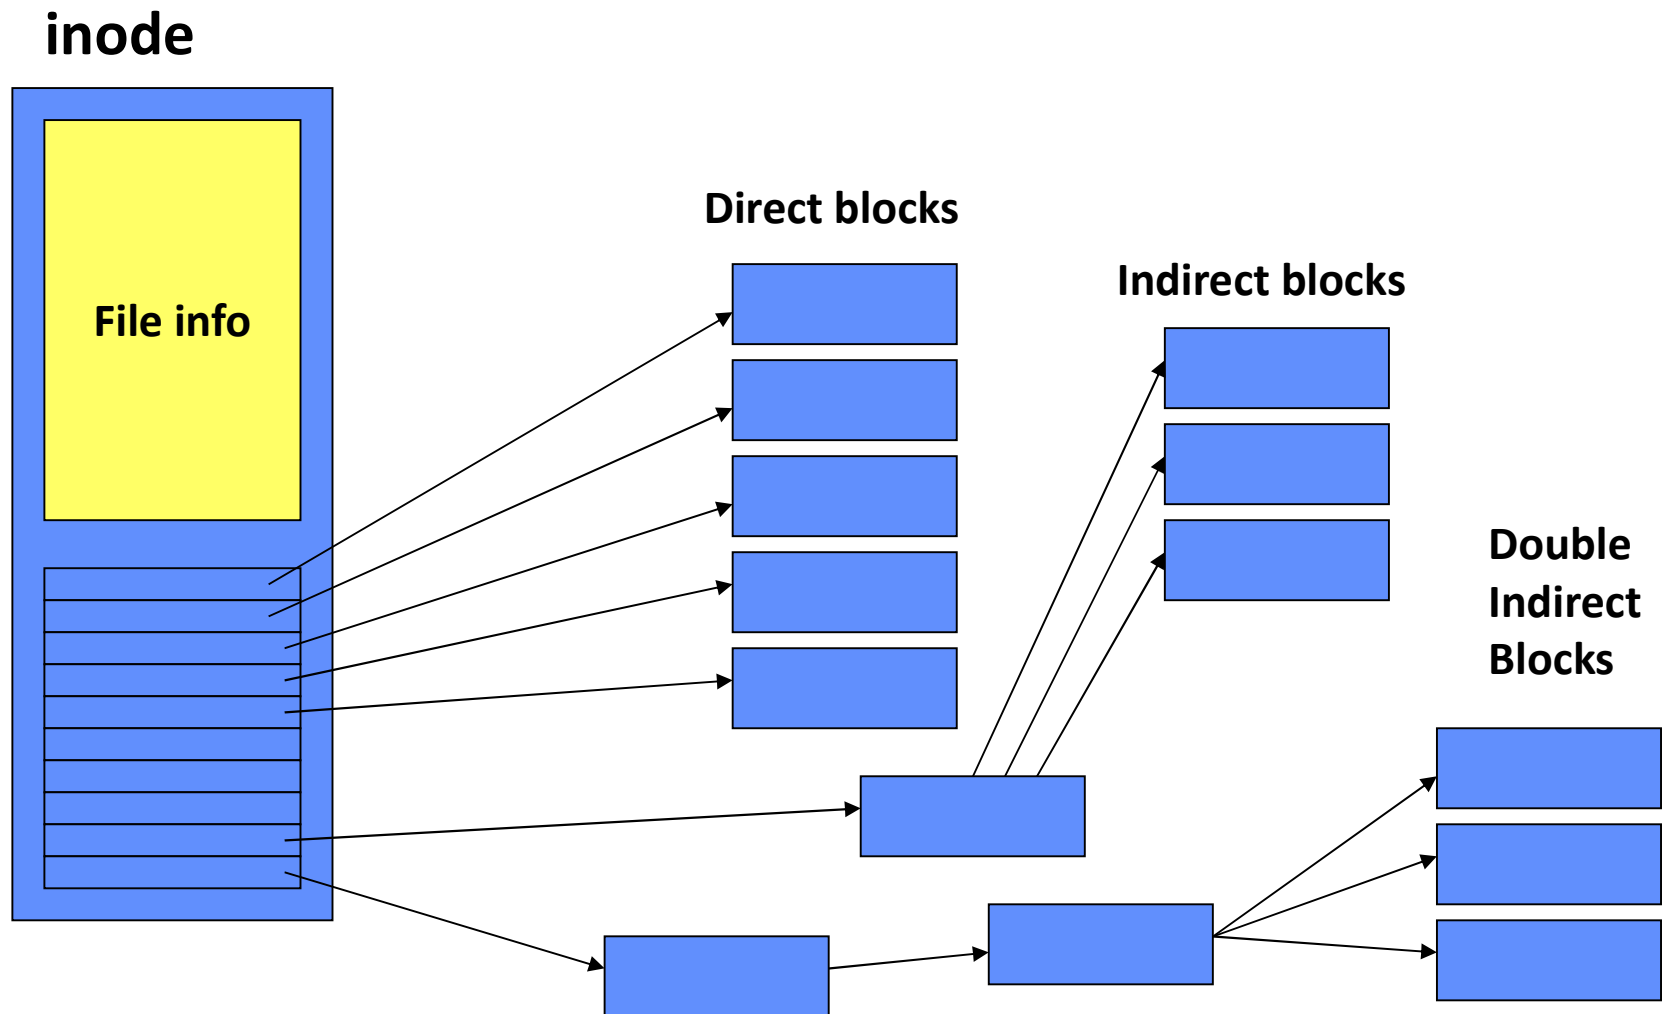


- A single file may be opened more than once
 - By the same process or different processes
 - Each open instance of a file gives a unique **file descriptor**
- Processes can share their **file descriptors**
- No restriction on concurrent file accesses by multiple processes
- If concurrent reading and writing is happening
 - Result may be unpredictable
 - So user processes must coordinate
- Files are associated with file names but they are referenced by **inode** – information node
 - **inode** is assigned an integer value unique to the file system – **inode** number or **i-number**

Regular Files (Linux) - *inode*

- When a file is created
 - System allocates a unique *inode* to the file
 - It creates an entry in the directory in which the file is created
 - Directory entry: an ordered pair (**inode number, filename**)
- The *inode* number for a file is used to access its attributes, including its contents on the disk for reading or writing (changing) them
- *An Inode* stores metadata associated with a file
 - Modification time stamp / access time
 - Owner (UID, GID)
 - Type (file, directory, device, pipe, ...)
 - Length
 - Pointers to data blocks that store file's contents
- Example: You created a file **Chapter 3** in your current working directory
 - System allocates *inode* number **53472**
 - Directory entry: (**52473, Chapter3**)

Inode diagram



Directories

Directories - Introduction

- Directories are handled as normal files
 - However, marked in *inode* as directory
- Accessing a file via its *inode* number is cumbersome and unsecure (security hole)
- Files are always opened from user space by a name
 - Not by *inode* number
- *Directories* are used to provide the names with which to access files
- A directory provides mapping of human-readable names to *inode* numbers
- A name and *inode* pair is called a *link*
- Directory is implemented as a simple table or a hash
 - Implemented and managed by Kernel
- The kernel directly uses this mapping to perform name-to-*inode* resolutions

Directories - Introduction

- A user application requests a file name to be opened
 - Kernel opens the directory containing the filename and searches for the given name
 - From the filename, the kernel obtains the *inode* number
 - *Inode* number points to *inode*
 - The *inode* contains metadata associated with the file and location of file's data on the disk
- Initially, there is only one directory on the disk, the *root directory*
 - Donated by */*.
- There are many directories in a file system
- Links inside of directories can point to the *inodes* of other directories
- Directories nest inside of other directories – hierarchy of directories

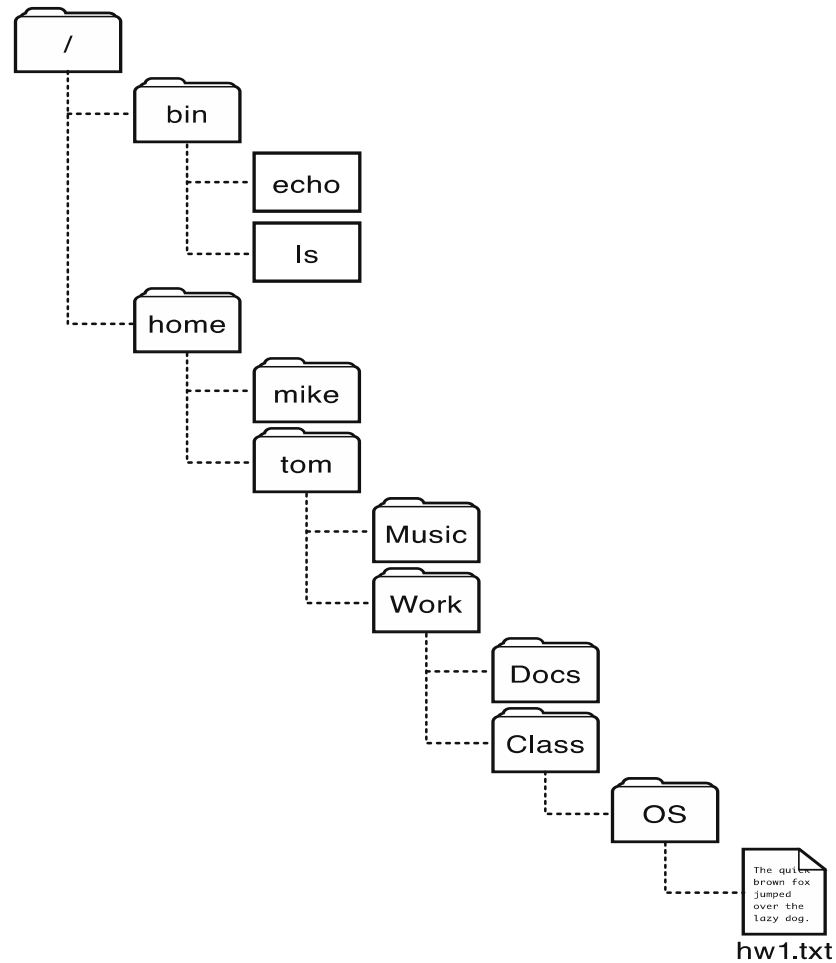
Directories - Introduction

- Users identify files via pathnames example
(*/home/blackbeard/concorde.png*)
- To open a given path, kernel walks each *directory entry* (*dentry*) in the pathname to find the *inode* of the next entry
- *dentry cache in Linux*
 - to store the results of directory resolutions
- Path
 - String that identifies a file or directory
 - Example: /home/tom/Work/Class/OS/hw1.txt or /home/tom
- Absolute Path
 - if it starts with “/”, the root directory
 - Example: /bin/ls
- Relative Path
 - relative to a process’s *current working directory*
 - Example: Work/Class/OS if the current directory is /home/tom
 - Equivalent to absolute path: /home/tom/Work/Class/OS

Directories - Introduction

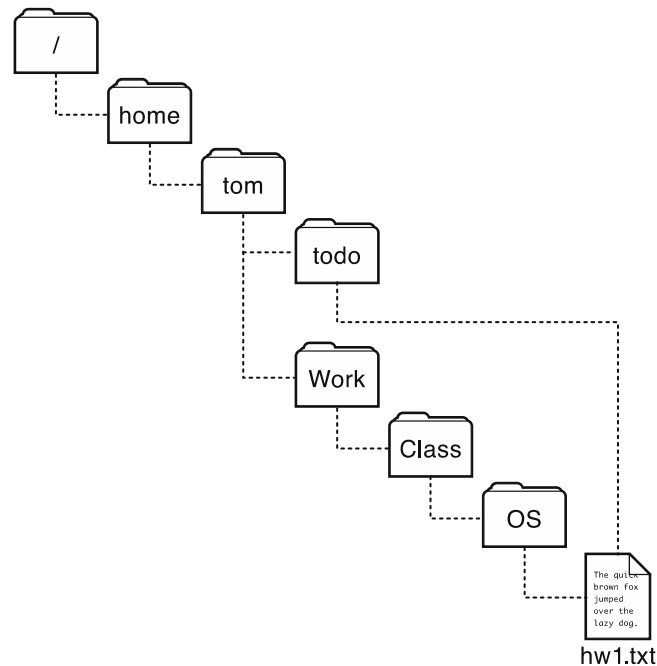
- Paths are resolved from left to right and the respective name is looked up in the directory
- As long as the current name is not the last in the path, it has to be a directory
 - Else it terminates with an error

Hierarchical organization of files using directories




Directory is (not always) a tree

- Sometimes it is convenient for a file to appear in two directories



Directory Structure

- Consider a current directory containing three files: Chapter 1, Chapter 2, and Chapter 3




Logical Structure

Contents of current directory

Inode #	File
16517779	.
16517772	..
16518415	Chapter1
16518416	Chapter2
16518417	Chapter3

**Contents of the disk block
containing current directory**

Directory Entry, inode, and file contents



ls -il command

```
[(base) Ravis-MacBook-Pro-2:sample ravimittal$ ls -il
total 29504
4391804511 -rw-r--r--@ 1 ravimittal  staff    2830690 Jun 18 16:53 Chapter 1.pptx
4391921390 -rw-r--r--@ 1 ravimittal  staff    1016998 Jun 22 18:04 Chapter 2.pptx
4393535828 -rw-r--r--@ 1 ravimittal  staff   11250437 Aug  8 15:34 Chapter 3.pptx
(base) Ravis-MacBook-Pro-2:sample ravimittal$
```

Hard and Soft (Symbolic) Links

File Sharing Via links

- Attributes of a Linux file are stored in its *inode* on disk
- When a file is opened, its *inode* is copied into the main memory
 - Speedy access to its contents
- How the use of an *inode* that results in a mechanism that allows you to access a file by using *multiple pathnames*?
 - To allow access to some files and directories through various other directories
- A link is a way to establish connection between
 - Files to be shared with directory enteries
- A file has N links, we mean that the file has N directory entries somewhere in the file system hierarchy
- Linux supports two types of links:
 - Hard links
 - Soft/symbolic links

Hard Links

- A hard link is a *pointer* to the *inode* of a file
- A hard link gives additional file name
- All hard links point to the same *inode*
- As long as the link counter is not 0, the file survives
- Hard links can be created for files in the same logical file system

In command

Syntax:

ln [options] existing-file new-file

ln [options] existing-file-list directory

Purpose: First syntax: Create a link to **existing-file** and name it **new-file**


Second syntax: Create links to the ordinary files in **existing-file-list** in **directory**; links have the same names as the original file

Commonly used options/features:

- f** Force creation of link; don't prompt if **new-file** already exists
- n** Don't create the link if **new-file** already exists
- s** Create a symbolic link to **existing-file** and name it **new-file**

```
[(base) Ravis-MacBook-Pro-2:sample ravimittal$ ls ]
chapter1.pptx  chapter2.pptx  chapter3.pptx
[(base) Ravis-MacBook-Pro-2:sample ravimittal$ ln chapter3.pptx chapter3.hard ]
[(base) Ravis-MacBook-Pro-2:sample ravimittal$ ls -il ]
total 31496
4391804511 -rw-r--r--@ 1 ravimittal  staff  2830690 Jun 18 16:53 chapter1.pptx
4393535828 -rw-r--r--@ 1 ravimittal  staff  11250437 Aug  8 15:34 chapter2.pptx
4391921390 -rw-r--r--@ 3 ravimittal  staff   1016998 Jun 22 18:04 chapter3.hard
4391921390 -rw-r--r--@ 3 ravimittal  staff   1016998 Jun 22 18:04 chapter3.pptx
(base) Ravis-MacBook-Pro-2:sample ravimittal$ █
```

Structure of current directory



(a)

Contents of current directory


Inode #	File
1076	.
2083	..
13059	Chapter1
17488	Chapter2
52473	Chapter3

(b)

Establishing a hard link

Chapter3 Chapter3.hard

Structure of current directory



(a)

Contents of current directory

Inode #	File
1076	.
2083	..
13059	Chapter1
17488	Chapter2
52473	Chapter3
52473	Chapter3.hard

(b)

Hard Link across directories


In memo6.hard memos/memo6

Contents of current directory

Inode #	File
1076	.
2083	..
13059	Chapter1
17488	Chapter2
52473	Chapter3
83476	memo6.hard

Contents of ~/memos

Inode #	File
1076	.
2083	..
83468	memo1
...	...
83476	memo6
...	...



Soft link or Symbolic link

- Many systems allows to use multiple names to refer to the same file
- A symbolic link is a directory mappings from a file name to another file name
- To open a file is opened via a symbolic link
 - File system first translates the name in the symbolic link to the target name and then uses the target name to open the file
- Example:
 - Create /a/b
 - Create a symbolic link from /c/d/ to /a/b
 - Unlinking /a/b, the file is no longer accessible and open("/c/d") will fail
- A MacOS file alias is similar to a symbolic link

Lecture Summary

- When a file is created
 - System allocates a unique *inode* to the file
- *Inode* stores metadata of the file on disk
- Directory stores file name and *inode* number
- Hard and Soft (Symbolic) links provide a way to access a file using multiple pathnames