

CS310 Operating Systems

Lecture 35 : File System - 3

Ravi Mittal

IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - Book: Modern Operating Systems by Andrew Tanenbaum and Herbert Bos,
 - Chapter 4
 - Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love
 - Book: Linux – The Textbook, by Sarwar, Koretsky
 - Book: Advanced Programming in the Unix Environment, by W Richard Stevens and Stephen Rago
 - Lecture Notes: IC221 Systems Programming,
 - <https://www.usna.edu/Users/cs/wcbrown/courses/IC221/classes/L09/Class.html>

Read the following:

- Book: Modern Operating Systems, by Andrew Tanenbaum and Herbert Bos
 - Chapter 4
- Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love
- Book: Linux – The Textbook, by Sarwar, Koretsky
- Book: Advanced Programming in the Unix Environment, by W Richard Stevens and Stephen Rago


Today's Class

- File System – Introduction (previous class)
- File System –Data Structures
- File System - Name Space
- Virtual File System
- Basic File System APIs

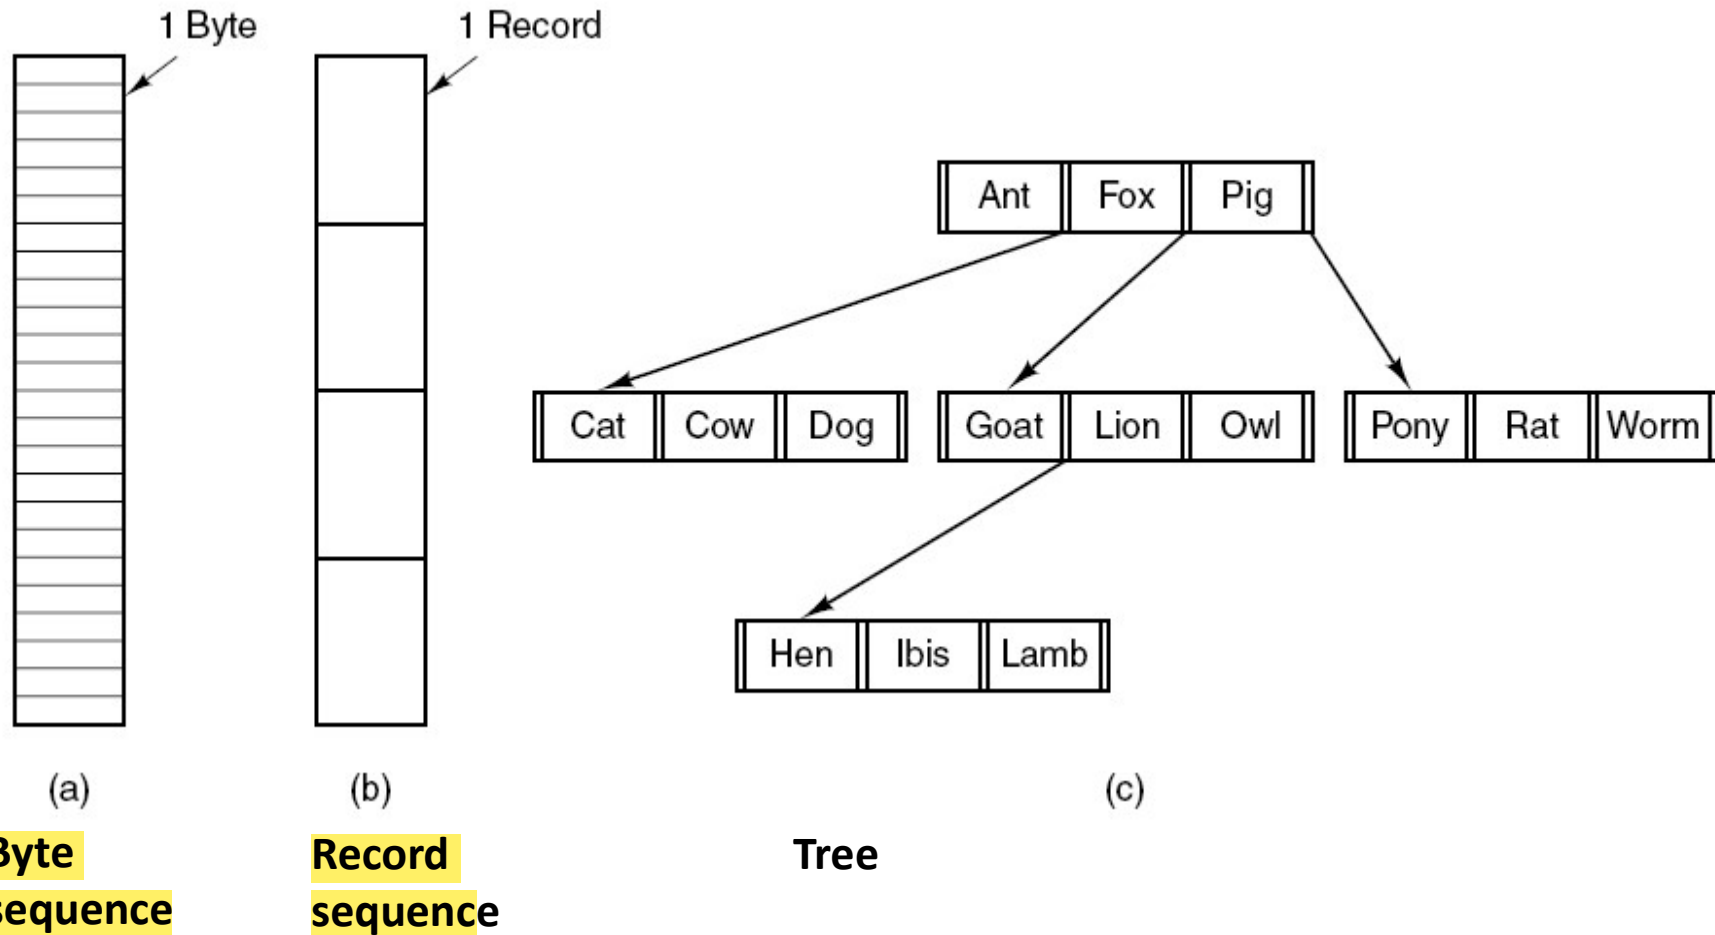
Previous Classes

File System

- Files are logical unit of information – Created by Processes
- A Disk may contain thousands of files – usually independent of each other
- Processes can read existing files and create new files
- Files remain in existence not affected by process creation and termination
- File must disappear only when the owner explicitly deletes it
- File are managed by Operating system
- File System deals with how files are
 - Structured
 - Named
 - Accessed
 - Used
 - Protected
 - Implemented, and
 - Managed



File Structure – Three Types



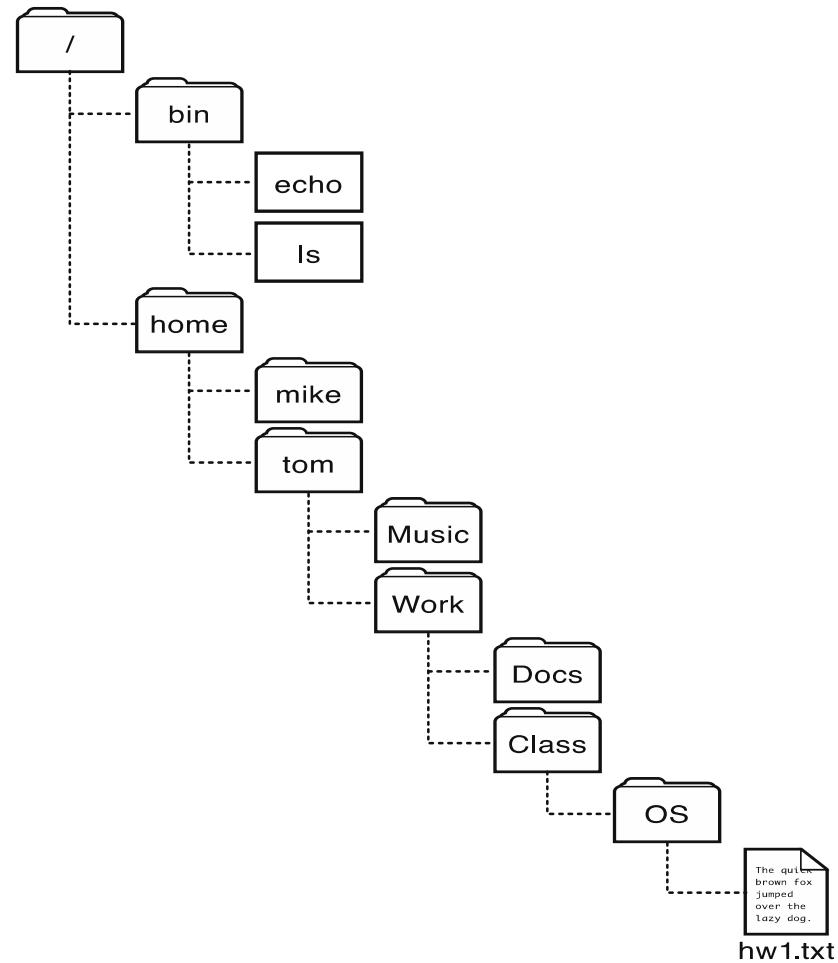
Regular Files (Linux)

- File reading or writing is done byte by byte
 - The **file position** increases
- File position may be set manually




- **File position** starts at zero
- Writing a byte to the middle of a file overwrites over the previous byte
- It is not possible to expand a file by writing into middle of it

Hierarchical organization of files using directories



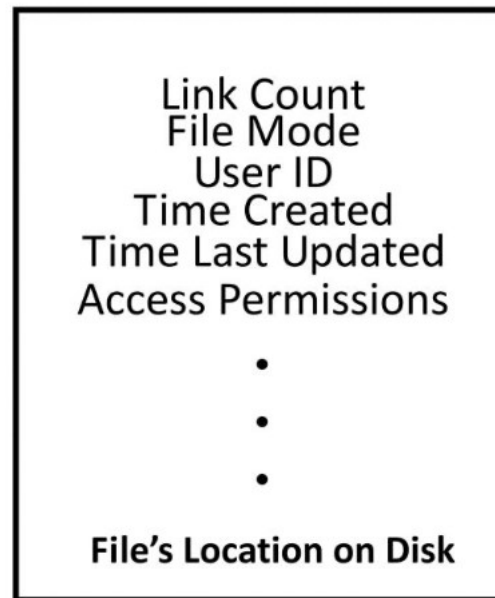
Directory is (not always) a tree

- Sometimes it is convenient for a file to appear in two directories

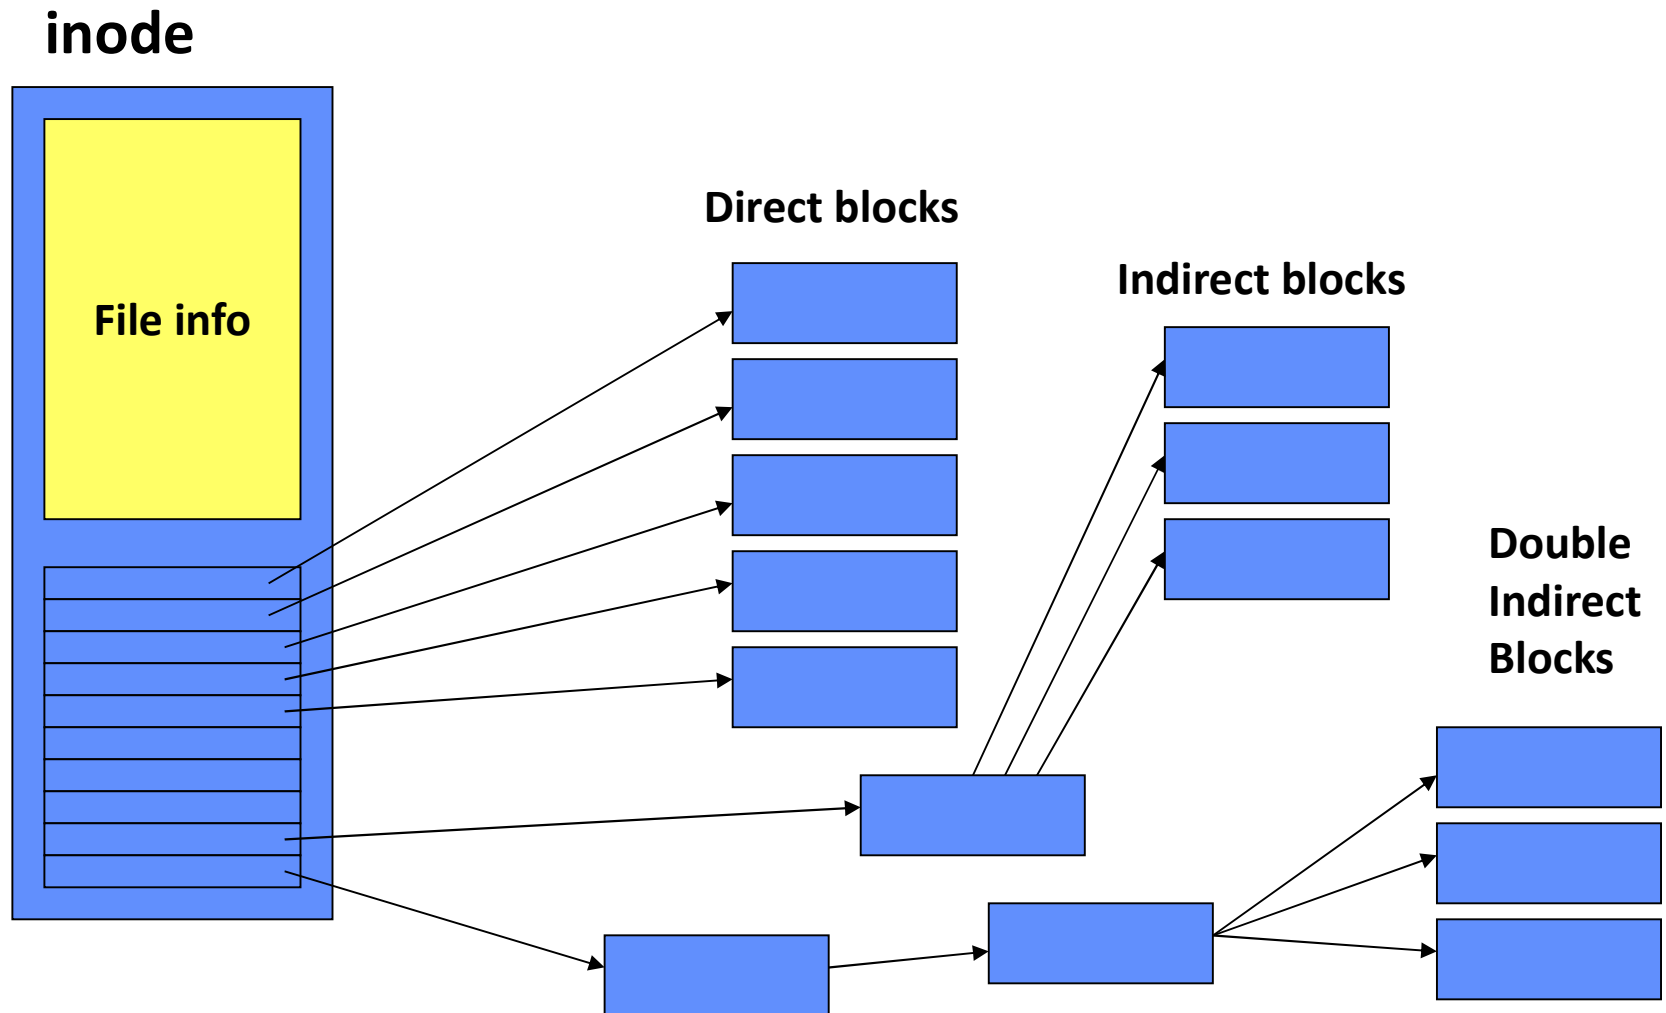


File System Data Structures (Unix)


Contents of *inode*



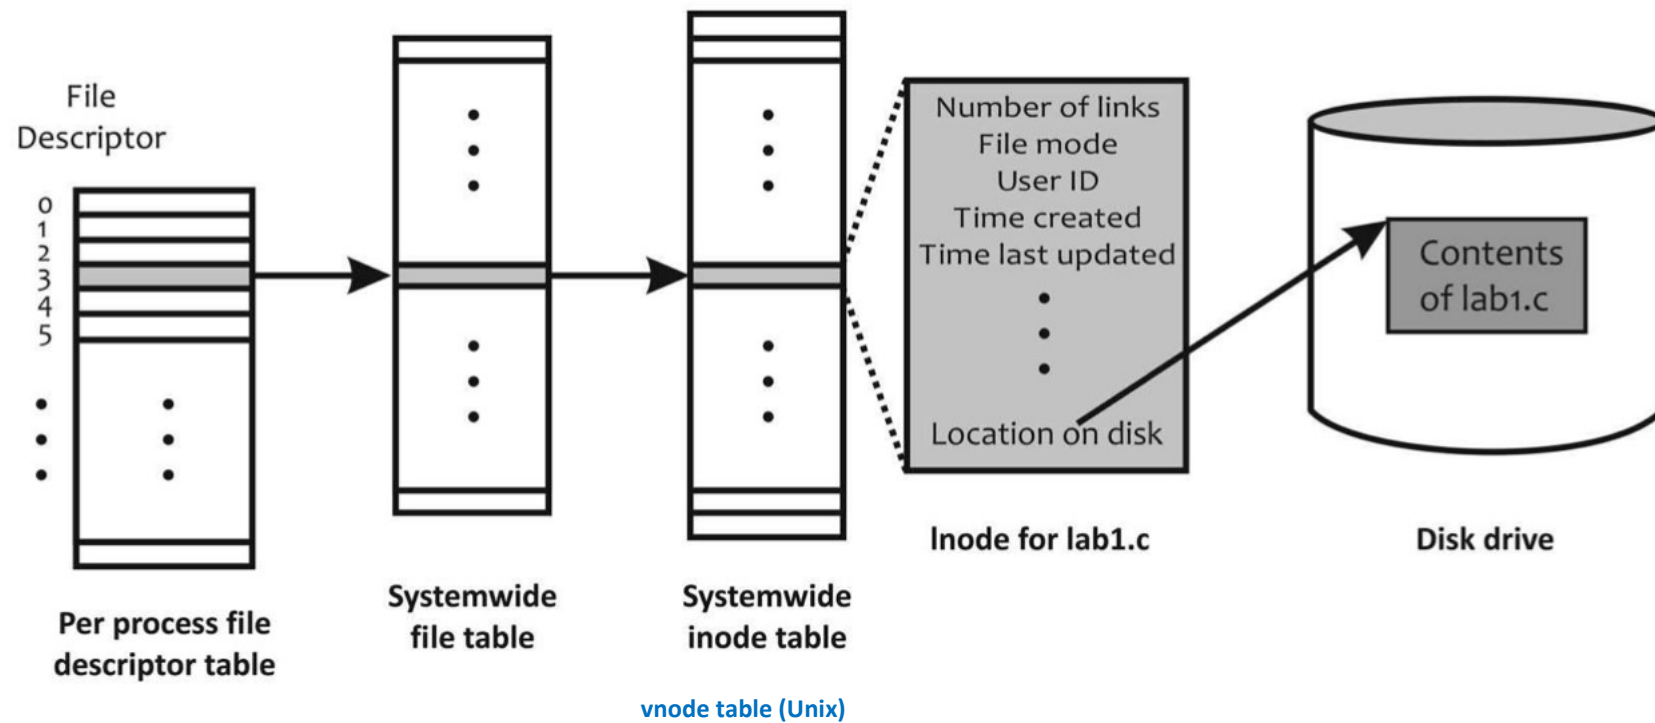
Inode diagram



Directory Entry, inode, and file contents

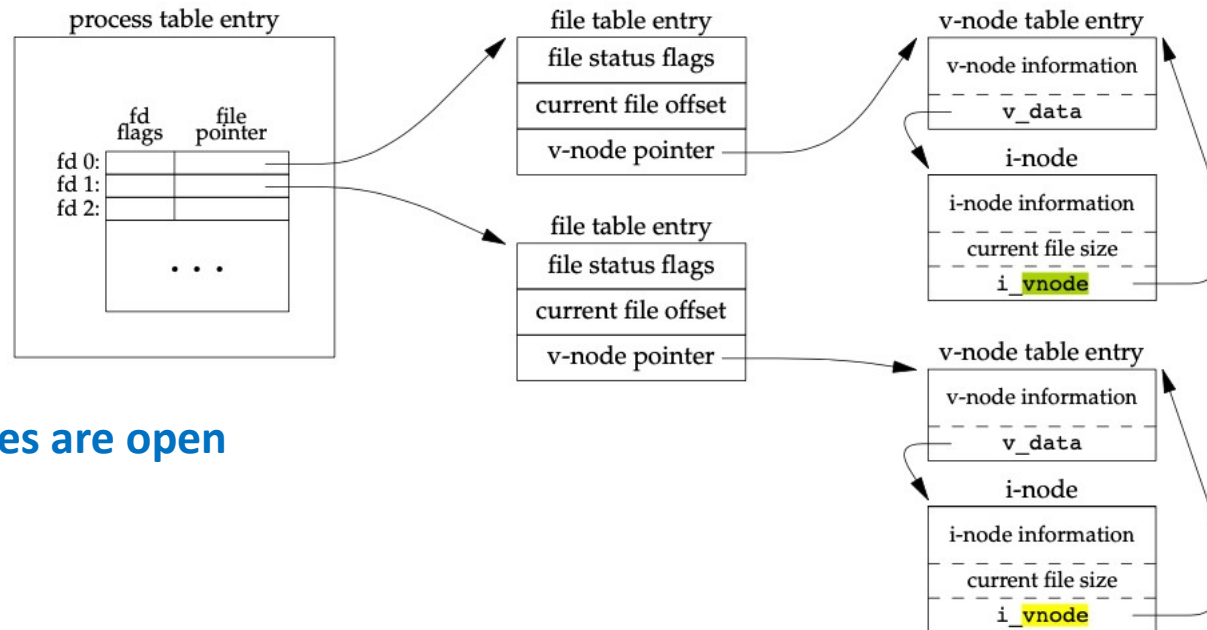


File descriptor, file table, inode, disk



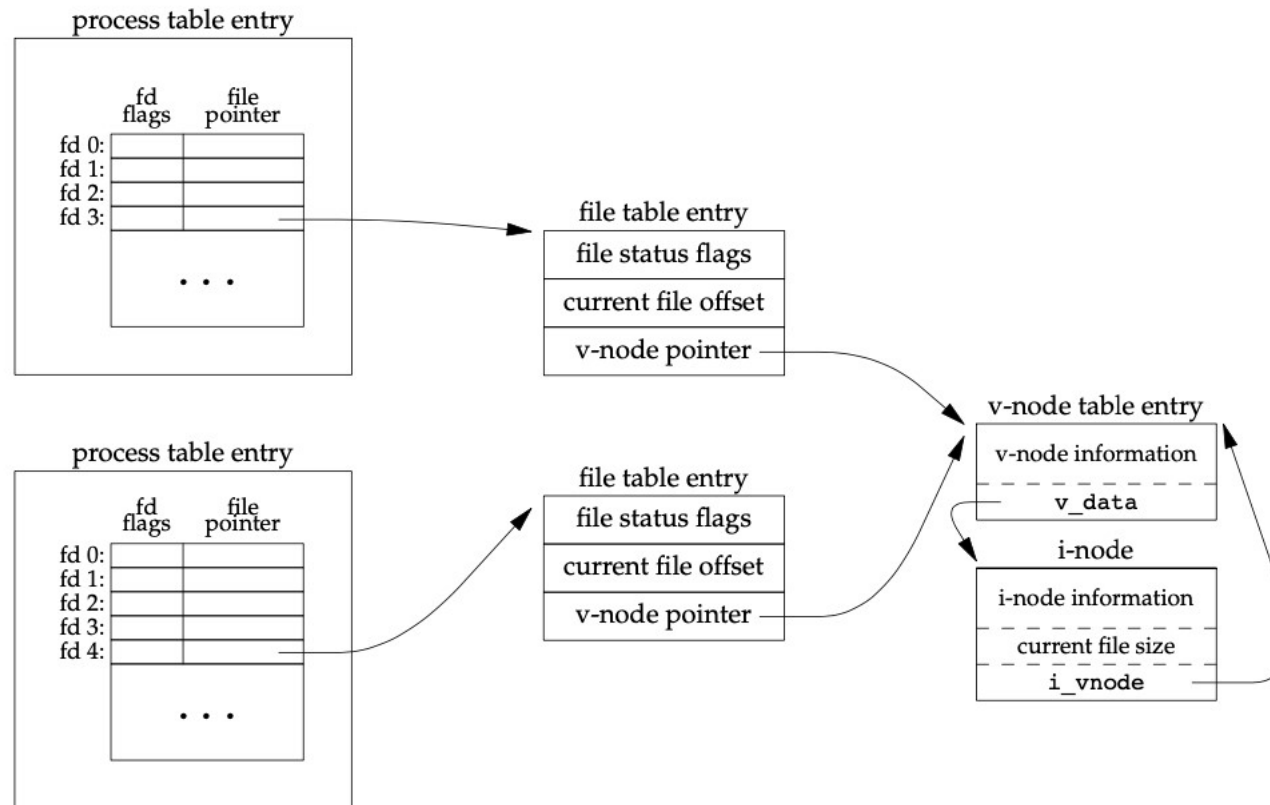
Kernel Data Structures for open files (Unix)

For a single process



Two files are open

Two processes using the same opened file



Big Picture

User Space

Process 1421
int fd 4

Process 1215
int fd 4

- Suppose stdin of process 1421 has been redirected to come from foo.txt.
- Suppose file dex. 4 of process 1215 is set to read from foo.txt.

Kernel Space

Process Table

pid 1215
file descriptor table

0
1
2
3
4
...

pid 1421
file descriptor table

0
1
2
3
4
...


Open file table

status	r
offset	10
vnode ptr	→
ref count	1
...	...
status	r
offset	2
vnode ptr	→
ref count	1
...	...
status	w
offset	0
vnode ptr	→
ref count	1
...	...

Vnode table

type	reg. file
fun ptrs	...
inode	→
ref count	2
...	...
type	reg. file
fun ptrs	...
inode	→
ref count	1
...	...

Physical Drive



File descriptor

- Kernel and the user process refer to each open connection by a number (type int) called a *file descriptor*
- The kernel uses *file descriptors* to perform file operations
 - Example: file read
- The kernel uses a *file descriptor* to index the per-process file descriptor table
 - Obtains a pointer to the systemwide file table
- **File descriptor Table (per process):**
 - It resides in Process Table entry (PCB)
 - PCB contains all *file descriptors* currently in use by a process
 - Each *file descriptor* points to a **file table entry**

System Open File Table

- The system needs to keep track of each **connection** to a file
- Since the same file may have many open **connections**, this connection is truly distinct from the concept of a file
- Distinct processes may even share a **connection**
 - eg. two processes wanted to write to the same log file, each adding new data to the end of the file)
 - The **connection** information cannot be kept inside the process table
- Kernel keeps a data structure called the *system open-file table* which has an entry for each **connection**
- Entry contains
 - **Connection Status**: read or write
 - Current offset in the file
 - A pointer to **vnode** entry for the file
 - **Vnode** is data structure representing the file
 - Linux: vnode → generic inode

Vnode table

- Each open file has a **vnode** structure
- **Vnode** contains
 - Information about type of the file
 - Pointers to functions that operate of file
 - Copy of the *inode* for the file
 - Physical information about the file
 - This information (about *inode*) is read from the disk when the file is opened
 - Reference Count: gives the number of hard links
 - A file can't be removed from the file system unless all of those references are removed

Physical Device: inodes etc

- The *inode* contains the owner of the file, the size of the file, pointers to where the actual data blocks for the file are located on disk
- A file may be broken up into many data blocks, which may be widely distributed across the physical drive
- The *inode* for a file contains the locations of each of the data blocks comprising the file

Special Files

File Types

- Regular Files (most common)
 - Text
 - Binary
- Directories
 - A file that contains the names and locations of other files
- Character Special and Block Special Files
 - Terminals (character special) and disks (block special)
- FIFO (named pipe)
 - A file type used for inter-process communication
- Socket
 - A file type used for network communication between processes

Special Files

- *Special files* are kernel objects that are represented as files
- Linux supports four types of files
 - block device files,
 - character device files,
 - named pipes, and
 - Unix domain sockets
- Special files are a way to let certain abstractions fit into the filesystem
 - Everything-is-a-file paradigm
- Linux provides a system call to create a special file

Device Files

- I/O devices are treated as special files called *device files*
- Hence, device access (in Unix/Linux systems) is performed via *device files*
 - Act and look like normal files residing on the file system
- The same system calls used to interact with regular files on disk can be used to directly interact with I/O devices
- *Device files* may be opened, read from, and written to
- This allows users to access and manipulate devices on the system
- Each type of device has its own special device file
- Two types
 - Character Devices
 - Block Devices

Character Device - Files

- Character devices are accessed via a special file called a *character device file* (node)
- A character device is accessed as a linear queue of bytes
- The device driver places bytes onto the queue, one by one
- User space reads the bytes in the order that they were placed on the queue
- Character devices are accessed via *character device files*
- Character devices provide access to data only as a stream, generally of characters (bytes)
- Example: *Keyboard, mice, printers* etc

Block Device - Files

- A block device, in contrast, is accessed as an array of bytes
- The device driver maps the bytes over a seekable device
- User space is free to access any valid bytes in the array, in any order
 - Read byte 12, byte 7, byte 15 etc
- Block devices are generally storage devices
 - Hard disks, CD-ROM drives, Flash Memory
- They are accessed via *block device files*
- Generally mounted as a filesystem

File System Name Space

Volume

- A volume is a collection of physical storage resources that form a logical storage device
 - Each instance of a file system manages files and directories for a volume
- A volume is an abstraction that corresponds to a logical disk

File System and namespaces

- Some operating systems separate different disks and drives into separate **namespaces**
 - a file on a floppy disk might be accessible via the pathname *A:\plank.jpg*
 - the hard drive is located at *C:*
 - DOS and Windows, which break the file namespace up into drive letters
 - This breaks the namespace up among device and partition boundaries
- Linux (Unix) provides a **global and unified namespace** of files and directories
 - In Unix, that same file on a floppy might be accessible via the pathname */media/floppy/plank.jpg*
 - Alongside files from other media


Mounting and Unmounting

- A *filesystem* is a collection of files and directories in a formal and valid hierarchy
- Filesystems may be individually added to and removed from the *global namespace* of files and directories
 - These operations are called *mounting* and *unmounting*
- Each filesystem is mounted to a specific location in the namespace, known as a *mount point*
- The root directory of the (added) filesystem is then accessible at this mount point
- For example, a CD might be mounted at */media/cdrom*, making the root of the filesystem on the CD accessible at */media/cdrom*

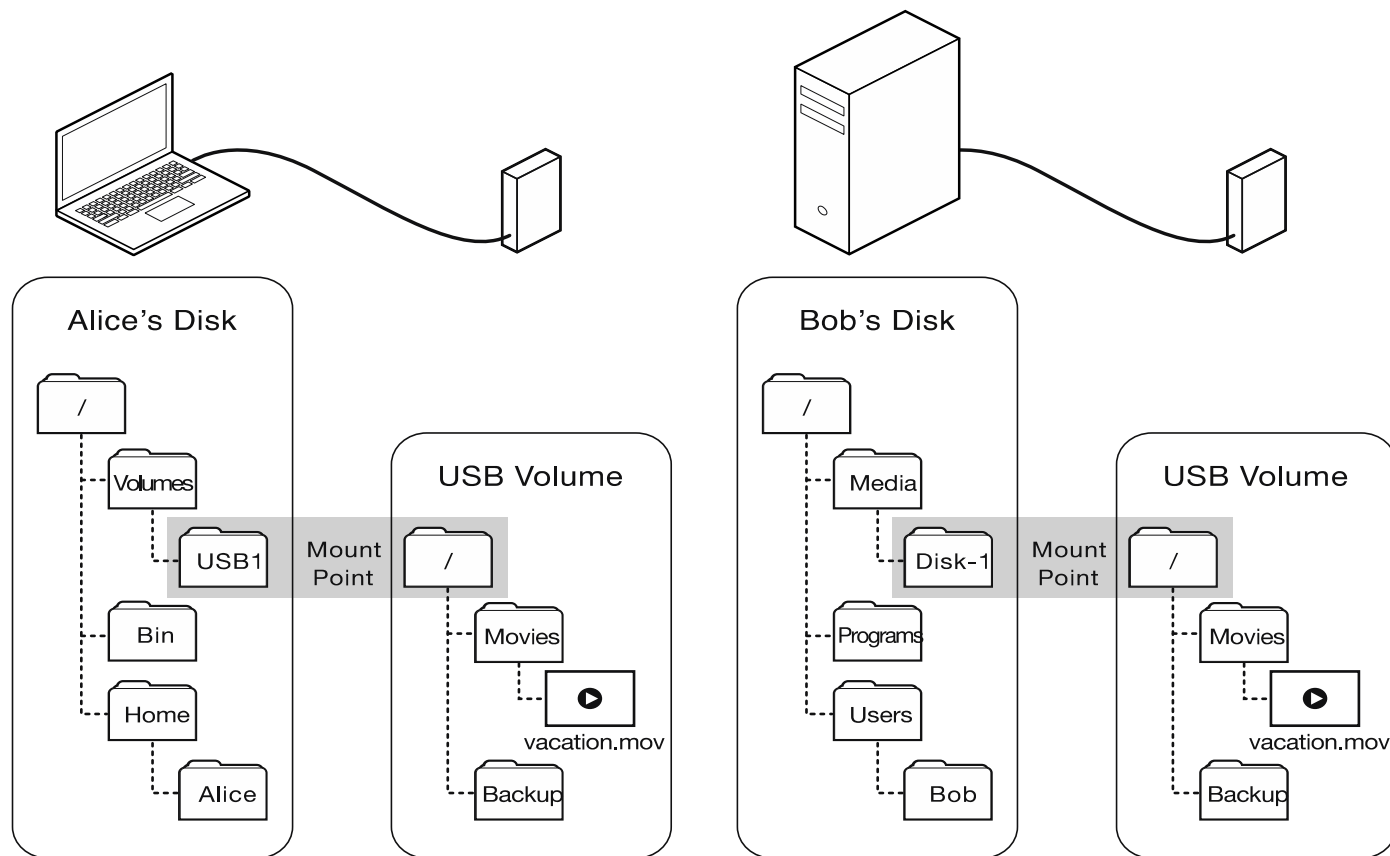
Mount (repeat)

- Mount allows multiple file systems on multiple volumes to form a single logical hierarchy
- Mounting a volume on an existing file system creates a mapping from some path in the existing file system to the root directory of the mounted volume's file system
- Mounting enables all mounted filesystems to appear as entries in a single tree

USB Volume



Mapping from some path in existing file system to the root directory of the mounted file system




Alice can access the vacation.mov movie using the path `/Volumes/usb1/Movies/vacation.mov`

Bob can access the movie using the path `/media/disk-1/Movies/vacation.mov`

Virtual File System (Linux/Unix)

Virtual File System (Unix/Linux)

- The Virtual File System (VFS) implements a generic file system interface between the actual file system implementation (in kernel) and accessing applications
- Applications can access different file systems on different media via a homogeneous set of UNIX system calls




Virtual File System

- VFS is a glue that enables syscalls such as `open()`, `read()`, and `write()` to work regardless of the filesystem or underlying physical medium
- System calls work *between* these different filesystems and media
 - In older operating systems, such as DOS, this would never have worked
- Generic Interface
 - Kernel implements abstraction layer around low-level interfaces
 - Linux to support different filesystems, even if they differ in supported features or behavior
 - FAT, Windows NTFS, Linux and Unix file systems

Virtual File System

- To the user, each file system looks the same
- Example: `write (fd, &buf, len);`
 - Write `len` bytes in file with descriptor `f` from buffer `buf`
 - It is translated into a system call `sys_write()` which determines the actual file writing method for the filesystem that holds file corresponding to `fd`
 - The generic write system call is then invokes the method that is part of the filesystem implementation



Basic File System APIs

File System API – file creation

- Creating files
 - `create()` creates
 - a new file with some metadata, and
 - a name for that file in a directory
 - `link()` creates a hard link
 - a new name for the same underlying file
 - `unlink()` removes a name for a file from its directory
 - If a file has multiple names or links, `unlink()` only removes the specified name, leaving the file accessible via other names
 - If the specified name is the last (or only) link to a file, then `unlink()` also deletes the underlying file and frees its resources
 - `mkdir()` and `rmdir()` create and delete directories

File System API – Open and Close

- `Open()`
 - To get a **file descriptor** it can use to refer to the opened file
 - **Path parsing and permission** checking can be done just when a file is opened
 - Not repeated on each read or write
 - The OS creates a data structure that stores information about the file: file id, read/write/exec, pointer to the processes current position in the file
 - The **file descriptor** can thus be thought of as a reference to the operating system's per-open-file data structure that the operating system will use for managing the process's access to the file
- `Close()`
 - Releases the open file record in the operating system

File System API – File Access

- While a file is open, an application can access the file's data in two ways
 - read()
 - Write()
 - Call to read and write starts from files current position
- Seek()
 - Supports random access within a file
 - Seek() call changes a process's **current position** for a specified open file
- Other APIs
 - munmap()
 - fsync()

Lecture Summary

- In file system (Unix / Linux) the following data structures are used
 - File descriptor
 - File Table entry (Process Control Block)
 - Open File Table
 - Vnode or generic inode table
 - Inode
- Devices are accessed with special files: Character special and block special files
- File Name Space
- Mounting and Unmounting
- Virtual File System (Linux / Unix)
- File APIs