

CS310 Operating Systems

Lecture 36 : File System - 4

Ravi Mittal

IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - Book: Modern Operating Systems by Andrew Tanenbaum and Herbert Bos,
 - Chapter 4
 - Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love
 - Book: Computer Systems, A programming Perspective, Bryant and O'Hallaron
 - Class presentation: University of California, Berkeley, CS162

Read the following:

- Book: Modern Operating Systems, by Andrew Tanenbaum and Herbert Bos
 - Chapter 4
- Book: Linux System Programming: talking directly to the kernel and C library, by Robert Love

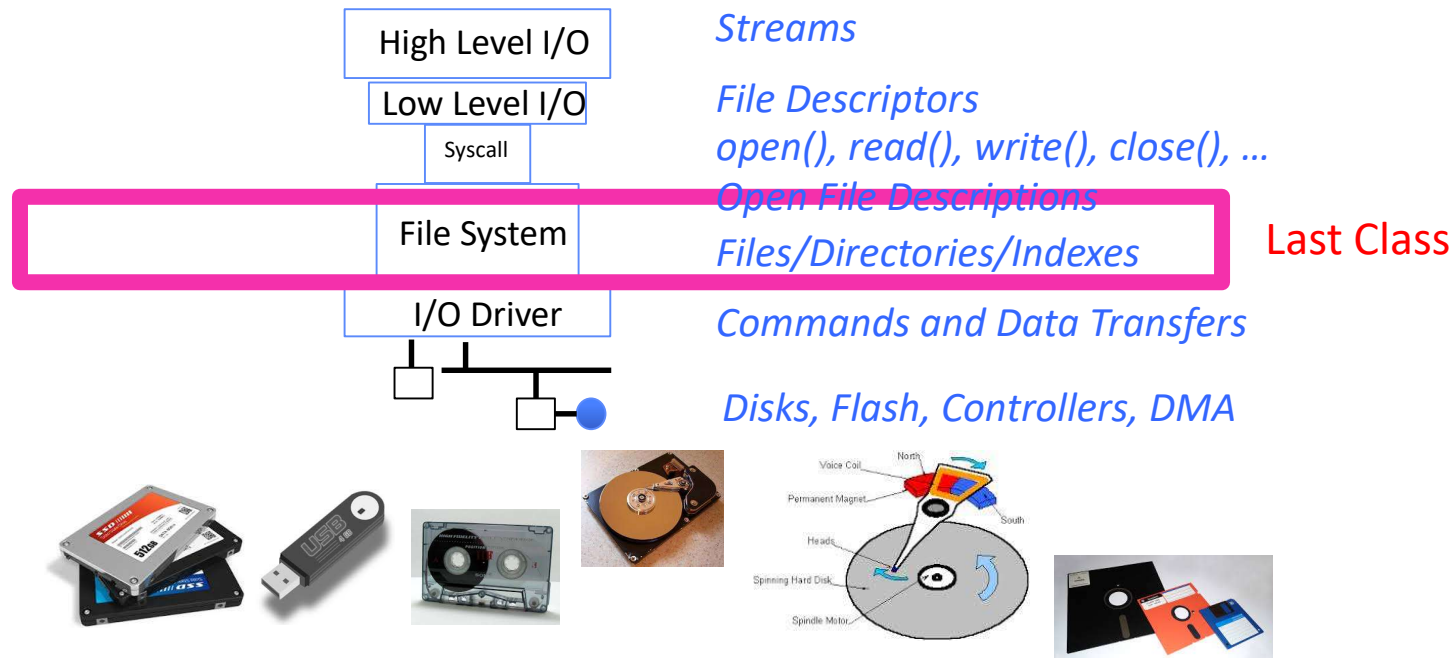
We will study..

- Where are we (File Systems)?
- C Library APIs
- Buffered I/O (File System): Introduction
- Buffered I/O File Operations

Where are we ?

I/O and Storage Layers

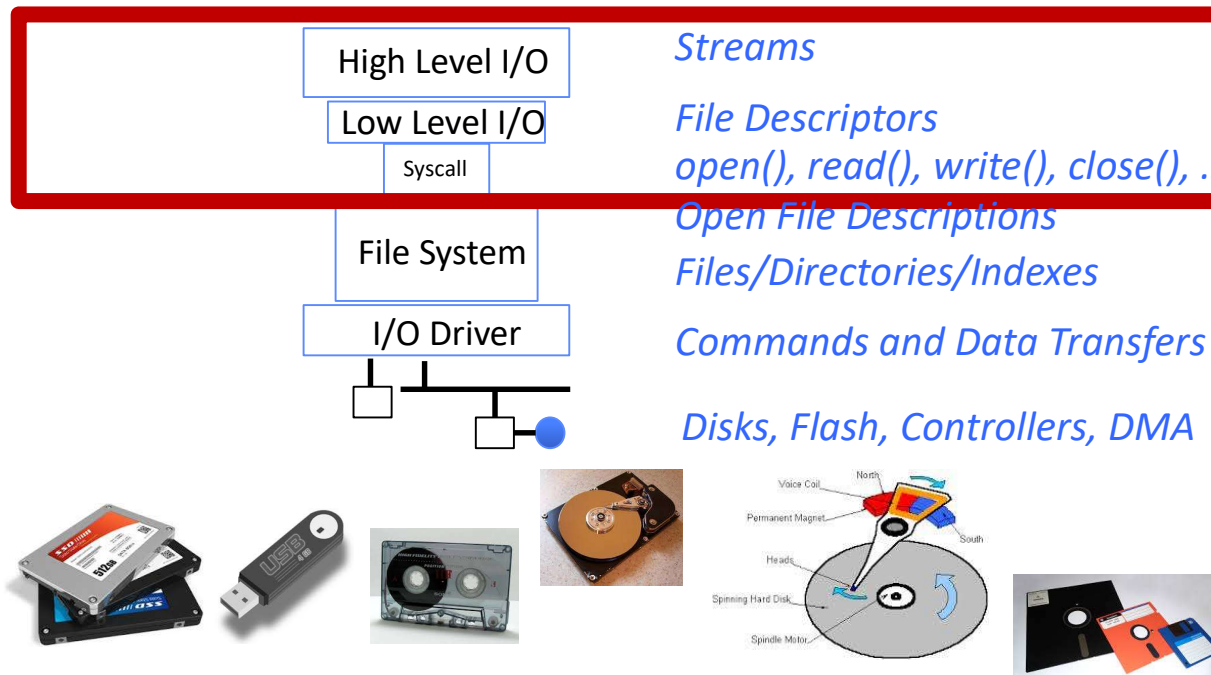
Application / Service



I/O and Storage Layers

Application / Service

We will study



File Access

- Sequential access
 - Read from the beginning
 - Can't skip around
 - Corresponds to magnetic tape
- Random access
 - Start from anywhere
 - Example: disks
 - Necessary for many applications
 - Database systems

The File Abstraction

- High-Level File I/O: Streams
- Low-Level File I/O: File Descriptors



We know

The File Abstraction

- High-Level File I/O: Streams
- Low-Level File I/O: File Descriptors



Today

C Library API

Buffered I/O - Introduction

C library High level APIs vs Linux Syscalls

- A C library function implementation
 - C interface the library provides to programmers to access **kernel related functions**
- If a C programmers directly uses file **syscalls**, he/she needs to read the **documentation**
- Each **syscall** has a **number of arguments** – makes programming complicated
 - A user, most of the time, may not use all arguments
 - User space programs are expected to find out arguments for each **syscall** by inspecting the documentation

C library High level APIs vs Linux Syscalls

- Example: `fopen()` is a library function which provides **buffered I/O services** for opening a file while `open()` is a system call that provides **non-buffered I/O services**
- Standard library file operations don't directly use **file descriptors**
- In standard I/O parlance, **an open file is called a stream**

`fopen("file", "r")` = `open("file", O_RDONLY)`

- However, Syscalls use **file descriptors**

`fd = open("file", O_RDONLY);`

User-Buffered I/O

- The *block* is an abstraction representing the smallest unit of storage on a filesystem
- Inside the kernel, all filesystem operations occur in terms of blocks
- No I/O operation may execute on an amount of data less than the block size
- If you only want to read a byte, you'll have to read a whole block
 - Even to read and modify a byte, you will have to perform operations on whole block
- Partial block operations are inefficient
- User applications don't work this way..
 - They use abstractions: byte, strings, etc - independent of block size
- Example: Reading a single byte 1024 times vs reading a single 1024 bytes once

User Buffered I/O

- To improve performance, data can be buffered internally by delaying writes and reading ahead
- In practice, blocks are usually 512, 1,024, 2,048, 4,096, or 8,192 bytes in size
- Buffers are usually multiple of block sizes: 4096 or 8192 bytes

User-Buffered I/O

- Buffering is done in the user space
 - Transparently in a library
- Writing
 - As data is written, it is stored in a **buffer** inside the program's address space
 - When the written data size reaches a set size – **buffer size** – entire buffer is written out (to disk) in a single write operation
 - Which means it's written to the underlying file descriptor
- Reading
 - Data is read (from disk) using **buffer-sized block aligned chunks**
 - Application's various sized read requests are served out from this buffer – say one byte at a time
 - When buffer is empty, another block-aligned chunk is read in
- Overall less system calls

User-Buffered I/O

- You can design and implement user buffering by hand in your own program
 - Many mission critical applications do this
- Vast majority of programs use
 - Popular standard I/O library (as a part of standard C library)
 - *iostream library* (as a part of *standard C++ library*)

Buffered I/O – File Operations

Standard I/O library – File Operations

- Standard I/O routines
 - Do not operate directly on file descriptors
 - Use their own unique identifier, known as the *file pointer*
- Inside C library, *file pointer* maps to a file descriptor
- *File pointer* is represented by a pointer to the FILE typedef
 - That is defined in <stdio.h>
- In standard I/O parlance, an open file is called a *stream*
- Streams may be opened for reading or writing or both

Opening a file



```
#include <stdio.h>
FILE *fopen( const char *path, const char *mode );
```

This function opens the file **path** with the behavior given by **mode** and associates a new stream with it.

Modes

Mode	Descriptions
Text	
r	Open existing file for reading. The stream is positioned at the start of the file
w	Open for writing. If the file exists, it is truncated to zero length. If the file doesn't exist it is created. Stream is positioned at the beginning of the file
a	Open for appending; created if does not exist. The stream is positioned at the end of the file
r+	Open existing file for reading & writing. Stream is positioned at the start of the file
w+	Open for reading & writing; truncated to zero if exists, create otherwise. The stream is positioned at the beginning of the file
a+	Open for reading & writing. Created if does not exist. Read from beginning, write as append. Read from the beginning and write from the end.

Opening a file

```
FILE *stream  
stream = fopen("/home/ravi/exam.txt", "r");  
If (!stream)  
    /*error */
```

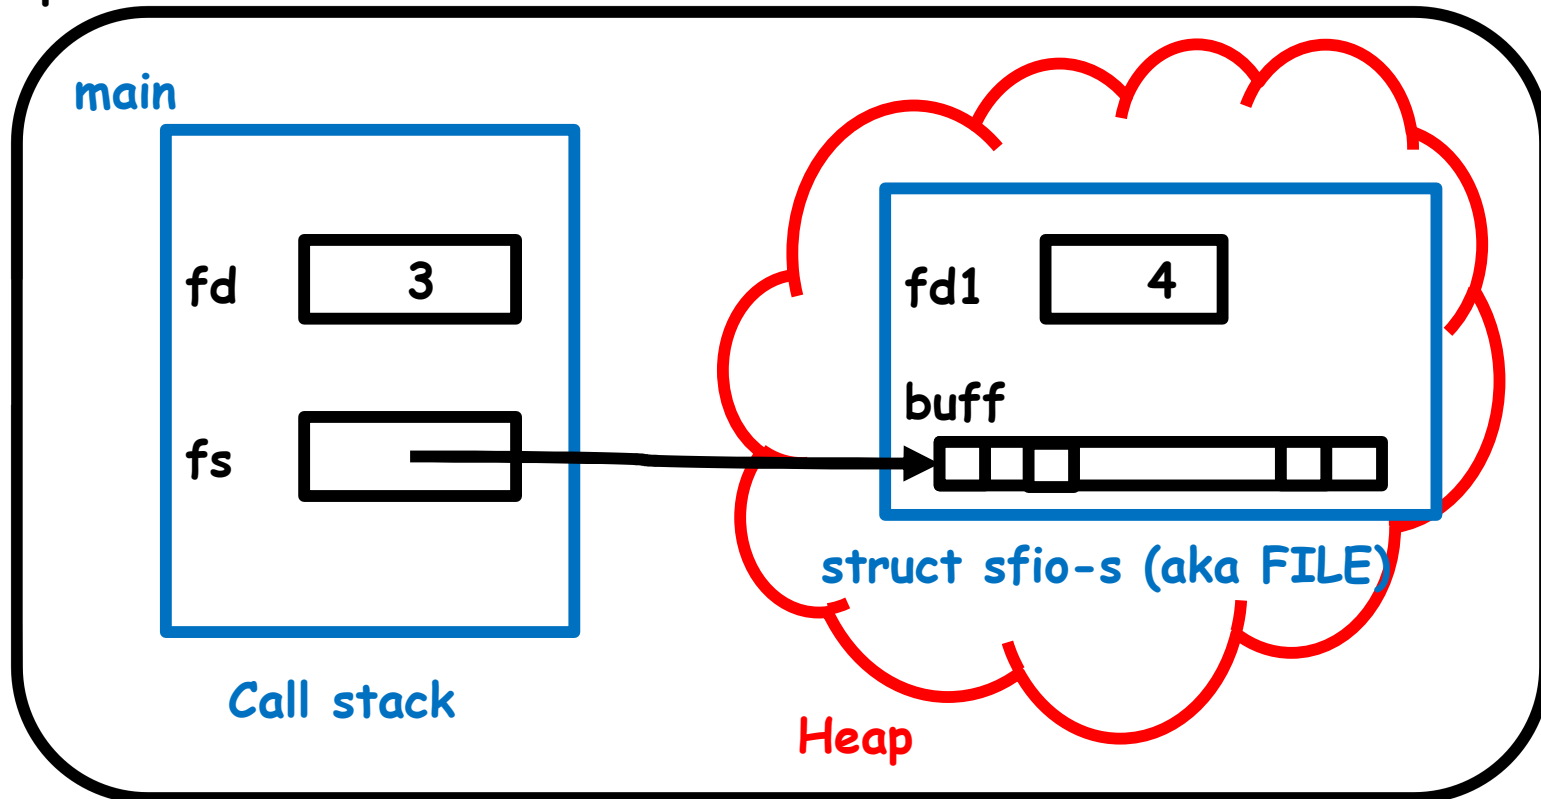
- Upon success, `fopen()` returns a valid FILE pointer
 - On failure it returns NULL and sets `errno` appropriately
 - `FILE*` represents a stream which is assigned to a variable – file pointer `stream`

read() syscall vs fread() API - User space

```
Int main(){  
    int fd = open("foo.txt", "O_RDONLY");        // for I/O syscall  
    FILE *fs = fopen("bar.txt", "w");            // for C lib I/O  
}
```

User space

Process 1101



Closing Streams

- The `fclose()` function closes a given stream:
`#include <stdio.h>`
`int fclose (FILE *stream);`
 - Any buffered and not-yet-written data is first flushed
 - On success, `fclose()` returns 0
 - On failure, it returns EOF and sets `errno` appropriately
- `fcloseall()` function closes all streams associated with the current process

C API Standard Streams

- Three predefined streams are opened implicitly when a program is executed
 - FILE `*stdin` – normal source of input, can be redirected
 - FILE `*stdout` – normal source of output, can be redirected
 - FILE `*stderr` – diagnostics and errors, can be redirected
- STDIN / STDOUT enable composition in Unix
- All can be redirected
 - `cat hello.txt | grep "World!"`
 - **cat's `stdout` goes to `grep's stdin`**

Reading Files – Multiple Functions

- Once file is open, we can read input from it
- We use one of three functions:
 - `fgetc`, `fgets`, `fread`
- `fgetc`
 - Useful when you want to read **one character** (eg letter) at a time

```
int fgetc(FILE * stream);
```

- This function reads the next character from stream
- returns it as an unsigned `char` cast to an `int`
- The return value of `fgetc()` must be stored in an `int`
- EOF is returned when end of file or error
- Reading the character advances the **current position** in the stream

Reading an entire line

- The function `fgets()` reads a string from a given stream:

```
#include <stdio.h>
```

```
char * fgets (char *str, int size, FILE *stream);
```

- This function reads up to *one less* than size bytes from stream
- Stores the results in `str`
- A null character (`\0`) is stored in the buffer after the last byte read in
- Reading stops after an EOF or a newline character is reached

Reading Binary Data

- Some developers want to read binary data such as C structures
- fread()

```
#include <stdio.h>
```

```
size_t fread (void *buf, size_t size, size_t nr, FILE *stream);
```

- Read up to **nr** elements of data, each of **size** bytes into the buffer pointed at by **buf**

Writing to a stream

- Writing a single character

```
#include <stdio.h>
int fputc (int c, FILE *stream);
```

- Writing a string of characters


```
#include <stdio.h>
int fputs (const char *str, FILE *stream);
```

- Writing Binary Data

```
#include <stdio.h>
size_t fwrite (void *buf, size_t size,
size_t nr, FILE *stream);
```

Random Access

```
int fseek(FILE *stream, long int offset, int whence);
```



- Sets file position pointer to a specific position
- *Stream*: pointer returned by `fopen`
- *Offset*: The position to seek to, relative to one of the positions specified by *whence*
- *whence*: The position from which to apply the offset; 3 positions
 - **SEEK_SET** – seek starts at beginning of file
 - **SEEK_CUR** – seek starts at current location in file
 - **SEEK_END** – seek starts at end of file

Many more – Stream Ops

```
// formatted
int fprintf(FILE *restrict stream, const char *restrict format,
            ...);
int fscanf(FILE *restrict stream, const char *restrict format,
            ...);
```

A sample program – character oriented streaming

```
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <ctype.h>
5  int main(int argc, char **argv) {
6
7  FILE* f = fopen("p1.c", "r");
8  int c;
9  int letters = 0;
10 while ( (c = fgetc(f)) != EOF){
11     if (isalpha(c)){
12         letters++;
13     }
14 }
15 printf("%s  has %d letters in it \n", "p1.c", letters );
16
17 }
18
```


Program using Syscalls

pfile2.c

```
2 // Creating a sequential file
3 #include <stdio.h>
4
5 int main(void){
6     FILE *cfPtr = NULL; // cfPtr = clients.txt file pointer
7
8     // fopen opens the file. Exit the program if unable to create the file
9     if ((cfPtr = fopen("clients.txt", "w")) == NULL) {
10         puts("File could not be opened");
11     }
12     else {
13         puts("Enter the account, name, and balance.");
14         puts("Enter EOF to end input.");
15         printf("%s", "? ");
16
17         int account = 0; // account number
18         char name[30] = ""; // account name
19         double balance = 0.0; // account balance
20
21         scanf("%d%29s%lf", &account, name, &balance);
22
23         // write account, name and balance into file with fprintf
24         while (!feof(stdin)) {
25             fprintf(cfPtr, "%d %s %.2f\n", account, name, balance);
26             printf("%s", "? ");
27             scanf("%d%29s%lf", &account, name, &balance);
28         }
29
30         fclose(cfPtr); // fclose closes file
31     }
32 }
```

```
(base) Ravis-MacBook-Pro-2:cp ravimittal$ ./pf
Enter the account, name, and balance.
Enter EOF to end input.
? 10 ravi 10.0
? 200 ram 50.50
? 300 sam 20.0
```

\$ cat clients.txt

```
10 ravi 10.00
200 ram 50.50
300 sam 20.00
```

EOF character

Linux/MAC OS : <Ctrl> d

Windows: <Ctrl> z enter

Lecture Summary

- Standard I/O is a user-buffering library provided as part of the standard C library
- Buffered file operations are useful when
 - You issue many system calls
 - Performance is crucial
 - Your access patterns are character- or line-based
 - You want interfaces to make such access easy without issuing extraneous system calls
 - You prefer a higher-level interface to the low-level Linux system calls