

CS310 Operating Systems

Lecture 44: Distributed System

Ravi Mittal

IIT Goa

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, University of California, Berkeley
 - Book: Computer Networking: A top-down Approach, by Kurose, and Ross
 - 15-440, Distributed Systems, Stanford University, Lecture 6, Class Notes
 - Book: Modern Operating System by Tanenbaum and Bos

Previous Classes

Today we will study

- Local Storage → Cloud Storage
- Socket Overview
- RPC Overview
- Distributed File System - Concepts

Local Storage → Cloud Storage

Access to Storage (Files) Today


- File system over SSD or HDD on your local machine
- File Server in your organization
 - Remote login (ssh), file transfer (scp) or mount
- Cloud storage
 - Accessed through web or app (drive, box, ...)
 - Mounted on your local machine
 - Replicated and/or Distributed

Cloud Storage Options

- Storage Account / Share is like disk “partition”
 - Holds file system: directory, index, free map, data blocks
- Access methods: mount, REST, file transfer, synch
- Security: credentials, encryption
- Performance: HDDs, SSDs
- Redundancy
 - Local RAID
 - Storage cluster in a Data Center
 - Zone redundant (across data centers)
 - Geographic regions


Geographic Replication: Cluster, Zone, Geo

- Highly durable: Hard to destroy all copies
- Highly available for reads: Just talk to any copy
- What about for writes? Need every copy online to update all together?




Centralized vs. Distributed

- **Centralized System**
 - Major functions performed on one physical computer
- **Distributed System**
 - Physically separate computers working together to perform a single task



Peer-to-Peer Model



Client/Server Model

Distributed Systems: Motivation

- The *promise* of distributed systems
 - *Higher availability*: one machine goes down, use another
 - *Better durability*: store data in multiple locations
 - *More security*: each piece easier to make secure
- Other advantages
 - Cheaper/easier to build lots of simple computers
 - Allows for adding more resources incrementally
 - Easier for users to collaborate

Distributed Systems: Goals/Requirements

- **Transparency:** the ability of the system to mask its complexity behind a simple interface
- Possible transparencies:
 - **Location:** Can't tell where resources are located
 - **Migration:** Resources may move without the user knowing
 - **Replication:** Can't tell how many copies of resource exist
 - **Concurrency:** Can't tell how many users there are
 - **Parallelism:** System may speed up large jobs by splitting them into smaller pieces
 - **Fault Tolerance:** System may hide various things that go wrong

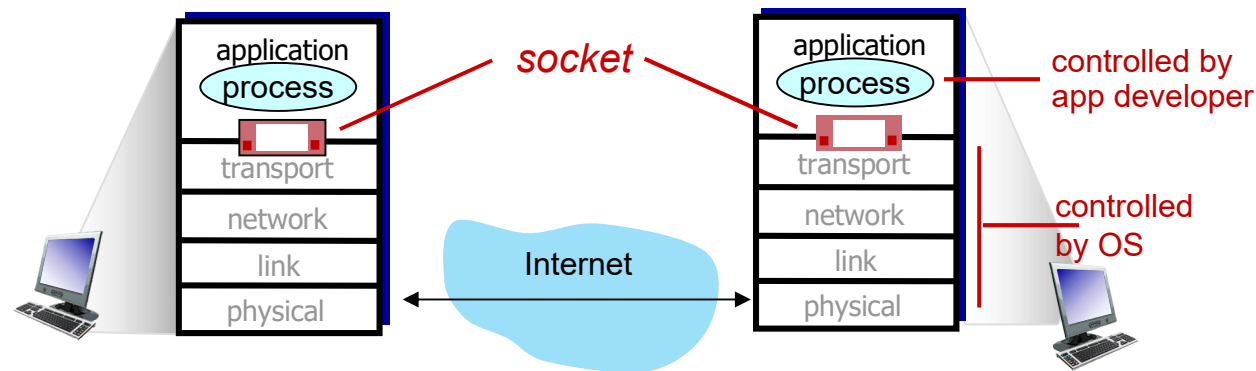
Examples of Transparency

- **RPC: Simple function-like interface**
 - Masks complexity of marshalling/unmarshalling, sending data, using sockets...
- **Sockets: Simple file-like interface**
 - Masks complexity of segmentation, retransmissions, windowing, etc.

Socket Overview

Sockets

- Process sends/receives messages to/from its socket
- A Socket is analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
 - two sockets involved: one on each side




Sockets

- **Socket**: an abstraction of a network I/O queue
 - Embodies one side of a communication channel
 - Same interface regardless of location of other end
 - Could be local machine (called “UNIX socket”) or remote machine (called “network socket”)
 - First introduced in 4.2 BSD UNIX: big innovation at time

Socket: Identify the Destination


- Addressing
 - IP address
- Multiplexing
 - port



Use of Sockets

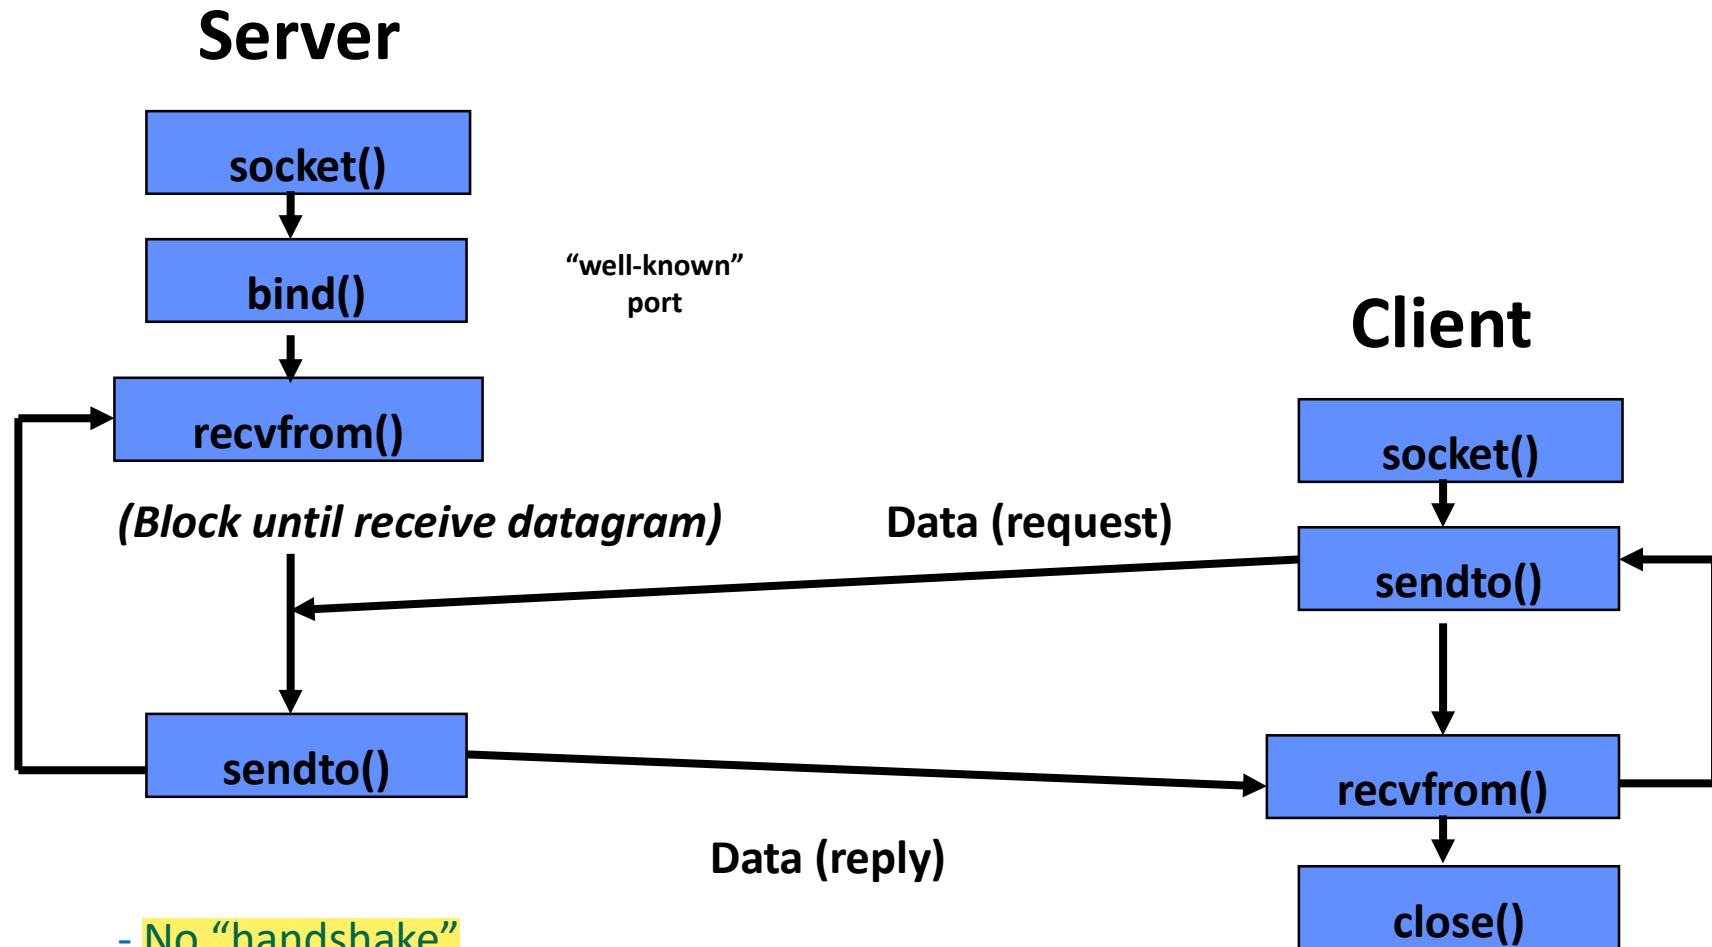
- How to use sockets
 - Setup socket
 - Where is the remote machine (IP address, hostname)
 - What service gets the data (port)
 - Send and Receive
 - Designed just like any other I/O in unix
 - send -- write
 - recv -- read
 - Close the socket
- Sockets
 - UDP Client Server
 - TCP Client Server

Recall: Socket Setup over TCP/IP



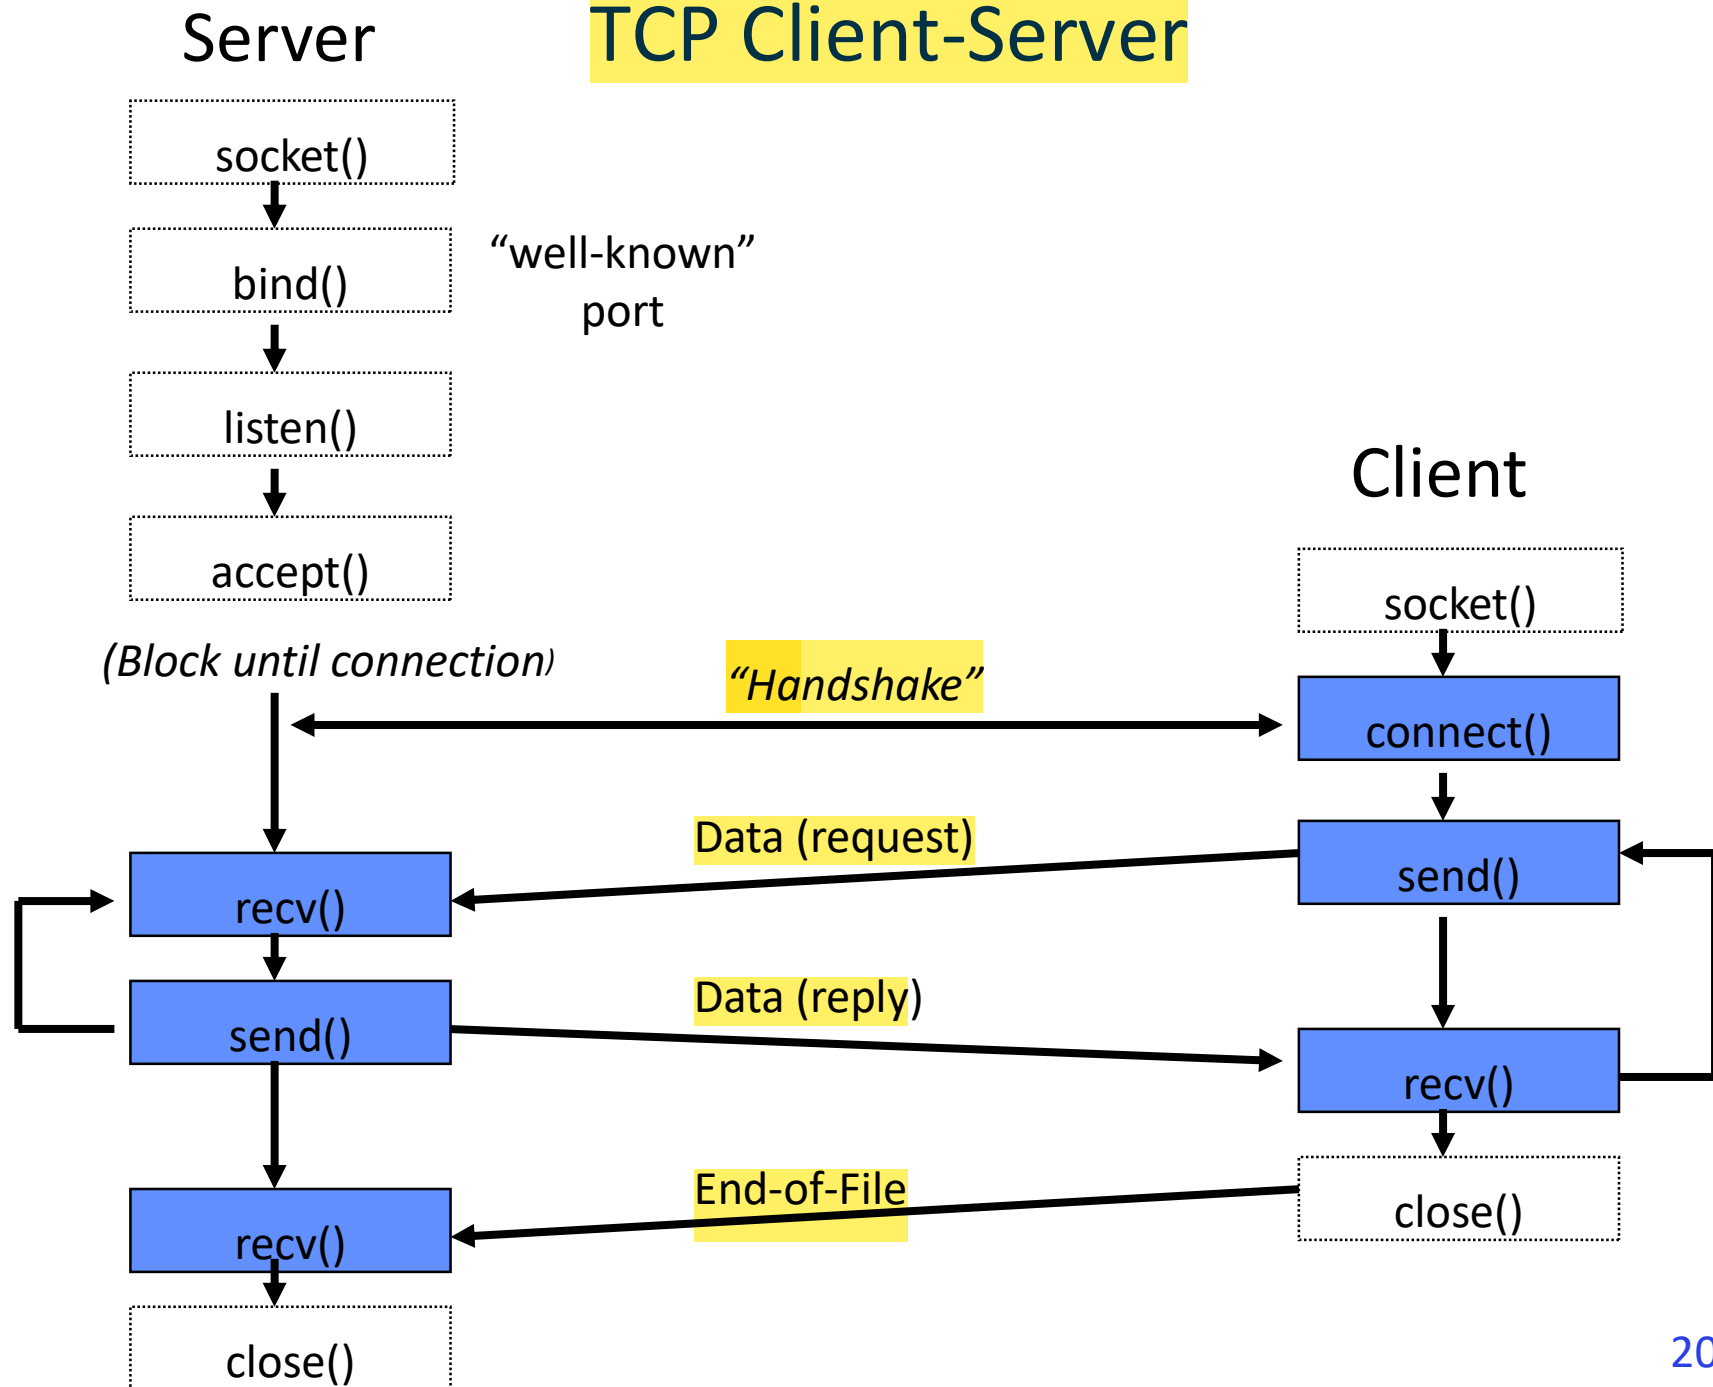
- Things to remember:
 - Connection involves 5 values:
[Client Addr, Client Port, Server Addr, Server Port, Protocol]
 - Often, Client Port “randomly” assigned
 - Server Port often “well known”
 - 80 (web), 443 (secure web), 25 (sendmail), etc
 - Well-known ports from 0—1023

UDP Client-Server (self reading)



- No "handshake"
- No simultaneous close
- No fork()/spawn() for concurrent servers!

TCP Client-Server



RPC Overview

RPC: Concept

- A type of client/server communication
- Attempts to make remote procedure calls look like local ones

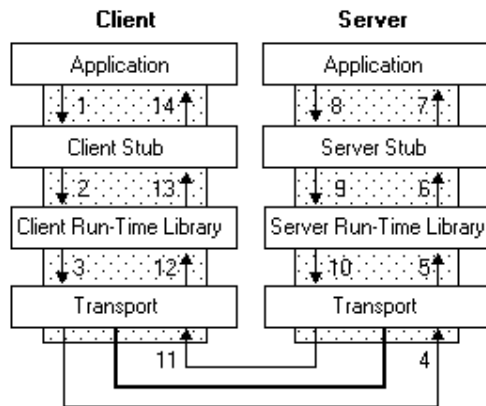


figure from Microsoft
MSDN

```
{ ...  
    foo()  
}  
  
void foo() {  
    invoke_remote_foo()  
}
```

Remote procedure call

- A remote procedure call makes a call to a remote service look like a local call
 - RPC makes transparent whether server is local or remote
 - RPC allows applications to become distributed transparently
 - RPC makes architecture of remote machine transparent


But it's not always simple

- Calling and called procedures run on different machines, with different address spaces
 - And perhaps different environments .. or operating systems ..
- Must convert to local representation of data
- Machines and network can fail

Stubs: obtaining transparency

- Compiler generates from API stubs for a procedure on the client and server
- Client stub
 - Marshals arguments into machine-independent format
 - Sends request to server
 - Waits for response
 - Unmarshals result and returns to caller
- Server stub
 - Unmarshals arguments and builds stack frame
 - Calls procedure
 - Server stub marshals results and sends reply

RPC Information Flow



RPC steps

- Step 1: Client calling client stub
 - Normal procedure call with parameters
- Step 2: the client stub packing the parameters in the message and making a system call to send the message
 - Packing parameters is called marshalling
- Step 3: Kernel sending the message from the client machine to the server machine
- Step 4: the Kernel passing the incoming packet to Server stub
- Step 5: the Server stub calling the server procedure
- Reply traces the same path

A few points !

- Client stub represents server procedure in client address space
- Client procedure is a normal procedure which calls client stub
 - Client stub has the same name as the server procedure
- Is Address Space an issue ?
 - Client procedure and client stub are in the same address space – parameters are passed in the usual way
 - The server procedure and server stub are in the same address space
- How to pass a pointer (as a parameter) to a remote procedure ?
- How to handle Global variables?

RPC Details (self study) – self study

- Equivalence with regular procedure call
 - Parameters \Leftrightarrow Request Message
 - Result \Leftrightarrow Reply message
 - Name of Procedure: Passed in request message
 - Return Address: mbox2 (client return mail box)
- Stub generator: Compiler that generates stubs
 - Input: interface definitions in an “interface definition language (IDL)”
 - Contains, among other things, types of arguments/return
 - Output: stub code in the appropriate source language
 - Code for client to pack message, send it off, wait for result, unpack result and return to caller
 - Code for server to unpack message, call procedure, pack results, send them off

RPC Details (2/3) – self study

- Cross-platform issues:
 - What if client/server machines are different architectures/languages?
 - Convert everything to/from some canonical form
 - Tag every item with an indication of how it is encoded (avoids unnecessary conversions)
- How does client know which mbox to send to?
 - Need to translate name of remote service into network endpoint (Remote machine, port, possibly other info)
 - **Binding**: the process of converting a user-visible name into a network endpoint
 - This is another word for “naming” at network level
 - Static: fixed at compile time
 - Dynamic: performed at runtime

RPC Details (3/3) – self study

- Dynamic Binding
 - Most RPC systems use dynamic binding via name service
 - Name service provides dynamic translation of service → mbox
 - Why dynamic binding?
 - Access control: check who is permitted to access service
 - Fail-over: If server fails, use a different one
- What if there are multiple servers?
 - Could give flexibility at binding time
 - Choose unloaded server for each new client
 - Could provide same mbox (router level redirect)
 - Choose unloaded server for each new request
 - Only works if no state carried from one call to next
- What if multiple clients?
 - Pass pointer to client-specific return mbox in request

Distributed File System

Distributed System Protocols are Built by Message Passing


- Sending/receiving messages is **atomic**
 - Each message is either fully received exactly once...
 - or not received at all (!)
- Interface:
 - **Mailbox**: temporary holding area for messages
 - Includes both destination location and queue
 - **Send(message, mbox)**
 - Send message to remote mailbox identified by mbox
 - **Receive(buffer, mbox)**
 - Wait until mbox has message, copy into buffer, and return
 - If threads sleeping on this mbox, wake up one of them

But, doesn't TCP give us reliable delivery?

- TCP provides a convenient interface to use an unreliable network...
- ... but it does *not* make the network reliable!
- Messages can still be lost if you use TCP
 - After many retransmissions, the OS “gives up” and breaks the connection
- Losing messages is fundamental problem in distributed systems
 - TCP's retransmissions turn packet losses into packet delays (even if it never “gives up”)
 - ***And very long delays look just like losses!***
 - TCP makes the network easy to use, and it can help improve performance
 - But TCP doesn't solve this fundamental problem (losing messages)

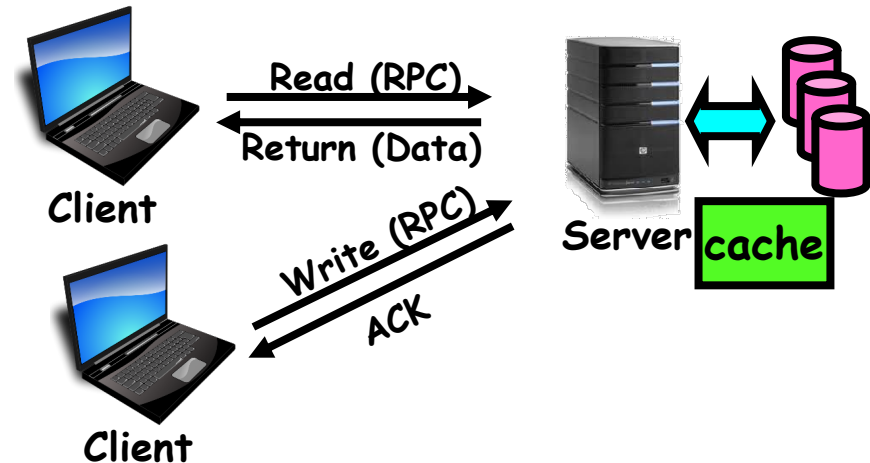
Distributed File Systems

- Transparent access to files stored on a remote disk
- *Mount* remote files into your local file system
 - Directory in local file system refers to remote files
 - e.g., `/home/oksi/162/` on laptop actually refers to `/users/oski` on campus file server



Simple Distributed File System

- Remote Disk: Opens, Reads, Writes, Closes forwarded to server
 - Use Remote Procedure Calls (RPC) to translate file system calls into remote requests
 - Server may cache files in memory to response more quickly
 - Server provides consistent view of file system to multiple clients
- Problem: performance (network slower than memory, server is bottleneck)



Lecture Summary

- Distributed systems are becoming very important in the present interconnected world
- Remote Procedure Call (RPC): Call procedure on remote machine
 - Provides same interface as procedure
 - Automatic packing and unpacking of arguments without user programming (in stub)
- Socket Programming
 - Connection involves 5 values:
[Client Addr, Client Port, Server Addr, Server Port, Protocol]